

SYMPHONY 5.6.18 User's Manual ¹

T.K. Ralphs²

M. Güzelsoy³

A. Mahajan⁴

January 15, 2022

¹This research was partially supported by NSF Grants CMMI-1435453, CMMI-0728011, DMI-0522796, DMI-0534862, DMS-9527124, CMMI-1130914, and Texas ATP Grant 97-3604-010. A revised version of Chapters 4 of this manual now appears in the Springer-Verlag book *Computational Combinatorial Optimization* edited by M. Jünger and D. Naddef, see <http://link.springer.de/link/service/series/0558/tocs/t2241.htm>

²Department of Industrial and Systems Engineering, Lehigh University, Bethlehem, PA 18017, ted@lehigh.edu, <http://coral.ielehigh.edu/~ted>

³SAS Institute, Cary, NC mguzelsoy@gmail.com

⁴Department of Industrial Engineering and Operations Research, IIT Bombay, India amahajan@iitb.ac.in, <http://www.ieor.iitb.ac.in/amahajan>

©2000-2022 Ted Ralphs

Acknowledgments

Many thanks are due to Laci Ladányi who worked with me on the development of a very early precursor of SYMPHONY called COMPSys many years ago now and who taught me much of what I then knew about programming. Thanks are due also to Marta Esö, who wrote an early draft of this manual for what was then COMPSys. Since those early days, many of my Ph.D students have worked on SYMPHONY and made substantial contributions. Menal Güzelsoy was instrumental in the development of SYMPHONY since version 4.0 and Ashutosh Mahajan, who has worked on SYMPHONY since version 5.0. In particular, Ashutosh and Menal did all of the work that went into improving SYMPHONY for release 5.2. I would also like to thank Matthew Galati and Ondrej Medek, who contributed to the development of SYMPHONY over the years, as well as Anahita Hassanzadeh and Suresh Bolusani. Finally, my sincere appreciation goes to Leslie Trotter, William Cook, Cornell University, Rice University, Lehigh University, the National Science Foundation, and the state of Texas, all of whom have supported the development of SYMPHONY over the past almost three decades.

Contents

1	Introduction	1
1.1	Introducing SYMPHONY 5.6.18	1
1.2	Features	1
1.3	A Brief History	3
1.4	Related Work	4
1.5	How to Use This Manual	4
1.6	Getting Additional Help	5
2	Installing SYMPHONY	7
2.1	Installing the Binary Distribution	7
2.1.1	Installation in Unix-like environments	8
2.1.2	Installation for Use With Microsoft Visual C++	9
2.2	Building From Source	9
2.2.1	External Dependencies	9
2.2.2	Building in Unix-like environments	10
2.2.3	Building Using Microsoft Visual C++	18
3	Using SYMPHONY	23
3.1	Using SYMPHONY Interactively	23
3.1.1	Unix-like Environments	23
3.1.2	Microsoft Windows	23
3.1.3	Main Menu	24
3.1.4	Set Menu	26

3.1.5	Display Menu	27
3.1.6	Sub Menu Browsing	28
3.2	Using SYMPHONY from the Command Line	29
3.3	Using the Callable Library	30
3.3.1	The C API	30
3.3.2	The C++ API	33
3.3.3	Linking to the Callable Library	34
3.4	Using the Callback Functions	35
4	Technical Details	37
4.1	Branch and Bound	37
4.2	Branch and Cut	37
4.3	Design of SYMPHONY	39
4.3.1	An Object-oriented Approach	40
4.3.2	Data Structures and Storage	41
4.3.3	Modular Implementation	43
4.3.4	Algorithm Summary	46
4.4	Details of the Implementation	47
4.4.1	The Master Module	47
4.4.2	The Node Processing Module	50
4.4.3	The Tree Management Module	54
4.4.4	The Cut Generation Module	55
4.4.5	The Cut Management Module	55
4.5	Parallelizing BCP	56
4.5.1	Parallel Configurations	57
4.5.2	Inter-process Communication	57
4.5.3	Fault Tolerance	58
5	Developing Custom Applications	59
5.1	Navigating the Source Code	59
5.2	Building an Application	60

5.2.1	Unix	60
5.2.2	Microsoft Visual C++	60
5.3	Writing the Callbacks	61
5.4	Data Structures	62
5.5	Parallel Implementation	63
5.5.1	Distributed-memory Architectures	63
5.5.2	Shared-memory Architectures	63
5.6	Debugging Your Application	64
5.6.1	The First Rule	64
5.6.2	Debugging with PVM	64
5.6.3	Checking the Validity of Cuts and Tracing the Optimal Path	64
5.6.4	Using the Interactive Graph Drawing Software	65
5.6.5	Other Debugging Techniques	65
5.7	Case Study: Implementing a Matching Solver	66
6	Reference	73
6.1	Callable Library C API	73
6.1.1	Primary Interface Functions	74
6.1.2	Parameter Query and Modification	89
6.1.3	Solver Status Query Functions	96
6.1.4	Data Query Functions	102
6.1.5	Data Modification Functions	128
6.1.6	Warm Starting Functions	145
6.1.7	Sensitivity Analysis Functions	151
6.2	Callable Library C++ API	155
6.3	User Callback API	158
6.3.1	Master module callbacks	158
6.3.2	LP module callbacks	174
6.3.3	Cut generator module callbacks	208
6.3.4	Cut pool module callbacks	214

6.3.5	Draw graph module callbacks	220
6.4	Run-time Parameters	224
6.4.1	Global parameters	224
6.4.2	Master module parameters	224
6.4.3	Draw Graph parameters	227
6.4.4	Tree Manager parameters	227
6.4.5	LP parameters	231
6.4.6	Cut Generator Parameters	238
6.4.7	Cut Pool Parameters	238
6.4.8	C++ Interface/OSI Parameters	239
	Bibliography	241

Chapter 1

Introduction

1.1 Introducing SYMPHONY 5.6.18

Welcome to the SYMPHONY Version 5.6.18 user's manual. Whether you are a new user or simply upgrading, this manual will help you get started with what we hope you will find to be a useful and powerful framework for solving mixed-integer linear programs (MILP) sequentially or in parallel. The subroutines in the SYMPHONY library comprise a state-of-the-art MILP solver with a modular design that makes it easy to customize for various problem settings. SYMPHONY works out of the box as a generic MILP solver that can be invoked from the command line, through an interactive shell, or by linking to the provided callable library, which has both C and C++ interfaces with a look and feel similar to that of other popular solvers (see Sections 6.1 and 6.2 for the library routines). Models can be read in MPS or GMPL (a subset of AMPL) format, as well as by interfacing with more powerful modeling environments, such as FlopC++ (also provided with the distribution). To develop a customized SYMPHONY application, various callbacks can be written and parameters set that modify the default behavior of the algorithm. Section 3.4 contains an overview of the API for these subroutines. Files containing function stubs are provided with the distribution.

SYMPHONY can be built on almost any platform and can be configured either for serial computation or in a wide variety of parallel modes. The parallel version can be built for either a fully distributed environment (network of workstations) or a shared-memory environment simply by changing a few configuration options (see Chapter 2). To run in a distributed environment, the user must have installed the *Parallel Virtual Machine* (PVM), available for free from Oak Ridge National Laboratories. To run in a shared-memory environment, the user must have installed an OpenMP compliant compiler (gcc 4.2 is currently the only compiler tested and fully supported).

1.2 Features

Starting in SYMPHONY 5.0, we introduced a number of new features that give SYMPHONY some unique capabilities. These include the ability to solve biobjective integer programs, the ability to

warms start the solution procedure, and the ability to perform basic sensitivity analyses. These capabilities have been further developed and enhanced with the introduction of Versions 5.1–5.6. Other new features and enhancements are listed below.

- SYMPHONY has an interactive optimizer that can be used through a command shell. In both the sequential and parallel configurations, the user can set parameters, load and solve instances interactively, and display results and statistics. For Windows users, this means that SYMPHONY can be invoked using the familiar procedure of “double-clicking” on the `symphony.exe` file in Windows Explorer.
- SYMPHONY supports automatic configuration using the new COIN-OR build system and the GNU autotools. Using the autotools, it is now possible to build SYMPHONY in most operating systems and with most common compilers without user intervention.
- Both the distributed and shared memory parallel configurations are fully implemented, tested, and supported. The user can now build and execute custom SYMPHONY applications in parallel, as well as solving generic MILPs in parallel “out of the box.”
- There are additional options for warm starting. The user can trim the warm starting tree before starting to resolve a problem. More specifically, the user can decide to initiate warm starting with a predefined partition of the final branch-and-cut tree resulting from a previous solution procedure. This partition can include either a number of nodes created first during the solution procedure or all of the nodes above a given level of the tree.
- The COIN-OR repository, the current host of SYMPHONY has recently undergone some significant improvements of its own that have resulted in improved services to users, detailed below.
 - SYMPHONY has a project management Web site, where users can submit trouble tickets, browse the source code interactively, and get up-to-date information on development. The address of the new site is <https://projects.coin-or.org/SYMPHONY>.
 - SYMPHONY is hosted using subversion, a version control system with features vastly improved over CVS, the previous hosting software. This has required some reorganization and renaming of the header files.
 - SYMPHONY is tightly integrated with other COIN-OR projects. Due to improved procedures for producing stable releases, it will now be much easier for us to determine the exact version of SYMPHONY and all other COIN projects you are using when you report a bug.
 - SYMPHONY is distributed with all COIN software needed to build a complete solver. Previously, other COIN software packages had to be downloaded and installed separately.

Two features have been deprecated and are no longer supported:

- The native interfaces to OSL and CPLEX are now deprecated and no longer supported. These solvers can be called through the COIN-OR OSI interface.

- Column generation functionality has also been officially deprecated. For now, there are a number of other software packages that offer better functionality than SYMPHONY for implementing branch and price algorithms.

For what's new specifically in Version 5.6.18, please check the `README` file that comes with the distribution.

1.3 A Brief History

Since the inception of optimization as a recognized field of study in mathematics, researchers have been both intrigued and stymied by the difficulty of solving many of the most interesting classes of discrete optimization problems. Even combinatorial problems, though conceptually easy to model as integer programs, have long remained challenging to solve in practice. The last two decades have seen tremendous progress in our ability to solve large-scale discrete optimization problems. These advances have culminated in the approach that we now call *branch and cut*, a technique (see [19, 29, 20]) which brings the computational tools of branch and bound algorithms together with the theoretical tools of polyhedral combinatorics. Indeed, in 1998, Applegate, Bixby, Chvátal, and Cook used this technique to solve a *Traveling Salesman Problem* instance with 13,509 cities, a full order of magnitude larger than what had been possible just a decade earlier [2] and two orders of magnitude larger than the largest problem that had been solved up until 1978. This feat becomes even more impressive when one realizes that the number of variables in the standard formulation for this problem is approximately the *square* of the number of cities. Hence, we are talking about solving a problem with roughly *100 million variables*.

There are several reasons for this impressive progress. Perhaps the most important is the dramatic increase in available computing power over the last decade, both in terms of processor speed and memory. This increase in the power of hardware has subsequently facilitated the development of increasingly sophisticated software for optimization, built on a wealth of theoretical results. As software development has become a central theme of optimization research efforts, many theoretical results have been “re-discovered” in light of their new-found computational importance. Finally, the use of parallel computing has allowed researchers to further leverage their gains.

Because of the rapidly increasing sophistication of computational techniques, one of the main difficulties faced by researchers who wish to apply these techniques is the level of effort required to develop an efficient implementation. The inherent need for incorporating problem-dependent methods (most notably for dynamic generation of variables and cutting planes) has typically required the time-consuming development of custom implementations. Around 1993, this led to the development by two independent research groups of software libraries aimed at providing a generic framework that users could easily customize for use in a particular problem setting. One of these groups, headed by Jünger and Thienel, eventually produced ABACUS (A Branch And CUt System) [21], while the other, headed by the authors, produced what was then known as COMPSys (Combinatorial Optimization Multi-processing System). After several revisions to enable more broad functionality, COMPSys became SYMPHONY (Single- or Multi-Process Optimization over Networks). A version of SYMPHONY written in C++, which we call COIN/BCP has also been produced at IBM under the COIN-OR project [24]. The COIN/BCP package takes substantially the

same approach and has the same functionality as SYMPHONY, but has extended SYMPHONY's capabilities in some areas.

1.4 Related Work

The 1990's witnessed a broad development of software for discrete optimization. Almost without exception, these new software packages were based on the techniques of branch, cut, and price. The packages fell into two main categories—those based on general-purpose algorithms for solving mixed-integer linear programs (MILPs) (without the use of special structure) and those facilitating the use of special structure by interfacing with user-supplied, problem-specific subroutines. We will call packages in this second category *frameworks*. There have also been numerous special-purpose codes developed for use in particular problem settings.

Of the two categories, MILP solvers are the most common. Among the dozens of offerings in this category are MINTO [27], MIPO [3], bc-opt [8], and SIP [26]. Generic frameworks, on the other hand, are far less numerous. The three frameworks we have already mentioned (SYMPHONY, ABACUS, and COIN/BCP) are the most full-featured packages available. Several others, such as MINTO, originated as MILP solvers but have the capability of utilizing problem-specific subroutines. CONCORDE [2, 1], a package for solving the *Traveling Salesman Problem* (TSP), also deserves mention as the most sophisticated special-purpose code developed to date.

Other related software includes several frameworks for implementing parallel branch and bound. Frameworks for general parallel branch and bound include PUBB [33], BoB [4], PPBB-Lib [35], and PICO [10]. PARINO [23] and FATCOP [6] are parallel MILP solvers.

1.5 How to Use This Manual

The manual is divided into six chapters. The first is the introduction, which you are reading now. Chapter 2 describes how to install SYMPHONY from either a source or binary distribution. If you have already managed to get SYMPHONY running using the instructions in the [README](#) file, you might want to skip to Chapter 3. However, keep in mind that the manual contains additional details for customizing your build. Chapter 3 contains an overview of how to use in all three major modes—as a black-box solver through the interactive shell or on the command line, as a callable library, and as a customizable framework. Chapter 4 contains further depth and a more complete technical description of the design and implementation of SYMPHONY. In Section 4.3, we describe the overall design of SYMPHONY without reference to the implementational details and with only passing reference to parallelism. In Section 4.4, we discuss the details of the implementation. In Section 4.5, we briefly discuss issues involved in parallel execution of SYMPHONY. Chapter 5 describes in detail how to develop a custom application using SYMPHONY. Note that it is not necessary to read Chapter 4 before undertaking development of a SYMPHONY application, but it may help. Chapter 6 contains reference material. Section 6.1 contains a description of the native C interface for the callable library. Section 6.2 contains a description of the interface for C++ environments. Section 6.3 contains a description of the user callback functions. SYMPHONY's parameters are described in Section 6.4. For reference use, the HTML version of this manual may

be more practical, as the embedded hyperlinks make it easier to navigate.

1.6 Getting Additional Help

The main point of entry for additional help, trouble-shooting, and problem-solving is the SYMPHONY Wiki and development Web site at

`https://github.com/coin-or/SYMPHONY`

There, bug reports can be submitted, previous bug reports searched, and general questions asked.

Chapter 2

Installing SYMPHONY

This chapter is concerned with detailed instructions for building and installing SYMPHONY, along with its associated libraries and applications. Note that the procedures outlined below may change over time, so please also check the current README documentation for the COIN-OR Optimization suite at

https://coin-or.github.io/user_introduction.html

documentation of the `coinbrew` build script (referenced below) at

<https://coin-or.github.io/coinbrew>.

2.1 Installing the Binary Distribution

For the users who only need to use SYMPHONY as a black-box solver or only need the SYMPHONY callable library in order to develop a custom application, there are binary distributions released for different compilers and platforms. Each distribution consists of an executable and libraries built from the source code of SYMPHONY version 5.6.18. It is important to note, however, that the pre-compiled binaries are missing some very useful functionality because of the fact that we are unable to distribute code linked to libraries licensed under the GNU General Public License (GPL). There are also a number of other potentially useful configurations (such as the parallel version) that we do not distribute in binary form. Building from source should be easy in most environments and will give you more flexibility. See Section 2.2 for more information on additional functionality available when building from source.

Binaries for many platforms are available as assets attached to releases on Github. For the current release, the URL is

<https://github.com/coin-or/SYMPHONY/releases/tag/releases%2F5.6.18>

In many cases, however, binaries can be installed more easily with a package manager.

- **Linux:** On Debian/Ubuntu, SYMPHONY is available in the package `coinor-symphony` and can be installed with apt. On Fedora, SYMPHONY is available in the package `coin-or-SYMPHONY`.
- **Mac OS X:** The easiest way to get SYMPHONY on Mac OS X is through Homebrew (<https://brew.sh>).
 - `brew tap coin-or-tools/coinor`
 - `brew install coin-or-tools/coinor/symphony`

2.1.1 Installation in Unix-like environments

The installation consists of unpacking the distribution into an appropriate installation directory, such as `/usr/local/symphony`. Unpack the distribution with

```
tar xzvf SYMPHONY-releases.\VER-XXX.tgz
```

where `XXX` indicates the platform and the version of compiler used to build the distribution. Switch into the root directory of the unpacked distribution.

Test the executable by going to the `bin` directory and typing

```
./symphony
```

This will start up the solver in interactive mode. Type `help` or `?` to see a list of available commands or see Chapter 3 for instructions on using the interactive solver.

To test the callable library, the distribution includes sample files in `examples` directory:

milp.c: This sample code is an implementation of a basic MILP solver using SYMPHONY's C callable functions with user defined input (see Section 3.3). To test the code, go to the `examples` directory and type

```
make milp
milp
```

You may need to customize some paths in the Makefile and to set the environment variable `LD_LIBRARY_PATH`. **milpOsi.c:** This sample code is an implementation of a basic MILP solver using SYMPHONY's C++ callable functions (through OsiSym interface) with user defined input (see Section 3.3.2). To test the code, go to `examples` directory and type,

```
make milpOsi
milpOsi
```

If everything is working properly, the libraries, executables and header files can be installed in appropriate system directories if desired. This must be done by hand. This step will be done automatically if building from source using the automatic scripts provided (see below). It is important to note, though, that

2.1.2 Installation for Use With Microsoft Visual C++

Download and unpack the archive `SYMPHONY-releases.5.6.18-XXX.zip` in an appropriate installation directory, such as `C:\Program Files (x86)\SYMPHONY`.

Test the executable by opening Windows Explorer and double-click on `symphony.exe` in the `bin` directory in the folder in which the distribution was unpacked. This should open a Window in which the interactive solver will appear. Type `help` or `?` to see a list of available commands or see Chapter 3 for instructions on using the interactive solver.

To test the callable library, the distribution includes sample codes along with associated project files in the `examples` directory.

milp.c: This sample code is an implementation of a basic MILP solver using SYMPHONY's C callable functions with user defined input (see Section 3.3). To build and test the code, open the SYMPHONY MSVC++ solution file and build the project `milp`, as described below in the next section.

2.2 Building From Source

SYMPHONY uses the COIN-OR build system and the GNU autotools to automate the build process. The build process should therefore be similar in all Unix-like environments. It is even possible to use this system to build SYMPHONY in Windows using the Microsoft Visual C++ compiler if you have MSys2 (<https://msys2.github.io/>), CYGWIN (<http://cygwin.org/>), or the Windows Subsystem for Linux (<https://docs.microsoft.com/en-us/windows/wsl/install-win10>) installed. The instructions below will lead you through the steps required to compile SYMPHONY as a generic MILP solver. This process will create (1) a generic callable library that allows SYMPHONY to be called from a C or C++ code and (2) an executable that can be used as a stand-alone application to solve MILPs written in either MPS or GMPL file format. SYMPHONY can be further customized by implementing one of more than 50 callback functions that change SYMPHONY's default behavior. For information on customizing SYMPHONY using callbacks, a quick start guide is provided below. More detailed information is provided in Chapter 5.

2.2.1 External Dependencies

2.2.1.1 The LP Engine

SYMPHONY requires the use of a third-party callable library to solve the LP relaxations once they are formulated. The LP solver is called through Open Solver Interface, also available from COIN-OR. The list of solvers with OSI interfaces currently numbers eight and includes both commercial and open source alternatives. By default, SYMPHONY now uses the COIN LP solver (Clp). However, if another LP solver is desired, this is possible by changing the configuration settings.

2.2.1.2 GMPL Parser

If you would like to be able to parse GMPL models (GMPL is a subset of AMPL), you will need to install GLPK (<http://www.gnu.org/software/glpk/>). To do so automatically, run the `get.Glpk` script in the `ThirdParty/Glpk` directory. After that, run `configure` with the additional argument `--with-gmpl` and Glpk will be built and linked automatically, enabling the ability to read GMPL files.

2.2.1.3 COIN Libraries

SYMPHONY uses other COIN libraries for certain functionality, such as parsing MPS files, solving linear programs, generating valid inequalities, and for general matrix utilities. The libraries required for this functionality are now included with the distribution.

2.2.1.4 GNU Libraries

If the `readline` and `history` libraries are available, the interactive solver will have command history and command completion available. This functionality is only available in Unix-like environments by configuring with the `--enable-gnu-packages` option (see below).

2.2.2 Building in Unix-like environments

2.2.2.1 Building the Standard Configuration

Building on Linux. Most Linux distributions come with all the required tools installed, but to be sure, check that you have

- `pkg-config`, `wget`, `gcc`, `g++`, `make`, `git`, `tar`
- BLAS and LAPACK
- Optional additions
 - Metis
 - Suitesparse
 - libamd
 - Cholmod
 - Nauty

To obtain the source code, the first step is to get the installer that will then fetch the source for SYMPHONY and all its dependencies. *You do not need to clone SYMPHONY first, just do the following!* Open a terminal and execute

```
git clone https://www.github.com/coin-or/coinbrew
```

Next, to check out source code for and build all the necessary projects (including dependencies), use the `coinbrew` script.

```
cd coinbrew
chmod u+x coinbrew
./coinbrew
```

(Note: The `chmod` command is only needed if the execute permission is not automatically set by git on cloning). Once you run the script, you will be prompted interactively to select a project to fetch and build. The rest should happen automatically. Alternatively, the following command-line incantation will execute the procedure non-interactively.

```
./coinbrew fetch build SYMPHONY@VER
```

The `coinbrew` script invokes a separate `configure` script that is auto-generated using the GNU autotools. This script takes a wide range of additional options and these options can be specified on the command-line when running `coinbrew`. For example, to build with debugging symbols, do

```
./coinbrew fetch build SYMPHONY@VER --enable-debug
```

To get help with additional options available in running the script, do

```
./coinbrew --help
```

The above procedures will build all required dependencies and SYMPHONY itself. Afterwards, the binaries will be installed in the directory `dist/bin` and the libraries in the directory `dist/lib`. If you wish to install in a different directory, such as `/usr/local`, then run the command

```
./coinbrew fetch build SYMPHONY@VER --prefix=/path/to/install/dir
```

After installation, you will also need to add `/path/to/install/dir/bin` to your `PATH` variable in your `.bashrc` and also add `/path/to/install/dir/lib` to your `LD_LIBRARY_PATH` if you want to link to COIN libraries.

Building on Windows. By far, the easiest way to build on Windows is with the GNU autotools and the GCC compilers. The first step is to install either

- Msys2 (<https://msys2.github.io/>)

- Windows Subsystem for Linux (<https://docs.microsoft.com/en-us/windows/wsl/install-win10>)

If you don't want to fool around with WSL (which is a great option if you already know your way around Unix), it is recommended to use MSys2, since it provides a minimal toolset that is easy to install. To get MSys2, download the installer from <https://msys2.github.io/>.

Following the above steps, you should have the `bash` command (with MSys2, be sure to run `msys2.shell.bat` or manually add `msys64\usr\bin`, `msys64\mingw32\bin`, and `msys64\mingw64\bin` to your Windows path).

Once you have bash installed and in your `PATH`, open a Windows terminal and type

```
bash
pacman -S make wget tar patch dos2unix diffutils git svn pkg-config zip unzip
```

To obtain the source code, the first step is to get the installer that will then fetch the source for SYMPHONY and all its dependencies. *You do not need to clone SYMPHONY first, just do the following!* Open a terminal and execute

```
git clone https://www.github.com/coin-or/coinbrew
```

Next, to check out source code for and build all the necessary projects (including dependencies), execute the script in the `coinbrew` subdirectory. To execute the script, do

```
cd coinbrew
chmod u+x coinbrew
./coinbrew
```

(Note: The `chmod` command is only needed if the execute permission is not automatically set by git on cloning). Once you run the script, you will be prompted interactively to select a project to fetch and build. the rest should happen automatically. Alternatively, the following command-line incantation will execute the procedure non-interactively.

```
./coinbrew fetch build SYMPHONY@VER
```

The `coinbrew` script invokes a separate `configure` script that is auto-generated using the GNU autotools. This script takes a wide range of additional options and these options can be specified on the command-line when running `coinbrew`. For example, to build with debugging symbols, do

```
./coinbrew fetch build SYMPHONY@VER --enable-debug
```

To get help with additional options available in running the script, do

```
./coinbrew --help
```

To use the resulting binaries and/or libraries, you will need to add the full path of the directory `build\bin` to your Windows executable search `PATH`, or, alternatively, copy the contents of the build directory to `C:\Program Files (x86)\SYMPHONY` and add the directory `C:\Program Files (x86)\SYMPHONY\bin` to your Windows executable search `PATH`. You may also consider adding `C:\Program Files (x86)\SYMPHONY\lib` to the `LIB` path and `C:\Program Files (x86)\SYMPHONY\include` to the `INCLUDE` path.

It is possible to use almost the exact same commands to build with the Visual Studio compilers. Before doing any of the above commands in the Windows terminal, first run the `vcvarsall.bat` script for your version of Visual Studio. Then follow all the steps above, but add `--enable-msvc` to the `coinbrew` command.

```
./coinbrew fetch build SYMPHONY@VER --enable-msvc
```

Building on OS X. OS X is a Unix-based OS and ships with many of the basic components needed to build COIN-OR, but it's missing some things. For examples, the latest versions of OS X come with the `clang` compiler but no Fortran compiler. You may also be missing the `wget` utility and `subversion` and `git` clients (needed for obtaining source code). The easiest way to get these missing utilities is to install Homebrew (see <http://brew.sh>). After installation, open a terminal and do

```
brew install gcc wget svn git
```

To obtain the source code, the first step is to get the installer that will then fetch the source for SYMPHONY and all its dependencies. *You do not need to clone SYMPHONY first, just do the following!* Open a terminal and execute

```
git clone https://www.github.com/coin-or/coinbrew
```

Next, to check out source code for and build all the necessary projects (including dependencies), execute the script in the `coinbrew` subdirectory. To execute the script, do

```
cd COIN-OR-OptimizationSuite
chmod u+x coi.install.sh
./coin.install.sh
```

(Note: The `chmod` command is only needed if the execute permission is not automatically set by git on cloning). Once you run the script, you will be prompted interactively to select a project to fetch and build. the rest should happen automatically. Alternatively, the following command-line incantation will execute the procedure non-interactively.

```
./coinbrew fetch build SYMPHONY@VER
```

With this setup, `clang` will be used for compiling C++ by default. If you want to use the `gcc` compiler provided by Homebrew, then replace the `build` command above with

```
./coinbrew fetch build SYMPHONY@VER CC=gcc-* CXX=g++-*
```

where `*` is the version number of `gcc`. The install script invokes a separate `configure` script that is auto-generated using the GNU autotools. This script takes a wide range of additional options and these options can be specified on the command-line when running `coinbrew`. For example, to build with debugging symbols, do

```
./coinbrew fetch build SYMPHONY@VER --enable-debug
```

To get help with additional options available in running the script, do

```
./coinbrew --help
```

If you wish to install in a different directory, such as `/usr/local`, then run the command

```
./coinbrew fetch build SYMPHONY@VER --prefix=/path/to/install/dir
```

After installation, you will also need to add `/path/to/install/dir/bin` to your `PATH` variable in your `.bashrc` and also add `/path/to/install/dir/lib` to your `DYLD_LIBRARY_PATH` if you want to link to COIN libraries.

2.2.2.2 Advanced Configuration Options

The `coin.install.sh` script supports many advanced configuration options. See Figure 2.1 for a list of some of the options the user may want to set.

In order to enable or disable an option, add the option as an argument to the `coin.install.sh` script. For instance, running

```
./coinbrew fetch build SYMPHONY@VER --enable-debug
```

will compile the source files with the debugging flag.

It is possible to use compilers other than the default (which is `g++`). For example, to perform an automated build of SYMPHONY using the MSVC++ compiler `cl` with GNU autotools in the CYGWIN environment configure with

```
./coinbrew fetch build SYMPHONY@VER --enable-msvc
```

To test out the optimizer manually, a sample MPS file called `sample.mps` and a sample GMPL/AMPL file called `sample.mod` together with its data file `sample.dat` are included with the distribution. You can use either the command-line or the interactive optimizer. To solve the sample MPS model, type

```
bin/symphony -F SYMPHONY/Datasets/sample.mps
```

To solve the GMPL model, use the `-F` switch to specify the file name and the `-D` for the data file name if the input is in GMPL/AMPL format, i.e., type

```
bin/symphony -F SYMPHONY/Datasets/sample.mod -D SYMPHONY/Datasets/sample.dat
```

For more MPS data files for further testing, see the MIPLIB library in the `Data/` subdirectory. To run the interactive optimizer, execute SYMPHONY without any command-line arguments, i.e., type

```
bin/symphony
```

and then type `help` or `?` to see a list of available commands. After the SYMPHONY library and the executable are compiled and tested, you can type

```
make clean
```

if you want to save disk space. That's it! Now you are ready to use SYMPHONY callable library or solve generic MILP problems through the executable.

2.2.2.3 Building for parallel architectures

Shared-memory architectures. To compile a shared memory version of SYMPHONY, simply use an OpenMP compliant compiler. Version 5.6.18 builds with OpenMP support enabled automatically when the compiler provides it. This should be the case for all recent versions of gcc and recent Microsoft compilers. Clang on OS X does not support OpenMP, though it is possible to build a version that does yourself. You can manually enable or disable OpenMP support with configure options by doing, e.g.,

```
./coinbrew fetch build SYMPHONY@VER --enable-openmp
```

In the Microsoft Visual Studio, enable OpenMP support in the properties for the SYMPHONY projects (it doesn't need to be enabled for the dependent projects). In a future version, this will also be made the default.

After configuring, follow the earlier instructions for building and testing. To invoke SYMPHONY from the command-line with multiple threads, specify the number of *additional* worker threads with the `-p` option, i.e., invoking SYMPHONY with

```
bin/symphony -p 2 -F SYMPHONY/Datasets/sample.mps
```

will utilize two worker threads to process subproblems in parallel. When more than 1 thread is used, the first thread is the “master” thread and only prints out periodic status messages, as well as doing some bookkeeping work. Therefore, it is recommended to always use at least one worker thread (`-p 1`). Starting in version 5.6.0, the number of threads used is automatically set equal to the number of available cores if no option is specified. For a fully sequential run with just one thread, invoke SYMPHONY with

```
bin/symphony -p 2 -F SYMPHONY/Datasets/sample.mps
```

Distributed-memory architectures To compile a distributed application, it is necessary that PVM be installed either in the system path or in the directory pointed to by the environment variable `PVM_ROOT` (this can be your home directory if PVM is not already installed on your network). To install PVM yourself, the current version can be obtained at <http://www.csm.ornl.gov/pvm/>. It should compile and install without problems on most architectures. You will have to make a few modifications to your `.cshrc` file, such as defining the `PVM_ROOT` environment variable, but this is all explained clearly in the PVM documentation. Note that all executables (or at least a link to them) must reside in the `$PVM_ROOT/bin/$PVM_ARCH` directory in order for parallel processes to be spawned correctly. The environment variable `PVM_ARCH` is set in your `.cshrc` file and should contain a string representing the current architecture type. To run a parallel application, you must first start up the daemon on each of the machines you plan to use in the computation. How to do this is also explained in the PVM documentation.

To configure for a parallel build with the default parallel configuration, invoke the configuration script as follows:

```
./coinbrew fetch build SYMPHONY@VER --with-pvm
```

Note that there are a number of different parallel configurations (see Chapter 4.3.3 for an overview of SYMPHONY’s parallel modules). The default configuration is to build two parallel modules, the first consisting of the master, tree management, and cut management modules, while the second consisting of the node processing, and cut generation modules. To build in another configuration, there are four configure flags that control which modules run as separate executables and which are called directly in serial fashion. The variables are as follows:

- `--with-cg`: If set, then the cut generator function will be called directly from the LP in serial fashion, instead of running as a separate executable. This is desirable if cut generation is quick and running it in parallel is not worth the price of the communication overhead.
- `--with-cp`: If set, then the cut pool(s) will be maintained as a data structure auxiliary to the tree manager.
- `--with-lp`: If set, then the LP functions will be called directly from the tree manager. When running the distributed version, this necessarily implies that there will only be one active

subproblem at a time, and hence the code will essentially be running serially. In the shared-memory version, however, the tree manager will be threaded in order to execute subproblems in parallel.

`--with-tm`: If set, then the tree will be managed directly from the master process. This is only recommended if a single executable is desired (i.e. the three other variables are also set to true). A single executable is extremely useful for debugging purposes.

These variables can be set in virtually any combination, though some don't really make much sense. Note that in a few user functions that involve process communication, there will be different versions for serial and parallel computation. This is accomplished through the use of `#ifdef` statements in the source code. This is well documented in the function descriptions and the in the source files containing the function stubs.

Once configured, follow the build instructions in Section 2.1.1 to build the code. Note that this will also compile the sequential version. Make sure there are links from your `$PVM_ROOT/bin/$PVM_ARCH` subdirectory to each of the executables in your `bin/` directory, as required by PVM. In order to keep track of the various possible configurations, executable and their corresponding libraries are named as follows. The name of each executable is `symphony`, along with a combination of the (abbreviated) names of the modules that were combined to produce that executable joined by underscores: `m` for the master module, `tm` for the tree management module, `lp` for the node processing module, `cg` for the cut generation module, and `cp` for the cut management module. For instance, in the default distributed configuration, the executables produced are `symphony_m_tm_cp` and `symphony_lp_cg`.

To test the parallel version, first start the PVM daemon by typing `pvm` on the command line and then typing `quit`. As above, invoke SYMPHONY using the sample MPS file called `sample.mps` included with the distribution. To specify the file name, use the `-F` command-line option, i.e., in the root directory, type

```
bin/symphony_m_EXT -F SYMPHONY/Datasets/sample.mps
```

where `EXT` is the extension to be added according to the chosen configuration of the modules.

2.2.2.4 Building SYMPHONY Applications

There are a number of sample applications available as examples of how to do customized development with SYMPHONY. These include customized solvers for the matching problem, the set partitioning problem (simple and advanced versions), the vehicle routing and traveling salesman problems, the mixed postman problem, the multi-criteria knapsack problem, and the capacitated network routing problem. These applications are contained in the `SYMPHONY/Applications/` subdirectory in the distribution. There is also a white paper that guides the user through the development of the MATCH solver in the `SYMPHONY/Doc/` directory. For detailed instructions on developing your own application with SYMPHONY, see Chapter 5.

In order to compile SYMPHONY's applications in Unix-like environments, you must first compile a version of the callable library with hooks for the callback functions.

```
./coinbrew fetch build SYMPHONY@\VER --with-application  
make  
make install
```

This will create the application library called `libSymAppl` to be used while building custom applications. Note that the generic sequential library and executable will also be made and installed.

After building the library, go to one of the application subdirectories in the `SYMPHONY/Applications/` directory and type `make` there to build the corresponding application. For more information, including the parallel configuration instructions, see the `INSTALL` file of the corresponding application.

2.2.3 Building Using Microsoft Visual C++

Here is a sketch outline of how to compile SYMPHONY in MS Windows with the MSVC++ compiler. These instructions will lead you through the steps required to compile SYMPHONY as a generic MILP solver. Note that the Windows version has some limitations. Detailed timing information is not currently provided. Support is only provided for running in sequential mode at this time.

First, obtain the source code by downloading from <https://www.coin-or.org/download/source/SYMPHONY/>. Unpack the archive to create the directory `SYMPHONY-5.6.18`. You now have three options. You can either build using the MSVC++ IDE, build on the command-line with MSVC++ compiler, or use the NMAKE utility.

2.2.3.1 Building with the MSVC++ Graphical Interface

These instructions are for MSVC++ Version 8. Instructions for other versions should be similar.

- Go to `MSVisualStudio\v10` directory and open the solution file `symphony.sln`.
- Note that there are a number of additional preprocessor definitions that control the functionality of SYMPHONY. These definitions are described in `sym.mak`, a Unix-style makefile included in the distribution. To enable the functionality associated with a particular definition, simply add it to the list of definitions of `libSymphony` project together with the required libraries and paths. For instance, if you want to enable GMPL reader option, you need to
 - add the directory of the header files of GLPK to the include files path
 - add `USE_GLMPL` to the defines
 - add the GLPK library to the solution
- Make sure that the project `symphony` is set as the startup project by choosing `Set as Startup Project` from the `Project` menu after selecting the `symphony` project in the Solution Explorer. Choose `Build Solution` from the Build menu. This should successfully build the SYMPHONY library and the corresponding executable.

- To test the executable, go to the **Debug** tab and choose **Start Without Debugging**. This should open a Window in which the interactive solver will appear. Type **help** or **?** to see a list of available commands or see Chapter 3 for instructions on using the interactive solver.

Note that there is some functionality missing from the Windows version. Most prominently, the timing functions do not work. In addition, the Windows version will only run in sequential mode for a variety of reasons. However, it should be relatively easy to get it running in parallel if you can get PVM working under Windows.

2.2.3.2 Building in a Windows Terminal

These instructions are for MSVC++ Version 10. Instructions for other versions should be similar, though project files for earlier versions are not well maintained.

- Open a command terminal (choose **Run** on the Start menu and type **cmd** in the dialogue box). Go to the **MSVisualStudio\v10** directory and type

```
devenv symphony.sln /Build "Win32|Release"
```

This will create both the 32-bit release version of SYMPHONY, including the library **libSymphony.lib** and the executable **symphony**. Of course, other configurations are supported as well and binaries will be created side-by-side in the appropriate directories according to platform and configuration selected. The library, together with the header files in **SYMPHONY\include**, can then be used to call SYMPHONY from any C/C++ code. The API for calling SYMPHONY is described in Section 3.3.

- Test the executable by opening Windows Explorer and double-clicking on **symphony.exe** in the **Debug** subdirectory. This should open a Window in which the interactive solver will appear. Type **help** or **?** to see a list of available commands or see Chapter 3 for instructions on using the interactive solver.
- If you modify the source code of SYMPHONY, type

```
devenv symphony.sln /Rebuild "Win32|Release"
```

in order to clean and rebuild everything.

2.2.3.3 Building with the MSVC++ compiler in MSys2

It is possible to perform an automated build of SYMPHONY using the MSVC++ compiler **cl** with GNU autotools in the MSys2 environment. To do so, follow the instructions for building in Unix-like environments (see Section 2.2.2), except when configuring, use the command

```
./coinbrew fetch build SYMPHONY@VER --enable-msvc
```

2.2.3.4 Building SYMPHONY Applications

As mentioned above, there are a number of sample applications available as examples of how to do development with SYMPHONY. These include solvers for the matching problem, the set partitioning problem (simple and advanced versions), the vehicle routing and traveling salesman problems, the mixed postman problem, multi-criteria knapsack problem and, capacitated network routing problem. These applications are contained in the `SYMPHONY/Applications/` subdirectory in the distribution. There is also a white paper that guides the user through the development of the MATCH solver in the `SYMPHONY/Doc/` directory. For instructions on developing your own application with SYMPHONY, see Chapter 5.

In order to compile SYMPHONY's applications in the Microsoft Visual C++ environment, obtain the source code as described earlier. As before, you then have three options. You can either build using the MSVC++ IDE, build on the command-line with MSVC++ executable, or use the NMAKE utility. The below instructions are for MSVC++ Version 6, but building in other versions should be similar. All of the following commands should be executed in the `SYMPHONY-5.6.18\Applications\XXX\MSVisualStudio\v10` directory, where `XXX` is the name of the application to be built.

Building With the MSVC++ Graphical Interface

- Open the solution file `xxx.sln`.
- The configuration steps are exactly the same as that described in Section 2.2.3.1. The only difference is that you are using the `xxx` project instead of the `symphony` project. Go through the same steps to configure.
- Once you have the proper settings, choose **Build xxx.exe** from the **Build** menu. This should successfully build the executable.
- To test the executable, right click on the `xxx` project, go to the **Debug** tab and set the program arguments to `-F ..\..\sample.xxx`. Note that command-line switches are Unix-style.
- Now choose **Execute** from the build menu and you have a working branch and bound solver.

Building in a Windows Terminal

- Open a command terminal (choose **Run** on the Start menu and type `cmd` in the dialogue box) and type

```
devenv xxx.sln /Build "Win32|Release"
```

where `xxx` is the name of the application. This will create the release version of the application executable, as above.

- To test the executable, type

```
Debug\xxx.exe -F ..\..\sample.xxx
```

- If the source files for the application are modified, type

```
devenv user.sln /Rebuild "Win32|Release"
```

in order to clean and rebuild everything.

Building with the MSVC++ compiler in MSys2 It is possible to build applications using an automated build of SYMPHONY with the MSVC++ compiler `cl` using the GNU autotools in the MSys2 environment. To do so, follow the instructions for building in Unix-like environments (see Section 2.2.2), except when configuring, use the command

```
./coinbrew fetch build SYMPHONY@VER --enable-msvc
```

Afterwards, you can build the individual applications using the “make” command, as usual in a Unix-like environment except that the compiler used will be `cl`.

<code>--enable-debug</code>	compile all projects with debug options set
<code>--enable-debug-symphony</code>	compile only SYMPHONY project with debug options
<code>--enable-msvc</code>	Under MSys2, compile so that executables do not depend on the CYGWIN DLLs
<code>--enable-static</code>	build static libraries
<code>--enable-static-executable</code>	create a complete static executable
<code>--enable-gnu-packages</code>	compile with GNU packages
	compile interactive optimizer with readline library
<code>--disable-cgl-cuts</code>	disable generic cut generation
<code>--enable-sensitivity-analysis</code>	compile in the sensitivity analysis features
<code>--enable-root-only</code>	process only the root node
<code>--enable-frac-branching</code>	compile in the fractional branching option
<code>--enable-tests</code>	perform additional sanity checks (for debugging purposes)
<code>--enable-tm-tests</code>	perform more tests
<code>--enable-trace-path</code>	additional debugging options
<code>--enable-cut-check</code>	additional debugging options
<code>--enable-statistics</code>	additional statistics
<code>--enable-pseudo-costs</code>	enable some experimental pseudo-cost branching tools
<code>--enable-draw-graph</code>	enable IGD graph drawing application
<code>--with-XXX-incdir</code>	specify the directory with the header files for the XXX package where XXX is one of LP solver packages: cplex, glpk, osl, soplex, xpress
<code>--with-XXX-lib</code>	specify the flags to link with the library XXX package where XXX is one of LP solver packages: cplex, glpk, osl, soplex, xpress
<code>--with-lp-solver[=lpsolver]</code>	specify the LP solver in small letters (default lpsolver=clp)
<code>--with-application</code>	compile the application library
<code>--enable-openmp</code>	compile in OpenMP features
<code>--with-pvm</code>	compile in parallel architecture (assuming that pvm is installed and the variable PVM_ROOT is defined.)
<code>--without-cg</code>	compile without cut generator module
<code>--without-cp</code>	compile without cut pool module
<code>--without-lp</code>	compile without LP solver module
<code>--without-tm</code>	compile without tree manager module

Figure 2.1: A list of useful configuration options

Chapter 3

Using SYMPHONY

3.1 Using SYMPHONY Interactively

3.1.1 Unix-like Environments

If you are planning to use the interactive optimizer in a Unix-like environment and you are building SYMPHONY from source, it is recommended that you run the configuration script (see Section 2.2) with the command-line argument that enables GNU packages, i.e.,

```
./coinbrew fetch build SYMPHONY@\VER --enable-gnu-packages
```

This will allow the interactive shell to behave exactly like a Linux terminal command line, i.e., it will keep the history of the used commands, will do command completion, etc. Note that you must have the required packages (`readline` and `history`) installed.

To use SYMPHONY's interactive shell, run the executable without any command line arguments, i.e., type

```
bin/symphony
```

You will enter a command shell environment where you will be prompted for inputs. The user interface consists of a *main menu*, where an instance is read in and solved, a *set menu*, where parameters are set, and a *display menu*, where results, statistics and parameter values are displayed.

3.1.2 Microsoft Windows

To invoke SYMPHONY's interactive solver in an Microsoft Windows environment, simply double-click on the `symphony.exe` file in Windows Explorer. This should open a terminal window in which the solver will run. Note that if you built SYMPHONY in CYGWIN without the option `--enable-dos-compile`, then you will have to have the CYGWIN DLL in your path in order for the executable to run.

3.1.3 Main Menu

Below is the main menu displayed at the beginning of a new session:

```
*****
*   This is SYMPHONY Version 5.6.18                               *
*   Copyright 2000-2022 Ted Ralphs                               *
*   All Rights Reserved.                                         *
*   Distributed under the Eclipse Public License 1.0             *
*****

***** WELCOME TO SYMPHONY INTERACTIVE MIP SOLVER *****

Please type 'help'/'?' to see the main commands!

SYMPHONY:
```

When you type `help` or `?`, a list of main commands is displayed:

```
SYMPHONY: help

List of main commands:

load      : read a problem in mps or ampl format
solve     : solve the problem
lpsolve   : solve the lp relaxation of the problem
set       : set a parameter
display   : display optimization results and stats
reset     : restart the optimizer
help      : show the available commands/params/options

quit/exit : leave the optimizer

SYMPHONY:
```

Following is an illustration of a session to read in a sample instance:

```
SYMPHONY: load
Name of the file: sample.mps
Coin0001I At line 1 NAME SAMPLE
Coin0001I At line 2 ROWS
Coin0001I At line 6 COLUMNS
Coin0001I At line 25 RHS
Coin0001I At line 28 BOUNDS
Coin0001I At line 34 ENDATA
```



```
Coin0002I Problem SAMPLE has 2 rows, 6 columns and 10 elements
SYMPHONY:
```

The format of the input file is recognized from the file extension. If there is none, you will be prompted to define the input format:

```
SYMPHONY: load
Name of the file: sample
Type of the file ('mps'/'ampl'/'gmpl'): mps
Coin0001I At line 1 NAME SAMPLE
Coin0001I At line 2 ROWS
Coin0001I At line 6 COLUMNS
Coin0001I At line 25 RHS
Coin0001I At line 28 BOUNDS
Coin0001I At line 34 ENDATA
Coin0002I Problem SAMPLE has 2 rows, 6 columns and 10 elements
SYMPHONY:
```

If the input is in AMPL/GMPL format, you will also be prompted to read in a data file (note again that in order to enable GMPL/AMPL reader, you have to install GLPK—see Section 2.2)):

```
SYMPHONY: load
Name of the file: sample.mod
Name of the data file: sample.dat
Reading model section from sample.mod...
32 lines were read
Reading data section from sample.dat...
68 lines were read
Generating nb...
Generating cost...
Model has been successfully generated
SYMPHONY:
```

After loading the instance, type **solve** to solve the corresponding integer program or **lpsolve** to solve its linear relaxation:

```
SYMPHONY: solve

***** Found Better Feasible Solution !
***** Cost: -40.000000

*****
* Optimal Solution Found                                     *
*****
```

```
SYMPHONY: lpsolve
```

```
***** Found Better Feasible Solution !
***** Cost: -43.000000
```

```
*****
* Optimal Solution Found                      *
*****
```

```
SYMPHONY:
```

As above, only the objective values of the feasible solutions found so far and the termination code of the solution process will be displayed (see Section 3.1.4 for displaying more output).

3.1.4 Set Menu

The **Set** submenu is used to set SYMPHONY's run-time parameters. To enter this submenu, type **set**:

```
SYMPHONY: set
Please type 'help'/'?' to see the list of parameters!
SYMPHONY\Set:
```

You can override the default value of a parameter by typing the name of the parameter. You will then be prompted to enter the new value of that parameter. For instance, in order to display more outputs during the solution process, you need to set the **verbosity** parameter (set to -1 by default for the interactive shell routines) to a nonnegative integer:

```
SYMPHONY\Set: verbosity
Value of the parameter: 3
Setting verbosity to: 3
SYMPHONY\Set:
```

A confirmation message will also be displayed. Note that typing **help** or **?** displays only a subset of the most commonly used run-time parameters. However, you are allowed to set any of the parameters given in Section 6.4. Additionally, you can set the values of parameters using a parameter file as an input. In such a file, the new value of each parameter must follow the name of that parameter. For instance, suppose that the **my_param** file consists of the following lines:

```
verbosity 3
node_selection_rule 3
time_limit 100
```

Then, type `param_file` to be prompted to read in the parameter file:

```
SYMPHONY\Set: param_file
Name of the parameter file: my_param
Setting verbosity to: 3
Setting node_selection_rule to: 3
Setting time_limit to: 100
SYMPHONY\Set:
```

At this point, you can return to the main menu by typing `back`, load an instance and solve it with updated run-time parameters.

3.1.5 Display Menu

The `Display` submenu is used to print out results and statistics of the solution process after a `solve` call. To enter this submenu and see available options, type `display` and then `help` or `?`:

```
SYMPHONY: display
Please type 'help'/'?' to see the display options!
SYMPHONY\Display: help
```

List of display options:

```
solution      : display the column values
obj           : display the objective value
stats         : display the statistics
parameter     : display the value of a parameter

back          : leave this menu
quit/exit     : leave the optimizer
```

```
SYMPHONY\Display:
```

Clearly, in order to display column solutions and the optimal solution value, you need to type `solution` and then `obj`:

```
SYMPHONY\Display: solution
Optimal Solution found!
+++++
Nonzero column names and values in the solution
+++++
COL00002      3.000
COL00006      1.000
```

```
SYMPHONY\Display: obj
Objective Value: -40.000000
SYMPHONY\Display:
```

You can also display the values of SYMPHONY's run-time parameters (see Section 6.4) by moving into **parameters** submenu:

```
SYMPHONY\Display: parameter
Please type 'help'/'?' to see the list of available parameters!
SYMPHONY\Display\Parameter:
```

For instance, in order to display the verbosity level, type **verbosity**:

```
SYMPHONY\Display\Parameter: verbosity
The value of verbosity: 3
SYMPHONY\Display\Parameter:
```

As in Set submenu, typing **help** or **?** will display only a subset of available run-time parameters. However, you are allowed to display the value of any of the parameters given in Section 6.4.

3.1.6 Sub Menu Browsing

SYMPHONY's interactive optimizer also allows the user to reach the lower level menu commands from the higher level menus. In other words, the user has the flexibility to use submenu commands without entering the corresponding submenu. As an instance, all three of the following sessions have the same result:

- SYMPHONY: display parameter verbosity
- SYMPHONY: display
Please type 'help'/'?' to see the display options!
SYMPHONY\Display: parameter verbosity
- SYMPHONY: display
Please type 'help'/'?' to see the display options!
SYMPHONY\Display: parameter
Please type 'help'/'?' to see the list of available parameters!
SYMPHONY\Display\Parameter: verbosity

This flexibility is also enabled for the **load** command and the Set submenu. The followings are all valid commands:

```
SYMPHONY: load sample.mps
```

```
SYMPHONY: load sample.mod sample.dat
```

```
SYMPHONY: set  
SYMPHONY\Set: verbosity 3
```

```
SYMPHONY: set verbosity 3  
SYMPHONY: set param_file my_param
```

3.2 Using SYMPHONY from the Command Line

For batch processing and scripting, SYMPHONY can also be called from the command line from a terminal in any operating system (note that in the Windows terminal, the path separator is `\` rather than `/`). When called from the command line, a number of command-line switches can be invoked to specify the file to be read and solved, as well as set parameters. Note that the switches are Unix-style, even in Windows). At a minimum, one must specify the name of the file to be read and solved. The following is the calling sequence to load in an instance file in MPS format and solve it.

```
./symphony -F sample.mps
```

To read and solve a model in LP format, the command would be

```
./symphony -L sample.lp
```

To read and solve a GMPL model and associated data file, the command would be

```
./symphony -F sample.mod -D sample.dat
```

In addition to specifying the name of the instance file, most of the common parameters can also be set on the command line by adding various switches. Calling SYMPHONY with just the argument `-h` will list all the options. To set parameters that cannot be set on the command line or to save parameter setting, it is possible to use a parameter file in which a group of parameters can be set. To invoke SYMPHONY with a parameter file, type `./symphony -f filename`, where `filename` is the name of the parameter file. The format of the file and a list of all parameters is given in Section 6.4.

The output level can be controlled through the use of the verbosity parameter, which can be invoked. Setting this parameter at different levels will cause different progress messages to be printed out. Level 0 only prints out the introductory and solution summary messages, along with status messages every 10 minutes. Level 1 prints out a message every time a new node is created. Level 3 prints out messages describing each iteration of the solution process. Levels beyond 3 print out even more detailed information. To get no output at all, the verbosity level must be set to -2.

3.3 Using the Callable Library

SYMPHONY's callable library consists of a complete set of subroutines for loading and modifying problem data, setting parameters, and invoking solution algorithms. The user invokes these subroutines through the API specified in the header file `symphony_api.h`. Some of the basic commands are described below. For the sake of brevity, the arguments have been left out.

3.3.1 The C API

`sym_open_environment()` Opens a new environment, and returns a pointer to it. This pointer then has to be passed as an argument to all other API subroutines (in the C++ interface, this pointer is maintained for the user).

`sym_parse_command_line()` Invokes the built-in parser for setting commonly used parameters, such as the file name which to read the problem data, via command-line switches. A call to this subroutine instructs SYMPHONY to parse the command line and set the appropriate parameters. This subroutine also sets all other parameter values to their defaults, so it should only be called when this is desired.

`sym_load_problem()` Reads the problem data and sets up the root subproblem. This includes specifying which cuts and variables are in the *core* (those that are initially present in every subproblem during the search process) and the additional cuts and variables to be initially active in the root subproblem. By default, SYMPHONY reads an MPS or GMPL file specified by the user, but the user can override this default by implementing a user callback that reads the data from a file in a customized format (see Section 3.4).

`sym_find_initial_bounds()` Invokes the user callback to find initial bounds using a custom heuristic.

`sym_solve()` Solves the currently loaded problem from scratch. This method is described in more detail in Section 4.4.1.1.

```

int main(int argc, char **argv)
{
    sym_environment *env = sym_open_environment();
    sym_parse_command_line(env, argc, argv);
    sym_load_problem(env);
    sym_solve(env);
    sym_close_environment(env);
}

```

Figure 3.1: Implementation of a generic MILP solver with the SYMPHONY C callable library.

sym_warm_solve() Solves the currently loaded problem from a warm start. This method is described in more detail in Section 4.4.1.2.

sym_mc_solve() Solves the currently loaded problem as a multicriteria problem. This method is described in more detail in Section 4.4.1.3.

sym_close_environment() Frees all problem data and deletes the environment.

As an example of the use of the library functions, Figure 3.1 shows the code for implementing a generic MILP solver with default parameter settings. To read in an MPS file called **sample.mps** and solve it using this program, the following command would be issued:

```
./symphony -F sample.mps
```

To read and solve a model in LP format, the command would be

```
./symphony -L sample.lp
```

The user does not have to invoke a command to read the input file. During the call to **sym_parse_command_line()**, SYMPHONY determines that the user wants to read in an MPS file. During the subsequent call to **sym_load_problem()**, the file is read and the problem data stored. To read an GMPL file, the user would issue the command

```
./symphony -F sample.mod -D sample.dat
```

Although the same command-line switch is used to specify the model file, the additional presence of the **-D** option indicates to SYMPHONY that the model file is in GMPL format and GLPK's GMPL parser is invoked [25]. Note that the interface and the code of Figure 3.1 is the same for both sequential and parallel computations. The choice between sequential and parallel execution modes

```

int main(int argc, char **argv)
{
    sym_environment *env = sym_open_environment();
    sym_parse_command_line(env, argc, argv);
    sym_load_problem(env);
    sym_set_int_param(env, "find_first_feasible", TRUE);
    sym_set_int_param(env, "node_selection_strategy", DEPTH_FIRST_SEARCH);
    sym_solve(env);
    sym_set_int_param(env, "find_first_feasible", FALSE);
    sym_set_int_param(env, "node_selection_strategy", BEST_FIRST_SEARCH);
    sym_warm_solve(env);
}

```

Figure 3.2: Implementation of a dynamic MILP solver with SYMPHONY.

is made at compile-time through modification of the makefile or the project settings, depending on the operating system.

To start the solution process from a warm start, the `sym_warm_solve()` command is used. SYMPHONY automatically records the warm start information resulting from the last solve call and restarts from that checkpoint if a call to `sym_warm_solve()` is made. Alternatively, external warm start information can be loaded manually. Figure 3.2 illustrates the use of the re-solve capability by showing the code for implementing a solver that changes from depth first search to best first search after the first feasible solution is found. The user can also modify problem data in between calls to the solver. Code for doing so is shown in Figure 3.3. In this example, the solver is allowed to process 100 nodes and then save the warm start information. Afterward, the original problem is solved to optimality, then is modified and re-solved from the saved checkpoint.

Finally, SYMPHONY now also has a bicriteria solve call. The applications of such a solver are numerous. Besides yielding the ability to closely examine the tradeoffs between competing objectives, the method can be used to perform detailed sensitivity analysis in a manner analogous to that which can be done with simplex based solvers for linear programs. As an example, suppose we would like to know exactly how the optimal objective function value for a given pure integer program depends on the value of a given objective function coefficient. Consider increasing the objective function coefficient of variable i from its current value. Taking the first objective function to be the original one and taking the second objective function to be the i^{th} unit vector, we can derive the desired sensitivity function by using the bicriteria solution algorithm to enumerate all supported solutions and breakpoints. This information can easily be used to obtain the desired function. Figure 3.4 shows the code for performing this analysis on variable 0.

In addition to the parts of the API we have just described, there are a number of standard subroutines for accessing and modifying problem data and parameters. These can be used between calls to the solver to change the behavior of the algorithm or to modify the instance being solved. These modifications are discussed in more detail in Section 4.4.1.2.

```

int main(int argc, char **argv)
{
    warm_start_desc *ws;
    sym_environment *env = sym_open_environment();
    sym_parse_command_line(env, argc, argv);
    sym_load_problem(env);
    sym_set_int_param(env, "node_limit", 100);
    sym_set_int_param(env, "keep_warm_start", TRUE);
    sym_solve(env);
    ws = sym_get_warm_start(env);
    sym_set_int_param(env, "node_limit", -1);
    sym_warm_solve(env);
    sym_set_obj_coeff(env, 0, 100);
    sym_set_obj_coeff(env, 200, 150);
    sym_set_warm_start(ws);
    sym_warm_solve(env);
}

```

Figure 3.3: Use of SYMPHONY’s warm start capability.

3.3.2 The C++ API

The Open Solver Interface (OSI) is a C++ class that provides a standard API for accessing a variety of solvers for mathematical programs. It is provided as part of the COIN-OR repository [24], along with a collection of solver-specific derived classes that translate OSI call into calls to the underlying libraries of the solvers. A code implemented using calls to the methods in the OSI base class can easily be linked with any solver for which there is an OSI interface. This allows development of solver-independent codes and eliminates many portability issues. The current incarnation of OSI supports only solvers for linear and mixed-integer linear programs, although a new version supporting a wider variety of solvers is currently under development.

We have implemented an OSI interface for SYMPHONY 5.6.18 that allows any solver built with

```

int main(int argc, char **argv)
{
    sym_environment *env = sym_open_environment();
    sym_parse_command_line(env, argc, argv);
    sym_load_problem(env);
    sym_set_obj2_coeff(env, 0, 1);
    sym_mc_solve(env);
}

```

Figure 3.4: Performing sensitivity analysis with SYMPHONY’s bicriteria solver.

```

int main(int argc, char **argv)
{
    OsiSymSolverInterface si;
    si.parseCommandLine(argc, argv);
    si.loadProblem();
    si.branchAndBound();
}

```

Figure 3.5: Implementation of a generic MILP solver with the SYMPHONY OSI interface.

SYMPHONY to be accessed through the OSI, including customized solvers and those configured to run on parallel architectures. To ease code maintenance, for each method in the OSI base class, there is a corresponding method in the callable library. The OSI methods are implemented simply as wrapped calls to the SYMPHONY callable library. When an instance of the OSI interface class is constructed, a call is made to `sym_open_environment()` and a pointer to the environment is stored in the class. Most subsequent calls within the class can then be made without any arguments. When the OSI object is destroyed, `sym_close_environment` is called and the environment is destroyed.

To fully support SYMPHONY's capabilities, we have extended the OSI interface to include some methods not in the base class. For example, we added calls equivalent to our `sym_parse_command_line()` and `sym_find_initial_bounds()`. Figure 3.5 shows the program of Figure 3.1 implemented using the OSI interface. Note that the code would be exactly the same for accessing any customized SYMPHONY solver, sequential or parallel.

Although we are using the OSI to access a MILP solver, the current version of the OSI is geared primarily toward support of solvers for linear programming (LP) problems. This is because LP solvers employing some version of the simplex algorithm support much richer functionality and a wider range of interface functions, due to their support of warm starting from previously saved checkpoints. This functionality is difficult to provide for MILP solvers. In SYMPHONY 5.6.18, we have implemented for MILPs some of the same functionality that has long been available for LP solvers. As such, our OSI interface supports warm starting and sensitivity analysis. The implementations of this functionality is straightforward at the moment, but will be improved in future versions.

3.3.3 Linking to the Callable Library

To link your program to the callable library, make sure you have included the header file `symphony.h` in all the source files that call SYMPHONY functions. Also, make sure that your include path contains the directory where all of SYMPHONY's header files are stored. Then simply include the appropriate SYMPHONY library in the set of libraries to be linked and make sure that the path to the library is in the library path. Example makefiles For Unix-like environments are included in the `Examples/` directory.

3.4 Using the Callback Functions

The user's main avenues for customization of SYMPHONY are the tuning of parameters and the implementation of one or more of over 50 user callback functions. The callback functions allow the user to override SYMPHONY's default behavior for many of the functions performed as part of its algorithm. The user has complete control over branching, cutting plane generation, management of the cut pool and the LP relaxation, search and diving strategies, etc. More detailed information about using the callback functions to develop custom applications is provided in Chapter 5.

Chapter 4

Technical Details

4.1 Branch and Bound

Branch and bound is the broad class of algorithms from which branch, cut, and price is descended. A branch and bound algorithm uses a divide and conquer strategy to partition the solution space into *subproblems* and then optimizes individually over each subproblem. For instance, let S be the set of solutions to a given problem, and let $c \in \mathbf{R}^S$ be a vector of costs associated with members of S . Suppose we wish to determine a least cost member of S and we are given $\hat{s} \in S$, a “good” solution determined heuristically. Using branch and bound, we initially examine the entire solution space S . In the *processing* or *bounding* phase, we relax the problem. In so doing, we admit solutions that are not in the feasible set S . Solving this relaxation yields a lower bound on the value of an optimal solution. If the solution to this relaxation is a member of S or has cost equal to \hat{s} , then we are done—either the new solution or \hat{s} , respectively, is optimal. Otherwise, we identify n subsets of S , S_1, \dots, S_n , such that $\cup_{i=1}^n S_i = S$. Each of these subsets is called a *subproblem*; S_1, \dots, S_n are sometimes called the *children* of S . We add the children of S to the list of *candidate subproblems* (those which need processing). This is called *branching*.

To continue the algorithm, we select one of the candidate subproblems and process it. There are four possible results. If we find a feasible solution better than \hat{s} , then we replace \hat{s} with the new solution and continue. We may also find that the subproblem has no solutions, in which case we discard, or *prune* it. Otherwise, we compare the lower bound to our global upper bound. If it is greater than or equal to our current upper bound, then we may again prune the subproblem. Finally, if we cannot prune the subproblem, we are forced to branch and add the children of this subproblem to the list of active candidates. We continue in this way until the list of active subproblems is empty, at which point our current best solution is the optimal one.

4.2 Branch and Cut

In many applications, the bounding operation is accomplished using the tools of linear programming (LP), a technique first described in full generality by Hoffman and Padberg [20]. This general class

Bounding Operation

Input: A subproblem \mathcal{S} , described in terms of a “small” set of inequalities \mathcal{L}' such that $\mathcal{S} = \{x^s : s \in \mathcal{F} \text{ and } ax^s \leq \beta \ \forall (a, \beta) \in \mathcal{L}'\}$ and α , an upper bound on the global optimal value.

Output: Either (1) an optimal solution $s^* \in \mathcal{S}$ to the subproblem, (2) a lower bound on the optimal value of the subproblem, or (3) a message **pruned** indicating that the subproblem should not be considered further.

Step 1. Set $\mathcal{C} \leftarrow \mathcal{L}'$.

Step 2. Solve the LP $\min\{cx : ax \leq \beta \ \forall (a, \beta) \in \mathcal{C}\}$.

Step 3. If the LP has a feasible solution \hat{x} , then go to Step 4. Otherwise, STOP and output **pruned**. This subproblem has no feasible solutions.

Step 4. If $c\hat{x} < \alpha$, then go to Step 5. Otherwise, STOP and output **pruned**. This subproblem cannot produce a solution of value better than α .

Step 5. If \hat{x} is the incidence vector of some $\hat{s} \in \mathcal{S}$, then \hat{s} is the optimal solution to this subproblem. STOP and output \hat{s} as s^* . Otherwise, apply separation algorithms and heuristics to \hat{x} to get a set of violated inequalities \mathcal{C}' . If $\mathcal{C}' = \emptyset$, then $c\hat{x}$ is a lower bound on the value of an optimal element of \mathcal{S} . STOP and return \hat{x} and the lower bound $c\hat{x}$. Otherwise, set $\mathcal{C} \leftarrow \mathcal{C} \cup \mathcal{C}'$ and go to Step 2.

Figure 4.1: Bounding in the branch and cut algorithm

of algorithms is known as *LP-based branch and bound*. Typically, the integrality constraints of an integer programming formulation of the problem are relaxed to obtain a *LP relaxation*, which is then solved to obtain a lower bound for the problem. In [29], Padberg and Rinaldi improved on this basic idea by describing a method of using globally valid inequalities (i.e., inequalities valid for the convex hull of integer solutions) to strengthen the LP relaxation. They called this technique *branch and cut*. Since then, many implementations (including ours) have been fashioned around the framework they described for solving the Traveling Salesman Problem.

As an example, let a combinatorial optimization problem $CP = (E, \mathcal{F})$ with *ground set* E and *feasible set* $\mathcal{F} \subseteq 2^E$ be given along with a cost function $c \in \mathbf{R}^E$. The incidence vectors corresponding to the members of \mathcal{F} are sometimes specified as the set of all incidence vectors obeying a (relatively) small set of inequalities. These inequalities are typically the ones used in the initial LP relaxation. Now let \mathcal{P} be the convex hull of incidence vectors of members of \mathcal{F} . Then we know by Weyl’s Theorem (see [28]) that there exists a finite set \mathcal{L} of inequalities valid for \mathcal{P} such that

$$\mathcal{P} = \{x \in \mathbf{R}^n : ax \leq \beta \ \forall (a, \beta) \in \mathcal{L}\}. \quad (4.1)$$

The inequalities in \mathcal{L} are the potential cutting planes to be added to the relaxation as needed. Unfortunately, it is usually difficult, if not impossible, to enumerate all of inequalities in \mathcal{L} or we could simply solve the problem using linear programming. Instead, they are defined implicitly and we use separation algorithms and heuristics to generate these inequalities when they are violated. In Figure 4.1, we describe more precisely how the bounding operation is carried out in branch and cut.

Once we have failed to either prune the current subproblem or separate the current fractional solution from \mathcal{P} , we are forced to branch. The branching operation is accomplished by specifying a

Branching Operation

Input: A subproblem \mathcal{S} and \hat{x} , the LP solution yielding the lower bound.

Output: S_1, \dots, S_p such that $\mathcal{S} = \cup_{i=1}^p S_i$.

Step 1. Determine sets $\mathcal{L}_1, \dots, \mathcal{L}_p$ of inequalities such that $\mathcal{S} = \cup_{i=1}^p \{x \in \mathcal{S} : ax \leq \beta \mid (a, \beta) \in \mathcal{L}_i\}$ and $\hat{x} \notin \cup_{i=1}^p S_i$.

Step 2. Set $S_i = \{x \in \mathcal{S} : ax \leq \beta \mid (a, \beta) \in \mathcal{L}_i \cup \mathcal{L}'\}$ where \mathcal{L}' is the set of inequalities used to describe \mathcal{S} .

Figure 4.2: Branching in the branch and cut algorithm

Generic Branch and Cut Algorithm

Input: A data array specifying the problem instance.

Output: The global optimal solution s^* to the problem instance.

Step 1. Generate a “good” feasible solution \hat{s} using heuristics. Set $\alpha \leftarrow c(\hat{s})$.

Step 2. Generate the first subproblem \mathcal{S}^I by constructing a small set \mathcal{L}' of inequalities valid for \mathcal{P} . Set $A \leftarrow \{\mathcal{S}^I\}$.

Step 3. If $A = \emptyset$, STOP and output \hat{s} as the global optimum s^* . Otherwise, choose some $\mathcal{S} \in A$. Set $A \leftarrow A \setminus \{\mathcal{S}\}$. Process \mathcal{S} .

Step 4. If the result of Step 3 is a feasible solution \bar{s} , then $c\bar{s} < c\hat{s}$. Set $\hat{s} \leftarrow \bar{s}$ and $\alpha \leftarrow c(\bar{s})$ and go to Step 3. If the subproblem was pruned, go to Step 3. Otherwise, go to Step 5.

Step 5. Perform the branching operation. Add the set of subproblems generated to A and go to Step 3.

Figure 4.3: Description of the generic branch and cut algorithm

set of hyperplanes which divide the current subproblem in such a way that the current solution is not feasible for the LP relaxation of any of the new subproblems. For example, in a combinatorial optimization problem, branching could be accomplished simply by fixing a variable whose current value is fractional to 0 in one branch and 1 in the other. The procedure is described more formally in Figure 4.2. Figure 4.3 gives a high level description of the generic branch and cut algorithm.

In the remainder of the manual, we often use the term *search tree*. This term derives from the common representation of the list of subproblems as the nodes of a graph in which each subproblem is connected only to its parent and its children. Storing the subproblems in such a form is an important aspect of our global data structures. Since the subproblems correspond to the nodes of this graph, they are sometimes referred to as *nodes in the search tree* or simply as *nodes*. The *root node* or *root* of the tree is the node representing the initial subproblem.

4.3 Design of SYMPHONY

SYMPHONY was designed with two major goals in mind—portability and ease of use. With respect to ease of use, we aimed for a “black box” design, whereby the user would not be required to know anything about the implementation of the library, but only about the user interface. With respect to portability, we aimed not only for it to be *possible* to use the framework in a wide

variety of settings and on a wide variety of hardware, but also for it to perform *effectively* in all these settings. Our primary measure of effectiveness was how well the framework would perform in comparison to a problem-specific (or hardware-specific) implementation written “from scratch.”

It is important to point out that achieving such design goals involves a number of very difficult tradeoffs. For instance, ease of use is quite often at odds with efficiency. In several instances, we had to give up some efficiency to make the code easy to work with and to maintain a true black box implementation. Maintaining portability across a wide variety of hardware, both sequential and parallel, also required some difficult choices. For example, solving large-scale problems on sequential platforms requires extremely memory-efficient data structures in order to maintain the very large search trees that can be generated. These storage schemes, however, are highly centralized and do not scale well to large numbers of processors.

4.3.1 An Object-oriented Approach

As we have already alluded to, applying BCP to large-scale problems presents several difficult challenges. First and foremost is designing methods and data structures capable of handling the potentially huge numbers of cuts and variables that need to be accounted for during the solution process. The dynamic nature of the algorithm requires that we must also be able to efficiently move cuts and variables in and out of the *active set* of each search node at any time. A second, closely-related challenge is that of effectively dealing with the very large search trees that can be generated for difficult problem instances. This involves not only the important question of how to store the data, but also how to move it between modules during parallel execution. A final challenge in developing a generic framework, such as SYMPHONY, is to deal with these issues using a problem-independent approach.

Describing a node in the search tree consists of, among other things, specifying which cuts and variables are initially *active* in the subproblem. In fact, the vast majority of the methods in BCP that depend on the model are related to generating, manipulating, and storing the cuts and variables. Hence, SYMPHONY can be considered an object-oriented framework with the central “objects” being the cuts and variables. From the user’s perspective, implementing a BCP algorithm using SYMPHONY consists primarily of specifying various properties of objects, such as how they are generated, how they are represented, and how they should be realized within the context of a particular subproblem.

With this approach, we achieved the “black box” structure by separating these problem-specific functions from the rest of the implementation. The internal library interfaces with the user’s subroutines through a well-defined Application Program Interface (API) (see Section 6.3) and independently performs all the normal functions of BCP—tree management, LP solution, and cut pool management, as well as inter-process communication (when parallelism is employed). Although there are default options for many of the operations, the user can also assert control over the behavior of the algorithm by overriding the default methods or by parameter setting.

Although we have described our approach as being “object-oriented,” we would like to point out that SYMPHONY is implemented in C, not C++. To avoid inefficiencies and enhance the modularity of the code (allowing for easy parallelization), we used a more “function-oriented” approach for the implementation of certain aspects of the framework. For instance, methods used for com-

municating data between modules are not naturally “object-oriented” because the type of data being communicated is usually not known by the message-passing interface. It is also common that efficiency considerations require that a particular method be performed on a whole set of objects at once rather than on just a single object. Simply invoking the same method sequentially on each of the members of the set can be extremely inefficient. In these cases, it is far better to define a method which operates on the whole set at once. In order to overcome these problems, we have also defined a set of *interface functions*, which are associated with the computational modules. These function is described in detail in Section 6.3.

4.3.2 Data Structures and Storage

Both the memory required to store the search tree and the time required to process a node are largely dependent on the number of objects (cuts and variables) that are active in each subproblem. Keeping this active set as small as possible is one of the keys to efficiently implementing BCP. For this reason, we chose data structures that enhance our ability to efficiently move objects in and out of the active set. Allowing sets of cuts and variables to move in and out of the linear programs simultaneously is one of the most significant challenges of BCP. We do this by maintaining an abstract *representation* of each global object that contains information about how to add it to a particular LP relaxation.

In the literature on linear and integer programming, the terms *cut* and *row* are typically used interchangeably. Similarly, *variable* and *column* are often used with similar meanings. In many situations, this is appropriate and does not cause confusion. However, in object-oriented BCP frameworks, such as SYMPHONY or ABACUS [21], a *cut* and a *row* are *fundamentally different objects*. A *cut* (also referred to as a *constraint*) is a user-defined representation of an abstract object which can only be realized as a row in an LP matrix *with respect to a particular set of active variables*. Similarly, a *variable* is a representation which can only be realized as a column of an LP matrix with respect to a *particular set of cuts*. This distinction between the *representation* and the *realization* of objects is a crucial design element and is what allows us to effectively address some of the challenges inherent in BCP. In the remainder of this section, we further discuss this distinction and its implications.

4.3.2.1 Variables

In SYMPHONY, problem variables are *represented* by a unique global index assigned to each variable by the user. This index represents each variable’s position in a “virtual” global list known only to the user. The main requirement of this indexing scheme is that, given an index and a list of active cuts, the user must be able to generate the corresponding column to be added to the matrix. As an example, in problems where the variables correspond to the edges of an underlying graph, the index could be derived from a lexicographic ordering of the edges (when viewed as ordered pairs of nodes).

This indexing scheme provides a very compact representation, as well as a simple and effective means of moving variables in and out of the active set. However, it means that the user must have a priori knowledge of all problem variables and a method for indexing them. For combinatorial models

such as the *Traveling Salesman Problem*, this does not present a problem. However, for some set partitioning models, for instance, the number of columns may not be known in advance. Even if the number of columns is known in advance, a viable indexing scheme may not be evident. Eliminating the indexing requirement by allowing variables to have abstract, user-defined representations (such as we do for cuts), would allow for more generality, but would also sacrifice some efficiency. A hybrid scheme, allowing the user to have both indexed and *algorithmic* variables (variables with user-defined representations) is planned for a future version of SYMPHONY.

For efficiency, the problem variables can be divided into two sets, the *base variables* and the *extra variables*. The base variables are active in all subproblems, whereas the extra variables can be added and removed. There is no theoretical difference between base variables and extra variables; however, designating a well-chosen set of base variables can significantly increase efficiency. Because they can move in and out of the problem, maintaining extra variables requires additional bookkeeping and computation. If the user has reason to believe a priori that a variable is “good” or has a high probability of having a non-zero value in some optimal solution to the problem, then that variable should be designated as a base variable. It is up to the user to designate which variables should be active in the root subproblem. Typically, when column generation is used, only base variables are active. Otherwise, all variables must be active in the root node.

4.3.2.2 Constraints

Because the global list of potential constraints (also called cuts) is not usually known a priori or is extremely large, constraints cannot generally be represented simply by a user-assigned index. Instead, each constraint is assigned a global index only after it becomes active in some subproblem. It is up to the user, if desired, to designate a compact *representation* for each class of constraints that is to be generated and to implement subroutines for converting from this compact representation to a matrix row, given the list of active variables. For instance, suppose that the set of nonzero variables in a particular class of constraints corresponds to the set of edges across a cut in a graph. Instead of storing the indices of each variable explicitly, one could simply store the set of nodes on one side (“shore”) of the cut as a bit array. The constraint could then be constructed easily for any particular set of active variables (edges).

Just as with variables, the constraints are divided into *core constraints* and *extra constraints*. The core constraints are those that are active in every subproblem, whereas the extra constraints can be generated dynamically and are free to enter and leave as appropriate. Obviously, the set of core constraints must be known and constructed explicitly by the user. Extra constraints, on the other hand, are generated dynamically by the cut generator as they are violated. As with variables, a good set of core constraints can have a significant effect on efficiency.

Note that the user is not *required* to designate a compact representation scheme. Constraints can simply be represented explicitly as matrix rows with respect to the global set of variables. However, designating a compact form can result in large reductions in memory use if the number of variables in the problem is large.

4.3.2.3 Search Tree

Having described the basics of how objects are represented, we now describe the representation of search tree nodes. Since the base constraints and variables are present in every subproblem, only the indices of the extra constraints and variables are stored in each node's description. A complete description of the current basis is maintained to allow a warm start to the computation in each search node. This basis is either inherited from the parent, computed during strong branching (see Section 4.4.2.3), or comes from earlier partial processing of the node itself (see Section 4.4.3.3). Along with the set of active objects, we must also store the identity of the object(s) which were branched upon to generate the node. The branching operation is described in Section 4.4.2.3.

Because the set of active objects and the status of the basis do not tend to change much from parent to child, all of these data are stored as differences with respect to the parent when that description is smaller than the explicit one. This method of storing the entire tree is highly memory-efficient. The list of nodes that are candidates for processing is stored in a heap ordered by a comparison function defined by the search strategy (see 4.4.3). This allows efficient generation of the next node to be processed.

4.3.3 Modular Implementation

SYMPHONY's functions are grouped into five independent computational modules. This modular implementation not only facilitates code maintenance, but also allows easy and highly configurable parallelization. Depending on the computational setting, the modules can be compiled as either (1) a single sequential code, (2) a multi-threaded shared-memory parallel code, or (3) separate processes running in distributed fashion over a network. The modules pass data to each other either through shared memory (in the case of sequential computation or shared-memory parallelism) or through a message-passing protocol defined in a separate communications API (in the case of distributed execution). an schematic overview of the modules is presented in Figure 4.4. In the remainder of the section, we describe the modularization scheme and the implementation of each module in a sequential environment.

4.3.3.1 The Master Module

The *master module* includes functions that perform problem initialization and I/O. This module is the only persistent module and stores all static problem data. The other modules are created only during a solve call and destroyed afterward. All calls to the API are processed through the master module. These functions of the master module implement the following tasks:

- Initialize the environment.
- Set and maintain parameter values.
- Read and store static problem data for instance to be solved.
- Compute an initial upper bound using heuristics.

The Modules of Branch, Cut, and Price




Figure 4.4: Schematic overview of the branch, cut, and price algorithm

- Perform problem preprocessing.
- Initialize the solution process, pass problem information to the solver modules and store the results after completion of the solve call.
- Track the status of associated processes during parallel solution calls.
- Act as a clearing house for output during the solution process.
- Store warm start information between solver calls.
- Service requests from the user through the API for problem data, problem modification, and parameter modification.

4.3.3.2 The Tree Management Module

The *tree manager* controls the overall execution of the algorithm. It tracks the status of its worker modules, as well as that of the search tree, and distributes the subproblems to be processed to the node processing module(s). Functions performed by the tree management module are:

- Receive data for the root node and place it on the list of candidates for processing.
- Receive data for subproblems to be held for later processing.
- Handle requests from linear programming modules to release a subproblem for processing.
- Receive branching object information, set up data structures for the children, and add them to the list of candidate subproblems.
- Keep track of the global upper bound and notify all node processing modules when it changes.
- Write current state information out to disk periodically to allow a restart in the event of a system crash.
- Keep track of run data and send it to the master program at termination.

4.3.3.3 The Node Processing Module

The *node processing* (NP) module is the most complex and computationally intensive of the five processes. Its job is to perform the bounding and branching operations. These operations are, of course, central to the performance of the algorithm. Functions performed by the LP module are:

- Inform the tree manager when a new subproblem is needed.
- Receive a subproblem and process it in conjunction with the cut generator and the cut pool.
- Decide which cuts should be sent to the global pool to be made available to other NP modules.
- If necessary, choose a branching object and send its description back to the tree manager.
- Perform the fathoming operation, including generating variables.

4.3.3.4 The Cut Generation Module

The *cut generator* performs only one function—generating valid inequalities violated by the current fractional solution and sending them back to the requesting LP process. Here are the functions performed by the cut generator module:

- Receive an LP solution and attempt to separate it from the convex hull of all solutions.
- Send generated valid inequalities back to the NP module.
- When finished processing a solution vector, inform the NP module not to expect any more cuts in case it is still waiting.

4.3.3.5 The Cut Management Module

The concept of a *cut pool* was first suggested by Padberg and Rinaldi [29], and is based on the observation that in BCP, the inequalities which are generated while processing a particular node in the search tree are also generally valid and potentially useful at other nodes. Since generating these cuts is usually a relatively expensive operation, the cut pool maintains a list of the “best” or “strongest” cuts found in the tree so far for use in processing future subproblems. Hence, the cut manager functions as an auxiliary cut generator. More explicitly, here are the functions of the cut pool module:

- Receive cuts generated by other modules and store them.
- Receive an LP solution and return a set of cuts which this solution violates.
- Periodically purge “ineffective” and duplicate cuts to control its size.

4.3.4 Algorithm Summary

Currently, SYMPHONY is what is known as a single-pool BCP algorithm. The term *single-pool* refers to the fact that there is a single central list of candidate subproblems to be processed, which is maintained by the tree manager. Most sequential implementations use such a single-pool scheme. However, other schemes may be used in parallel implementations. For a description of various types of parallel branch and bound, see [16].

The user begins by initializing the SYMPHONY environment and can then invoke subroutines for reading in parameters and problem data, finding an initial upper bound, and designating the initial set of active cuts and variables in the root node. Once the user invokes a solve routine, a tree manager is created to manage the solution process. The tree manager module in turn sets up the cut pool module(s), the linear programming module(s), and the cut generator module(s). Currently, there are three solve calls supported by the API. The first call is the *initial solve* (see Section 4.4.1.1), which solves the problem from scratch without using warm start information. The second type of solve call is a *warm solve*, which solves the problem using previously computed warm

start information (see Section 4.4.1.2). Finally, there is a *multicriteria solve* call which is used to enumerate efficient solutions to a given multicriteria MILP (see Section 4.4.1.3).

During the solution process, the tree manager functions control the execution by maintaining the list of candidate subproblems and sending them to the NP modules as they become idle. The NP modules receive nodes from the tree manager, process them, branch (if required), and send back the identity of the chosen branching object to the tree manager, which in turn generates the children and places them on the list of candidates to be processed (see Section 4.4.2.3 for a description of the branching operation). A schematic summary of the algorithm is shown in Figure 4.4. The preference ordering for processing nodes is a run-time parameter. Typically, the node with the smallest lower bound is chosen to be processed next since this strategy minimizes the overall size of the search tree. However, at times, it is advantageous to *dive* down in the tree. The concepts of *diving* and *search chains*, introduced in Section 4.4.3, extend the basic “best-first” approach.

We mentioned earlier that cuts and variables can be treated in a somewhat symmetric fashion. However, it should be clear by now that our current implementation favors the implementation of branch and cut algorithms, where the computational effort spent generating cuts dominates that of generating variables. Our methods of representation also clearly favor such problems. In a future version of the software, we plan to erase this bias by adding additional functionality for handling variable generation and storage. This is the approach already taken by of COIN/BCP [24]. For more discussion of the reasons for this bias and the differences between the treatment of cuts and variables, see Section 4.4.2.2.

4.4 Details of the Implementation

4.4.1 The Master Module

The primary functions performed by the master module were listed in Section 4.3.3.1. Here, we describe the implementational details of the various solve calls.

4.4.1.1 Initial Solve

Calling the initial solve method solves a given MILP from scratch, as described above. The first action taken is to create an instance of the tree manager module that will control execution of the algorithm. If the algorithm is to be executed in parallel on a distributed architecture, the master module spawns a separate tree manager process that will autonomously control the solution process. The tree manager in turn creates the modules for processing the nodes of the search tree, generating cuts, and maintaining cut pools. These modules work in concert to execute the solution process. When it makes sense, sets of two or more modules, such as a node processing module and a cut generation module may be combined to yield a single process in which the combined modules work in concert and communicate with each other through shared memory instead of across the network. When running as separate process, the modules communicate with each other using a standard communications protocol. Currently, the only option supported is PVM, but it would be relatively easy to add an MPI implementation.

The overall flow of the algorithm is similar to other branch and bound implementations and is detailed below. A priority queue of candidate subproblems available for processing is maintained at all times and the candidates are processed in an order determined by the search strategy. The algorithm terminates when the queue is empty or when another specified condition is satisfied. A new feature in SYMPHONY 5.6.18 is the ability to stop the computation based on exceeding a given time limit, exceeding a given limit on the number of processed nodes, achieving a target percentage gap between the upper and lower bounds, or finding the first feasible solution. After halting prematurely, the computation can be restarted after modifying parameters or problem data. This enables the implementation of a wide range of dynamic and on-line solution algorithms, as we describe next.

4.4.1.2 Solve from Warm Start

Among the utility classes in the COIN-OR repository is a base class for describing the data needed to warm start the solution process for a particular solver or class of solvers. To support this option for SYMPHONY, we have implemented such a warm start class for MILPs. The main content of the class is a compact description of the search tree at the time the computation was halted. This description contains complete information about the subproblem corresponding to each node in the search tree, including the branching decisions that lead to the creation of the node, the list of active variables and constraints, and warm start information for the subproblem itself (which is a linear program). All information is stored compactly using SYMPHONY's native data structures, which store only the differences between a child and its parent, rather than an explicit description of every node. This approach reduces the tree's description to a fraction of the size it would otherwise be. In addition to the tree itself, other relevant information regarding the status of the computation is recorded, such as the current bounds and best feasible solution found so far. Using the warm start class, the user can save a warm start to disk, read one from disk, or restart the computation at any point after modifying parameters or the problem data itself. This allows the user to easily implement periodic checkpointing, to design dynamic algorithms in which the parameters are modified after the gap reaches a certain threshold, or to modify problem data during the solution process if needed.

Modifying Parameters. The most straightforward use of the warm start class is to restart the solver after modifying problem parameters. To start the computation from a given warm start when the problem data has not been modified, the tree manager simply traverses the tree and adds those nodes marked as candidates for processing to the node queue. Once the queue has been reformed, the algorithm is then able to pick up exactly where it left off. Code for using the resolve command was shown in Figure 3.2. The situation is more challenging if the user modifies problem data in between calls to the solver. We address this situation next.

Modifying Problem Data. If the user modifies problem data in between calls to the solver, SYMPHONY must make corresponding modifications to the leaf nodes of the current search tree to allow execution of the algorithm to continue. In principle, any change to the original data that does not invalidate the subproblem warm start data, i.e., the basis information for the LP relaxation, can be accommodated. Currently, SYMPHONY can only handle modifications to the rim vectors

of the original MILP. Methods for handling other modifications, such as the addition of columns or the modification of the constraint matrix itself, will be added in the future. To initialize the algorithm, each leaf node, regardless of its status after termination of the previous solve call, must be inserted into the queue of candidate nodes and reprocessed with the changed rim vectors. After this reprocessing, the computation can continue as usual. Optionally, the user can “trim the tree” before resolving. This consists of locating nodes whose descendants are all likely to be pruned in the resolve and eliminating those descendants in favor of processing the parent node itself. This ability could be extended to allow changes that invalidate the warm start data of some leaf nodes.

The ability to resolve after modifying problem data has a wide range of applications in practice. One obvious use is to allow dynamic modification of problem data during the solve procedure, or even after the procedure has been completed. Implementing such a solver is simply a matter of periodically stopping to check for user input describing a change to the problem. Another obvious application is in situations where it is known a priori that the user will be solving a sequence of very similar MILPs. This occurs, for instance, when implementing algorithms for multicriteria optimization, as we describe in Section 4.4.1.3. One approach to this is to solve a given “base problem” (possibly limiting the size of the warm start tree), save the warm start information from the base problem and then start each subsequent call from this same checkpoint. Code for implementing this was shown in Figure 3.3.

4.4.1.3 Bicriteria Solve

For those readers not familiar with bicriteria integer programming, we briefly review the basic notions here. For clarity, we restrict the discussion here to pure integer programs (ILPs), but the principles are easily generalized. A bicriteria ILP is a generalization of a standard ILP presented earlier that includes a second objective function, yielding an optimization problem of the form

$$\begin{array}{ll} \text{vmin} & [cx, dx], \\ \text{s.t.} & Ax \leq b, \\ & x \in \mathbb{Z}^n. \end{array} \quad (4.2)$$

The operator *vmin* is understood to mean that solving this program is the problem of generating *efficient* solutions, which are these feasible solutions p to 4.2 for which there does not exist a second distinct feasible solution q such that $cq \leq cp$ and $dq \leq dp$ and at least one inequality is strict. Note that (4.2) does not have a unique optimal solution value, but a set of pairs of solution values called *outcomes*. The pairs of solution values corresponding to efficient solutions are called *Pareto outcomes*. Surveys of methodology for enumerating the Pareto outcomes of multicriteria integer programs are provided by Climaco et al. [7] and more recently by Ehrgott and Gandibleux [11, 12] and Ehrgott and Wiecek [13].

The bicriteria ILP (4.2) can be converted to a standard ILP by taking a nonnegative linear combination of the objective functions [17]. Without loss of generality, the weights can be scaled so they sum to one, resulting in a family of ILPs parameterized by a scalar $0 \leq \alpha \leq 1$, with the bicriteria objective function replaced by the *weighted sum objective*

$$(\alpha c + (1 - \alpha)d)x. \quad (4.3)$$

Each selection of weight α produces a different single-objective problem. Solving the resulting ILP produces a Pareto outcome called a *supported outcome*, since it is an extreme point on the convex lower envelope of the set of Pareto outcomes. Unfortunately, not all efficient outcomes are supported, so it is not possible to enumerate the set of Pareto outcomes by solving a sequence of ILPs from this parameterized family. To obtain all Pareto outcomes, one must replace the weighted sum objective (4.3) with an objective based on the *weighted Chebyshev norm* studied by Eswaran et al. [14] and Solanki [34]. If x^c is a solution to a weighted sum problem with $\alpha = 1$ and x^d is the solution with $\alpha = 0$, then the weighted Chebyshev norm of a feasible solution p is

$$\max\{\alpha(cp - cx^c), (1 - \alpha)(dp - dx^d)\}. \quad (4.4)$$

Although this objective function is not linear, it can easily be linearized by adding an artificial variable, resulting in a second parameterized family of ILPs. Under the assumption of *uniform dominance*, Bowman showed that an outcome is Pareto if and only if it can be obtained by solving some ILP in this family [5]. In [31], the authors presented a method for enumerating all Pareto outcomes by solving a sequence of ILPs in this parameterized family. By slightly perturbing the objective function, they also showed how to relax the uniform dominance assumption. Note that the set of all supported outcomes, which can be thought of as an approximation of the set of Pareto outcomes, can be similarly obtained by solving a sequence of ILPs with weighted sum objectives.

SYMPHONY 5.6.18 contains a generic implementation of the algorithm described in [31], along with a number of methods for approximating the set of Pareto outcomes. To support these capabilities, we have extended the OSI interface so that it allows the user to define a second objective function. Of course, we have also added a method for invoking this bicriteria solver called `multiCriteriaBranchAndBound()`. Relaxing the uniform dominance requirement requires the underlying ILP solver to have the ability to generate, among all optimal solutions to a ILP with a primary objective, a solution minimizing a given secondary objective. We added this capability to SYMPHONY through the use of optimality cuts, as described in [31].

Because implementing the algorithm requires the solution of a sequence of ILPs that vary only in their objective functions, it is possible to use warm starting to our advantage. Although the linearization of (4.4) requires modifying the constraint matrix from iteration to iteration, it is easy to show that these modifications cannot invalidate the basis. In the case of enumerating all supported outcomes, only the objective function is modified from one iteration to the next. In both cases, we save warm start information from the solution of the first ILP in the sequence and use it for each subsequent computation.

4.4.2 The Node Processing Module

The NP module is at the core of the algorithm, as it performs the processing and bounding operations for each subproblem. A schematic diagram of the node processing loop is presented in Fig. 4.5. The details of the implementation are discussed in the following sections.




Figure 4.5: Overview of the node processing loop

4.4.2.1 The LP Engine

SYMPHONY requires the use of a third-party callable library (referred to as the *LP engine* or *LP library*) to solve the LP relaxations once they are formulated. As with the user functions, SYMPHONY communicates with the LP engine through an API that converts SYMPHONY's internal data structures into those of the LP engine. Currently, the framework will only work with advanced, simplex-based LP engines, such as CPLEX [9], since the LP engine must be able to accept an advanced basis, and provide a variety of data to the framework during the solution process. The internal data structures used for maintaining the LP relaxations are similar to those of CPLEX and matrices are stored in the standard column-ordered format.

4.4.2.2 Managing the LP Relaxation

The majority of the computational effort of BCP is spent solving LPs and hence a major emphasis in the development was to make this process as efficient as possible. Besides using a good LP engine, the primary way in which this is done is by controlling the size of each relaxation, both in terms of number of active variables and number of active constraints.

The number of constraints is controlled through use of a local pool and through purging of ineffective constraints. When a cut is generated by the cut generator, it is first sent to the local cut pool. In each iteration, up to a specified number of the strongest cuts (measured by degree of violation) from the local pool are added to the problem. Cuts that are not strong enough to be added to the relaxation are eventually purged from the list. In addition, cuts are purged from the LP itself when they have been deemed ineffective for more than a specified number of iterations, where ineffective is defined as either (1) the corresponding slack variable is positive, (2) the corresponding slack variable is basic, or (3) the dual value corresponding to the row is zero (or very small). Cuts that have remained effective in the LP for a specified number of iterations are sent to the global pool where they can be used in later search nodes. Cuts that have been purged from the LP can be made active again if they later become violated.

The number of variables (columns) in the relaxation is controlled through *reduced cost fixing* and *dynamic column generation*. Periodically, each active variable is *priced* to see if it can be fixed by reduced cost. That is, the LP reduced cost is examined in an effort to determine whether fixing that variable at one of its bounds would remove improving solutions; if not, the variable is fixed and removed from consideration. If the matrix is *full* at the time of the fixing, meaning that all unfixed variables are active, then the fixing is permanent for that subtree. Otherwise, it is temporary and only remains in force until the next time that columns are dynamically generated.

Because SYMPHONY was originally designed for combinatorial problems with relatively small numbers of variables, techniques for performing dynamic column generation are somewhat unrefined. Currently, variables are priced out sequentially by index, which can be costly. To improve the process of pricing variables, we plan to increase the symmetry between our methods for handling variables and those for handling cuts. This includes (1) allowing user-defined, abstract representations for variables, (2) allowing the use of “variable generators” analogous to cut generators, (3) implementing both global and local pools for variables, (4) implementing heuristics that help determine the order in which the indexed variables should be priced, and (5) allowing for methods

of simultaneously pricing out large groups of variables. Much of this is already implemented in COIN/BCP.

Because pricing is computationally burdensome, it currently takes place only either (1) before branching (optional), or (2) when a node is about to be pruned (depending on the phase—see the description of the two-phase algorithm in Sect. 4.4.3.3). To use dynamic column generation, the user must supply a subroutine which generates the column corresponding to a particular user index, given the list of active constraints in the current relaxation. When column generation occurs, each column not currently active that has not been previously fixed by reduced cost is either priced out immediately, or becomes active in the current relaxation. Only a specified number of columns may enter the problem at a time, so when that limit is reached, column generation ceases. For further discussion of column generation, see Sect. 4.4.3.3, where the two-phase algorithm is described.

Since the matrix is stored in compressed form, considerable computation may be needed to add and remove rows and columns. Hence, rows and columns are only physically removed from the problem when there are sufficiently many to make it “worthwhile.” Otherwise, deleted rows and columns remain in the matrix but are simply ignored by the computation. Note that because ineffective rows left in the matrix increase the size of the basis unnecessarily, it is usually advisable to adopt an aggressive strategy for row removal.

4.4.2.3 Branching

Branching takes place whenever either (1) both cut generation and column generation (if performed) have failed; (2) “tailing off” in the objective function value has been detected; or (3) the user chooses to force branching. Branching can take place on cuts or variables and can be fully automated or fully controlled by the user, as desired. Branching can result in as many children as the user desires, though two is typical. Once it is decided that branching will occur, the user must either select the list of candidates for *strong branching* (see below for the procedure) or allow SYMPHONY to do so automatically by using one of several built-in strategies, such as branching on the variable whose value is farthest from being integral. The number of candidates may depend on the level of the current node in the tree—it is usually best to expend more effort on branching near the top of the tree.

After the list of candidates is selected, each candidate is *pre-solved*, by performing a specified number of iterations of the dual simplex algorithm in each of the resulting subproblems. Based on the objective function values obtained in each of the potential children, the final branching object is selected, again either by the user or by built-in rule. This procedure of using exploratory LP information in this manner to select a branching candidate is commonly referred to as *strong branching*. When the branching object has been selected, the NP module sends a description of that object to the tree manager, which then creates the children and adds them to the list of candidate nodes. It is then up to the tree manager to specify which node the now-idle NP module should process next. This issue is further discussed below.

4.4.3 The Tree Management Module

4.4.3.1 Managing the Search Tree

The tree manager’s primary job is to control the execution of the algorithm by deciding which candidate node should be chosen as the next to be processed. This is done using either one of several built-in rules or a user-defined rule. Usually, the goal of the search strategy is to minimize overall running time, but it is sometimes also important to find good feasible solutions early in the search process. In general, there are two ways to decrease running time—either by decreasing the size of the search tree or by decreasing the time needed to process each search tree node.

To minimize the size of the search tree, the strategy is to select consistently that candidate node with the smallest associated lower bound. In theory, this strategy, sometimes called *best-first*, will lead the smallest possible search tree. However, we need to consider the time required to process each search tree node as well. This is affected by both the quality of the current upper bound and by such factors as communication overhead and node set-up costs. When considering these additional factors, it is sometimes be more effective to deviate from the best-first search order. We discuss the importance of such strategies below.

4.4.3.2 Search Chains and Diving

One reason for not strictly enforcing the search order is because it is somewhat expensive to construct a search node, send it to an NP module, and set it up for processing. If, after branching, we choose to continue processing one of the children of the current subproblem, we avoid the set-up cost, as well as the cost of communicating the node description of the retained child subproblem back to the tree manager. This is called *diving* and the resulting chain of nodes is called a *search chain*. There are a number of rules for deciding when an NP module should be allowed to dive. One such rule is to look at the number of variables in the current LP solution that have fractional values. When this number is low, there may be a good chance of finding a feasible integer solution quickly by diving. This rule has the advantage of not requiring any global information. We also dive if one of the children is “close” to being the best node, where “close” is defined by a chosen parameter.

In addition to the time saved by avoiding reconstruction of the LP in the child, diving has the advantage of often leading quickly to the discovery of feasible solutions, as discussed above. Good upper bounds not only allow earlier pruning of unpromising search chains, but also should decrease the time needed to process each search tree node by allowing variables to be fixed by reduced cost.

4.4.3.3 The Two-Phase Algorithm

If no heuristic subroutine is available for generating feasible solutions quickly, then a unique two-phase algorithm can also be invoked. In the two-phase method, the algorithm is first run to completion on a specified set of core variables. Any node that would have been pruned in the first phase is instead sent to a pool of candidates for the second phase. If the set of core variables is small, but well-chosen, this first phase should be finished quickly and should result in a near-optimal

solution. In addition, the first phase will produce a list of useful cuts. Using the upper bound and the list of cuts from the first phase, the root node is *repriced*—that is, it is reprocessed with the full set of variables and cuts. The hope is that most or all of the variables not included in the first phase will be priced out of the problem in the new root node. Any variable thus priced out can be eliminated from the problem globally. If we are successful at pricing out all of the inactive variables, we have shown that the solution from the first phase was, in fact, optimal. If not, we must go back and price out the (reduced) set of extra variables in each leaf of the search tree produced during the first phase. We then continue processing any node in which we fail to price out all the variables.

In order to avoid pricing variables in every leaf of the tree, we can *trim the tree* before the start of the second phase. Trimming the tree consists of eliminating the children of any node for which each child has lower bound above the current upper bound. We then reprocess the parent node itself. This is typically more efficient, since there is a high probability that, given the new upper bound and cuts, we will be able to prune the parent node and avoid the task of processing each child individually.

4.4.4 The Cut Generation Module

To implement the cut generator process, the user must provide a function that accepts an LP solution and returns cuts violated by that solution to the NP module. In parallel configurations, each cut is returned immediately to the NP module, rather than being passed back as a group once the function exits. This allows the LP to begin adding cuts and solving the current relaxation before the cut generator is finished if desired. Parameters controlling if and when the LP should begin solving the relaxation before the cut generator is finished can be set by the user.

SYMPHONY generates generic cutting planes using the Cut Generator Library, also available from COIN-OR. The CGL can be used to generate cuts in cases where problem-specific cutting planes are not available or not implemented yet.

4.4.5 The Cut Management Module

4.4.5.1 Maintaining and Scanning the Pool

The cut manager’s primary job is to receive a solution from an NP module and return cuts from the pool that are violated by it. The cuts are stored along with two pieces of information—the level of the tree on which the cut was generated, known simply as the *level* of the cut, and the number of times it has been checked for violation since the last time it was actually found to be violated, known as the number of *touches*. The number of touches can be used as a simplistic measure of its effectiveness. Since the pool can get quite large, the user can choose to scan only cuts whose number of touches is below a specified threshold and/or cuts that were generated on a level at or above the current one in the tree. The idea behind this second criterion is to try to avoid checking cuts that were not generated “nearby” in the tree, as they are less likely to be effective. Any cut generated at a level in the tree below the level of the current node must have been generated in a different part of the tree. Although this is admittedly a naive method, it does seem to work reasonably well.

On the other hand, the user may define a specific measure of quality for each cut to be used instead. For example, the degree of violation is an obvious candidate. This measure of quality must be computed by the user, since the cut pool module has no knowledge of the cut data structures. The quality is recomputed every time the user checks the cut for violation and a running average is used as the global quality measure. The cuts in the pool are periodically sorted by this measure and only the highest quality cuts are checked each time. All duplicate cuts, as well as all cuts whose number of touches exceeds or whose quality falls below specified thresholds, are periodically purged from the pool to keep it as small as possible.

4.4.5.2 Using Multiple Pools

For several reasons, it may be desirable to have multiple cut pools. When there are multiple cut pools, each pool is initially assigned to a particular node in the search tree. After being assigned to that node, the pool services requests for cuts from that node and all of its descendants until such time as one of its descendants gets assigned to another cut pool. After that, it continues to serve all the descendants of its assigned node that are not assigned to other cut pools.

Initially, the first cut pool is assigned to the root node. All other cut pools are unassigned. During execution, when a new node is sent to be processed, the tree manager must determine which cut pool the node should be serviced by. The default is to use the same cut pool as its parent. However, if there is currently an idle cut pool process (either it has never been assigned to any node or all the descendants of its assigned node have been processed or reassigned), then that cut pool is assigned to this new node. All the cuts currently in the cut pool of its parent node are copied to the new pool to initialize it, after which the two pools operate independently on their respective subtrees. When generating cuts, the NP module sends the new cuts to the cut pool assigned to service the node during whose processing the cuts were generated.

The primary motivation behind the idea of multiple cut pools is two-fold. First, we want simply to limit the size of each pool as much as possible. By limiting the number of nodes that a cut pool has to service, the number of cuts in the pool will be similarly limited. This not only allows cut storage to spread over multiple processors, and hence increases the available memory, but at the same time, the efficiency with which the cut pool can be scanned for violated cuts is also increased. A secondary reason for maintaining multiple cut pools is that it allows us to limit the scanning of cuts to only those that were generated in the same subtree as the current search node. As described above, this helps focus the search and should increase the efficiency and effectiveness of the search. This idea also allows us to generate locally valid cuts, such as the classical Gomory cuts (see [28]).

4.5 Parallelizing BCP

Because of the clear partitioning of work that occurs when the branching operation generates new subproblems, branch and bound algorithms lend themselves well to parallelization. As a result, there is already a significant body of research on performing branch and bound in parallel environments. We again point the reader to the survey of parallel branch and bound algorithms by Gendron and Crainic [16], as well as other references such as [10, 18, 32, 22].

In parallel BCP, as in general branch and bound, there are two major sources of parallelism. First, it is clear that any number of subproblems on the current candidate list can be processed simultaneously. Once a subproblem has been added to the list, it can be properly processed before, during, or after the processing of any other subproblem. This is not to say that processing a particular node at a different point in the algorithm won't produce different results—it most certainly will—but the algorithm will terminate correctly in any case. The second major source of parallelism is to parallelize the processing of individual subproblems. By allowing separation to be performed in parallel with the solution of the linear programs, we can theoretically process a node in little more than the amount of time it takes to solve the sequence of LP relaxations. Both of these sources of parallelism can be easily exploited using the SYMPHONY framework.

The most straightforward parallel implementation, which is the one we currently employ, is a master-slave model, in which there is a central manager responsible for partitioning the work and parceling it out to the various slave processes that perform the actual computation. The reason we chose this approach is because it allows memory-efficient data structures for sequential computation and yet is conceptually easy to parallelize. Unfortunately, this approach does have limited scalability. For further discussions on the scalability of BCP algorithms and approaches to improving it, see [30] and [36].

4.5.1 Parallel Configurations

SYMPHONY supports numerous configurations, ranging from completely sequential to fully parallel, allowing efficient execution in many different computational settings. As described in the previous section, there are five modules in the standard distributed configuration. Various subsets of these modules can be combined to form separate executables capable of communicating with each other across a network. When two or more modules are combined, they simply communicate through shared-memory instead of through message-passing. However, they are also forced to run in sequential fashion in this case, unless the user chooses to enable threading using an OpenMP compliant compiler (see next section).

As an example, the default distributed configuration includes a separate executable for each module type, allowing full parallelism. However, if cut generation is fast and not memory-intensive, it may not be worthwhile to have the NP module and its associated cut generator work independently, as this increases communication overhead without much potential benefit. In this case, the cut generator functions can be called directly from the NP module, creating a single, more efficient executable.

4.5.2 Inter-process Communication

SYMPHONY can utilize any third-party communication protocol supporting basic message-passing functions. All communication subroutines interface with SYMPHONY through a separate communications API. Currently, PVM [15] is the only message-passing protocol supported, but interfacing with another protocol is a straightforward exercise.

Additionally, it is possible to configure the code to run in parallel using threading to process multiple search tree nodes simultaneously. Currently, this is implemented using OpenMP compiler directives

to specify the parallel regions of the code and perform memory locking functions. Compiling the code with an OpenMP compliant compiler will result in a shared-memory parallel executable. For a list of OpenMP compliant compilers and other resources, visit <http://www.openmp.org>.

4.5.3 Fault Tolerance

Fault tolerance is an important consideration for solving large problems on computing networks whose nodes may fail unpredictably. The tree manager tracks the status of all processes and can restart them as necessary. Since the state of the entire tree is known at all times, the most that will be lost if an NP module or cut generator is killed is the work that had been completed on that particular search node. To protect against the tree manager itself or a cut pool being killed, full logging capabilities have been implemented. If desired, the tree manager can write out the entire state of the tree to disk periodically, allowing a warm restart if a fault occurs. Similarly, the cut pool process can be warm-started from a log file. This not only allows for fault tolerance but also for full reconfiguration in the middle of solving a long-running problem. Such reconfiguration could consist of anything from adding more processors to moving the entire solution process to another network.

Chapter 5

Developing Custom Applications

5.1 Navigating the Source Code

To develop an application with SYMPHONY, you need to first understand how the source files are organized. Note that in this chapter, all path names are given Unix-style. When you unpack the SYMPHONY source distribution, you will notice at the root level a number of files associated with the automatic configuration system, as well as a number of subdirectories, each of which corresponds to a library used by SYMPHONY for some specific functionality. The files associated with SYMPHONY itself are located in the `SYMPHONY` subdirectory. Within the SYMPHONY subdirectory are a number of other subdirectories, including one called `src` containing the source files for SYMPHONY itself.

Also in the main `SYMPHONY/` subdirectory, there is a subdirectory called `Applications/` (see Sections 2.2.2.4 and 2.2.3.4) for instructions on building the applications). The `Applications/` subdirectory contains the source code for a number of sample applications developed with SYMPHONY, as well as function stubs for developing a custom application using SYMPHONY's callbacks. The subdirectory `SYMPHONY/Applications/USER` contains the files needed for implementing the callbacks and is a template for developing an application. In this directory and its subdirectories, which mirror the subdirectories of SYMPHONY itself, each file contains function stubs that can be filled in to create a new custom application. There is a separate subdirectory for each module—master (`Master/`), tree management (`TreeManager/`), cut generation (`CutGen/`), cut management (`CutPool/`), and node processing (`LP/`). Within each subdirectory, there is a file, initially called `USER/*/user_*.c`, where `*` is the name of the module. The primary thing that you, as the user, need to understand to build a custom application is how to fill in these stubs. That is what the second part of this chapter is about. Before describing that, however, we will discuss how to build your application.

5.2 Building an Application

Note that the template application can be built and will work without modifying any of the source files. In this case, SYMPHONY will behave according to default settings.

5.2.1 Unix

First, download and build SYMPHONY as described in Section 2.2.2.4. This will generate the required library and makefiles for each application. After this, typing `make` in the `SYMPHONY/Applications/USER/` subdirectory should successfully build the executable. For more information, including the parallel configuration instructions, see the `SYMPHONY/Applications/USER/INSTALL` file.

5.2.2 Microsoft Visual C++

First, download `SYMPHONY-5.6.18` and unpack the archive if it is required. You now have three options. You can either compile on the command-line using the automated DEVENV build system or NMAKE utility or you can use the provided project and solution files. For all of the following options, first go to the `SYMPHONY\Applications\USER\MSVisualStudio \v8` directory.

5.2.2.1 Using the MSDEV Utility

- Open a command line terminal and type

```
devenv user.sln /Build "Win32|Release"
```

This will create both the release version of the USER application, including the executable `user` and the SYMPHONY library needed for linking with applications.

- To test the executable, type

```
Debug\user.exe -F ..\..\sample.user
```

- If USER source files are modified, type

```
devenv user.sln /make all /rebuild
```

in order to clean and rebuild everything.

5.2.2.2 Using the MSVC++ IDE

- Open the solution file `user.sln`.
- The configuration steps are exactly the same with the MSVC++ section of `SYMPHONY`. The only difference is that, you have the `user` project instead of the `symphony` project. Go through the related steps of section 2.2.3 to see how to get USER executable.

- Once you have the proper settings, choose **Build user.exe** from the **Build** menu. This should successfully build the executable.
- To test the executable, right click on the **user** project, go to the **Debug** tab and set the program arguments to `-F ../..\sample.mps`. Note that command-line switches are Unix-style.
- Now choose **Execute** from the build menu and you have a working branch and bound solver! After successful compilation, you can fill in the user callback functions as describe in Section 5.

5.3 Writing the Callbacks

For each module, all callback functions are invoked from so-called *wrapper functions* that provide the interface and also performs a default action if the user chooses not to override it. Although SYMPHONY is written in C, the wrapper functions provide a C++-style interface in which the user can either accept the default action or override it. Each wrapper function is named `*_u()`, where `*` is the name of the corresponding callback function, and is defined in a file called `*_wrapper.c`. The wrapper function first collects the necessary data and hands it to the user by calling the user function. Based on the return value from the user, the wrapper then performs any necessary post-processing. All callback functions have default options, so that SYMPHONY now acts as a generic MILP solver out of the box.

In Section 6.3, the callback functions are described in detail. The name of every callback function starts with `user_`. There are three kinds of arguments:

IN: An argument containing information that the user might need to perform the function.

OUT: A pointer to an argument in which the user should return a result (requested data, decision, etc.) of the function.

INOUT: An argument which contains information the user might need, but also for which the user can change the value.

The return values for most function are as follows:

Return values:

<code>USER_ERROR</code>	Error in the user function. Printing an error message is the user's responsibility. Depending on the work the user function was supposed to do, the error might be ignored (and some default option used), or the process aborts.
<code>USER_SUCCESS</code>	The user function was implemented and executed correctly.
<code>USER_DEFAULT</code>	This option means that the user function was not implemented and that SYMPHONY should either execute a default subroutine (the default is one of the built-in options—SYMPHONY decides which one to use based on initial parameter settings and the execution of the algorithm) or else do nothing, if execution of the subroutine is optional.
<code>built_in_option1</code> <code>built_in_option2</code> ...	The specified built-in option will be used.

Notes:

- Sometimes an output is optional. This is always noted in the function descriptions.
- If an array has to be returned (i.e., the argument is `type **array`), then (unless otherwise noted) the user has to allocate space for the array itself and set `*array` to be the array allocated. If an output array is optional and the user is not returning any values in that array, then the user *must not* set `*array` because this is how SYMPHONY decides which optional arrays are filled up.
- Some built-in options are implemented so that the user can invoke them directly from the callback function. This might be useful if, for example, the user wants to use different built-in options at different stages of the algorithm.

5.4 Data Structures

The user can define her own data structure for each module to maintain problem data and any other information the user needs access to in order to implement functions to customize the solver. A pointer to this data structure is maintained by SYMPHONY and is passed to the user as an argument to each user function. The pointer must be initially passed using the `sym_set_user_data()` command. Since SYMPHONY knows nothing about this data structure, it is up to the user to allocate it and maintain it. The user must also implement a function to free it. The functions for freeing the user data structures in each module are called `user_free_*`, where `*` is the module. These functions are called by SYMPHONY at the time when other data structures for the modules are being freed and the module is being closed. By default, for sequential computation, there is one common user data structure for all modules and the pointer to that data structure is passed to all user functions, regardless of the module. This setup should work fine for most sequential applications. In parallel, however, pointers cannot be shared between modules and data must be explicitly passed. In this case, it is sometimes more efficient to maintain in each module only the data necessary to perform the functions of that module.

5.5 Parallel Implementation

5.5.1 Distributed-memory Architectures

While the implementation of SYMPHONY strives to shield the user from having to know anything about communications protocols or the specifics of inter-process communication, it may be necessary for the user to pass information from one module to another in order to implement a parallel application. For instance, the user may want to pass data describing the problem instance to the LP process after reading them in from a file in the master process. For the purpose of passing user data from the master process to other processes, a customization function called `user_send*_data()` is provided in the master module, along with a corresponding function called `user_receive*_data()` in the module `*`. These two functions work in tandem to transport the user's data from the master, where it can be read in from a file, to the proper module for processing. There are also a number of other tandem pairs of *send* and *receive* functions that are used to transport user data from place to place.

All data are sent in the form of arrays of either type `char`, `int`, or `double`, or as strings. To send an array, the user has simply to invoke the function `send_XXX_array(XXX *array, int length)` where `XXX` is one of the previously listed types. To receive that array, there is a corresponding function called `receive_?_array(? *array, int length)`. When receiving an array, the user must first allocate the appropriate amount of memory. In cases where variable length arrays need to be passed, the user must first pass the length of the array (as a separate array of length one) and then the array itself. In the receive function, this allows the length to be received first so that the proper amount of space can be allocated before receiving the array itself. Note that data must be received in exactly the same order as it was passed, as data is read linearly into and out of the message buffer. The easiest way to ensure this is done properly is to simply copy the send statements into the receive function and change the function names. It may then be necessary to add some allocation statements in between the receive function calls.

5.5.2 Shared-memory Architectures

In the shared memory configuration, it is not necessary to use message passing to move information from one module to another since memory is globally accessible. In the few cases where the user would ordinarily have to pass information using message passing, it is easiest and most efficient to simply copy the information to the new location. This copying gets done in the *send* function and hence the *receive* function is never actually called. This means that the user must perform all necessary initialization, etc. in the send function. This makes it a little confusing to write source code which will work for all configurations. However, the confusion should be minimized by looking at the sample applications, especially the VRP solver, which works in all configurations, sequential, distributed parallel, and shared parallel.

5.6 Debugging Your Application

Much of this section applies to Unix operating systems. However, it may also be useful for Windows users.

5.6.1 The First Rule

SYMPHONY has many built-in options to make debugging easier. The most important one, however, is the following rule. **It is easier to debug the fully sequential version than the fully distributed version.** Debugging parallel code is not terrible, but it is more difficult to understand what is going on when you have to look at the interaction of several different modules running as separate processes. This means multiple debugging windows which have to be closed and restarted each time the application is re-run. For this reason, it is highly recommended to develop code that can be compiled serially even if you eventually intend to run in a fully distributed environment. This does make the coding marginally more complex, but believe me, it's worth the effort. The vast majority of your code will be the same for either case. Make sure to use the configuration flag to `--enable-debug` while building (see Section 2.2).

5.6.2 Debugging with PVM

If you wish to venture into debugging your distributed application, then you simply need to set the parameter `*_debug`, where `*` is the name of the module you wish to debug, to “1” in the parameter file. This will tell PVM to spawn the particular process or processes in question under a debugger. What PVM actually does in this case is to launch the script `$PVM_ROOT/lib/debugger`. You will undoubtedly want to modify this script to launch your preferred debugger in the manner you deem fit. If you have trouble with this, please send e-mail to the list serve (see Section 1.6).

It's a little tricky to debug interacting parallel processes. The main difficulty is in that the order of operations is difficult to control. Random interactions can occur when processes run in parallel due to varying system loads, process priorities, etc. Therefore, it may not always be possible to duplicate errors. To force runs that you should be able to reproduce, make sure the parameter `no_cut_timeout` appears in the parameter file or start SYMPHONY with the `-a` option. This will keep the cut generator from timing out, a major source of randomness. Furthermore, run with only one active node allowed at a time (set `max_active_nodes` to “1”). This will keep the tree search from becoming random. These two steps should allow runs to be reproduced. You still have to be careful, but this should make things easier.

5.6.3 Checking the Validity of Cuts and Tracing the Optimal Path

Sometimes the only evidence of a bug is the fact that the optimal solution to a particular problem is never found. This is usually caused by either (1) adding an invalid cut, or (2) performing an invalid branching. There are two options available for discovering such errors. The first is for checking the validity of added cuts. This checking must, of course, be done by the user, but SYMPHONY can facilitate such checking. To do this, the user must fill in the function

`user_check_validity_of_cut()` (see Section 6.3.3). THIS function is called every time a cut is passed from the cut generator to the LP and can function as an independent verifier. To do this, the user must pass (through her own data structures) a known feasible solution. Then for each cut passed into the function, the user can check whether the cut is satisfied by the feasible solution. If not, then there is a problem! Of course, the problem could also be with the checking routine. To enable this functionality, the user must configure SYMPHONY with the flag `--enable-cut-check` (see Section 2.2).

Tracing the optimal path can alert the user when the subproblem which admits a particular known feasible solution (at least according to the branching restrictions that have been imposed so far) is pruned. This could be due to an invalid branching. Note that this option currently only works for branching on binary variables. To use this facility, the user must fill in the function `user_send_feas_sol()` (see Section 6.3.1). All that is required is to pass out an array of user indices that are in the feasible solution that you want to trace. Each time the subproblem which admits this feasible solution is branched on, the branch that continues to admit the solution is marked. When one of these marked subproblems is pruned, the user is notified. To enable this functionality, the user must configure SYMPHONY with the flag `--enable-trace-path` (see Section 2.2).

5.6.4 Using the Interactive Graph Drawing Software

The Interactive Graph Drawing (IGD) software package is included with SYMPHONY and SYMPHONY facilitates its use through interfaces with the package. The package, which is a Tcl/Tk application, is extremely useful for developing and debugging applications involving graph-based problems. Given display coordinates for each node in the graph, IGD can display support graphs corresponding to fractional solutions with or without edge weights and node labels and weights, as well as other information. Furthermore, the user can interactively modify the graph by, for instance, moving the nodes apart to “disentangle” the edges. The user can also interactively enter violated cuts through the IGD interface.

To use IGD, you must have installed PVM since the drawing window runs as a separate application and communicates with the user’s routines through message passing. To compile the graph drawing application, type `make dg` in the SYMPHONY root directory. The user routines in the file `user_dg.c` can be filled in, but it is not necessary to fill anything in for basic applications.

After compiling `dg`, the user must write some subroutines that communicate with `dg` and cause the graph to be drawn. Regrettably, this is currently a little more complicated than it needs to be and is not well documented. However, by looking at the sample application, it should be possible to see how it is done. To enable graph drawing, put the line `do_draw_graph 1` into the parameter file or use the `-d` command line option. It can be difficult to get IGD to work. If you are interested in using it and cannot get it to work, feel free to contact me.

5.6.5 Other Debugging Techniques

Another useful built-in function is `write_mps()`, which will write the current LP relaxation to a file in MPS format. This file can then be read into the LP solver interactively or examined by hand for

errors. Many times, CPLEX gives much more explicit error messages interactively than through the callable library. The form of the function is

```
void write_mps(LPdata *lp_data, char *fname)
```

where `fname` is the name of the file to be written. If SYMPHONY is forced to abandon solution of an LP because the LP solver returns an error code, the current LP relaxation is automatically written to the file `matrix.[bc_index].[iter_num].mps` where `bc_index` is the index of the current subproblem and `iter_num` is the current iteration number. The `write_mps()` function can be called using breakpoint code to examine the status of the matrix at any point during execution.

Logging is another useful feature. Logging the state of the search tree can help isolate some problems more easily. See Section 6.4.4 for the appropriate parameter settings to use logging.

5.7 Case Study: Implementing a Matching Solver

This section was contributed by Michael Trick a few years ago and is a walkthrough of the steps for developing a very simple application using SYMPHONY. Rather than presenting the code in its final version, we will go through the steps that a user would go through. Note that some of the code is lifted from the vehicle routing application. This code is designed to be a sequential code. The MATCH application discussed here is part of the SYMPHONY distribution and the source code can be found in the `SYMPHONY/Applications/MATCH` directory.

The goal is to create a minimum matching on a complete graph. Initially, we will just formulate this as an integer program with one variable for each possible pair that can be matched. Then we will include a set of constraints that can be added by cut generation.

We begin with the template code in the `USER` subdirectory included with SYMPHONY. This gives stubs for each user callback routine. First, I need to define a data structure for describing an instance of the matching problem. We use the template structure `USER_PROBLEM` in the file `include/user.h` for this purpose. To describe an instance, we just need the number of nodes and the cost matrix. In addition, we also need a way of assigning an index to each possible assignment. Here is the data structure:

```
typedef struct USER_PROBLEM{
    int          numnodes;
    int          cost[MAXNODES][MAXNODES];
    int          match1[MAXNODES*(MAXNODES-1)/2];
    int          match2[MAXNODES*(MAXNODES-1)/2];
    int          index[MAXNODES][MAXNODES];
}user_problem;
```

The fields `match1`, `match2`, and `index` will be used later in the code in order to map variables to the corresponding assignment and vice versa.

Next, we need to read in the problem instance. We could implement this function within the `user_io()` callback function (see the file `user_master.c`). However, in order to show how it can

be done explicitly, we will define our own function `match_read_data()` in `user_main.c` to fill in the user data structure and then use `sym_set_user_data()` to pass this structure to SYMPHONY. The template code already provides basic command-line options for the user. The “-F” flag is used to specify the location of a data file, from which we will read in the data. The datafile contains first the number of nodes in the graph (`nnodes`) followed by the pairwise cost matrix (nnode by nnode). We read the file in with the `match_read_data()` routine in `user_main.c`:

```
int match_read_data(user_problem *prob, char *infile)
{
    int i, j;
    FILE *f = NULL;

    if ((f = fopen(infile, "r")) == NULL){
        printf("main(): user file %s can't be opened\n", infile);
        return(ERROR__USER);
    }

    /* Read in the costs */
    fscanf(f, "%d", &(prob->numnodes));
    for (i = 0; i < prob->numnodes; i++)
        for (j = 0; j < prob->numnodes; j++)
            fscanf(f, "%d", &(prob->cost[i][j]));

    return (FUNCTION_TERMINATED_NORMALLY);
}
```

We can now construct the integer program itself. This is done by specifying the constraint matrix and the rim vectors in sparse format. We will have a variable for each possible assignment (i, j) with $i < j$. We have a constraint for each node i , so it can only be matched to one other node.

We define the IP in our other helper function `match_load_problem()` in `user_main.c`. In the first part of this routine, we will build a description of the IP, and then in the second part, we will load this representation to SYMPHONY through `sym_explicit_load_problem()`. Note that we could instead create a description of each subproblem dynamically using the `user_create_subproblem()` callback (see `user_lp.c`), but this is more complicated and unnecessary here.

```
int match_load_problem(sym_environment *env, user_problem *prob){

    int i, j, index, n, m, nz, *matbeg, *matind;
    double *matval, *lb, *ub, *obj, *rhs, *rngval;
    char *sense, *is_int;

    /* set up the initial LP data */
    n = prob->numnodes*(prob->numnodes-1)/2;
    m = 2 * prob->numnodes;
```

```

nz = 2 * n;

/* Allocate the arrays */
matbeg  = (int *) malloc((n + 1) * ISIZE);
matind  = (int *) malloc((nz) * ISIZE);
matval  = (double *) malloc((nz) * DSIZE);
obj      = (double *) malloc(n * DSIZE);
lb      = (double *) calloc(n, DSIZE);
ub      = (double *) malloc(n * DSIZE);
rhs      = (double *) malloc(m * DSIZE);
sense   = (char *) malloc(m * CSIZE);
rngval  = (double *) calloc(m, DSIZE);
is_int  = (char *) malloc(n * CSIZE);

/* Fill out the appropriate data structures -- each column has
   exactly two entries */
index = 0;
for (i = 0; i < prob->numnodes; i++) {
    for (j = i+1; j < prob->numnodes; j++) {
        prob->match1[index] = i; /*The first component of assignment 'index'*/
        prob->match2[index] = j; /*The second component of assignment 'index'*/
        /* So we can recover the index later */
        prob->index[i][j] = prob->index[j][i] = index;
        obj[index] = prob->cost[i][j]; /* Cost of assignment (i, j) */
        is_int[index] = TRUE;
        matbeg[index] = 2*index;
        matval[2*index] = 1;
        matval[2*index+1] = 1;
        matind[2*index] = i;
        matind[2*index+1] = j;
        ub[index] = 1.0;
        index++;
    }
}
matbeg[n] = 2 * n;

/* set the initial right hand side */
for (i = 0; i < m; i++) {
    rhs[i] = 1;
    sense[i] = 'E';
}

/* Load the problem to SYMPHONY */
sym_explicit_load_problem(env, n, m, matbeg, matind, matval, lb, ub,
                          is_int, obj, 0, sense, rhs, rngval, true);

```

```

    return (FUNCTION_TERMINATED_NORMALLY);

}

```

Now, we are ready to gather everything in the `main()` routine in `user_main()`. This will involve to create a SYMPHONY environment and a user data structure, read in the data, create the corresponding IP, load it to the environment and ask SYMPHONY to solve it (`CALL_FUNCTION` is just a macro to take care of the return values):

```

int main(int argc, char **argv)
{
    int termcode;
    char * infile;

    /* Create a SYMPHONY environment */
    sym_environment *env = sym_open_environment();

    /* Create the data structure for storing the problem instance.*/
    user_problem *prob = (user_problem *)calloc(1, sizeof(user_problem));

    CALL_FUNCTION( sym_set_user_data(env, (void *)prob) );
    CALL_FUNCTION( sym_parse_command_line(env, argc, argv) );
    CALL_FUNCTION( sym_get_str_param(env, "infile_name", &infile));
    CALL_FUNCTION( match_read_data(prob, infile) );
    CALL_FUNCTION( match_load_problem(env, prob) );
    CALL_FUNCTION( sym_solve(env) );
    CALL_FUNCTION( sym_close_environment(env) );
    return(0);
}

```

OK, that's it. That defines an integer program, and if you compile and optimize it, the rest of the system will come together to solve this problem. Here is a data file to use:

```

6
0 1 1 3 3 3
1 0 1 3 3 3
1 1 0 3 3 3
3 3 3 0 1 1
3 3 3 1 0 1
3 3 3 1 1 0

```

The optimal value is 5. To display the solution, we need to be able to map back from variables to the nodes. That was the use of the `node1` and `node2` parts of the `USER_PROBLEM`. We can now use `user_display_solution()` in `user_master.c` to print out the solution:

```

int user_display_solution(void *user, double lpetol, int varnum, int *indices,
                        double *values, double objval)
{
    /* This gives you access to the user data structure. */
    user_problem *prob = (user_problem *) user;
    int index;

    for (index = 0; index < varnum; index++){
        if (values[index] > lpetol) {
            printf("%2d matched with %2d at cost %6d\n",
                prob->node1[indices[index]],
                prob->node2[indices[index]],
                prob->cost[prob->node1[indices[index]]]
                [prob->node2[indices[index]]]);
        }
    }

    return(USER_SUCCESS);
}

```

We will now update the code to include a crude cut generator. Of course, We could go for a Gomory-Hu type odd-set separation (ala Grötschel and Padberg) but for the moment, let's just check for sets of size three with more than value 1 among them (such a set defines a cut that requires at least one edge out of any odd set). We can do this by brute force checking of triples, as follows:

```

int user_find_cuts(void *user, int varnum, int iter_num, int level,
                  int index, double objval, int *indices, double *values,
                  double ub, double etol, int *num_cuts, int *alloc_cuts,
                  cut_data ***cuts)
{
    user_problem *prob = (user_problem *) user;
    double edge_val[200][200]; /* Matrix of edge values */
    int i, j, k, cutind[3];
    double cutval[3];

    int cutnum = 0;

    /* Allocate the edge_val matrix to zero (we could also just calloc it) */
    memset((char *)edge_val, 0, 200*200*ISIZE);

    for (i = 0; i < varnum; i++) {
        edge_val[prob->node1[indices[i]]][prob->node2[indices[i]]] = values[i];
    }
}

```

```

for (i = 0; i < prob->nnodes; i++){
    for (j = i+1; j < prob->nnodes; j++){
        for (k = j+1; k < prob->nnodes; k++) {
            if (edge_val[i][j]+edge_val[j][k]+edge_val[i][k] > 1.0 + etol) {
                /* Found violated triangle cut */
                /* Form the cut as a sparse vector */
                cutind[0] = prob->index[i][j];
                cutind[1] = prob->index[j][k];
                cutind[2] = prob->index[i][k];
                cutval[0] = cutval[1] = cutval[2] = 1.0;
                cg_add_explicit_cut(3, cutind, cutval, 1.0, 0, 'L',
                                   TRUE, num_cuts, alloc_cuts, cuts);

                cutnum++;
            }
        }
    }
}

return(USER_SUCCESS);
}

```

Note the call of `cg_add_explicit_cut()`, which tells SYMPHONY about any cuts found. If we now solve the matching problem on the sample data set, the number of nodes in the branch and bound tree should just be 1 (rather than 3 without cut generation).

Chapter 6

Reference

6.1 Callable Library C API

This chapter specifies the interface for using SYMPHONY's callable library. These function calls can be used to build custom applications that call SYMPHONY as a subroutine, as described in Section 3.3. All callable library function begin with the prefix `sym_`. To call these function from an application, include the header file `symphony.h` and then link with the SYMPHONY library as described in Section 2. In general, if an array is requested, such as the array of lower bounds on the variables, for instance, the user is responsible for allocating an array of appropriate size and passing it to SYMPHONY. SYMPHONY will then fill up the array.

6.1.1 Primary Interface Functions

▷ **sym_open_environment**

```
sym_environment *sym_open_environment()
```

Description:

This routine is used to get a new SYMPHONY environment to be passed as an argument to all other API subroutines. This routine also invokes the callback function `user_initialize()` (see Section 6.3.1).

Return values:

NULL	Error. Environment could not be initialized. None of the other API subroutines can be called after this point.
sym_environment *	Pointer to a successfully opened environment

▷ **sym_create_copy_environment**

```
sym_environment *sym_create_copy_environment(sym_environment *env)
```

Description:

This routine is used to copy the given environment.

Arguments:

sym_environment *env IN Pointer to the SYMPHONY environment.

Return values:

NULL An empty environment is passed in.

SYM_ENVIRONMENT * Pointer to the copy of the environment.

▷ **sym_parse_command_line**

```
int sym_parse_command_line(sym_environment *env, int argc, char **argv)
```

Description:

This routine parses the command line arguments. It must be called whenever the user specifies any of SYMPHONY's built-in command-line switches. For instance, this is the case when the user specifies the location of an MPS, LP, or GMPL file using the `-F` or `-L` switch or when the user specifies the location of a parameter file with the `-f` switch. This command also invokes the user callback function `user_readparams()` (see Section 6.3.1).

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int argc</code>	IN	The number of command line arguments.
<code>char **argv</code>	IN	Array of pointers to these arguments.

Return values:

<code>ERROR__USER</code>	Error. User error detected in <code>user_readparams()</code> function.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_find_initial_bounds**

```
int sym_find_initial_bounds(sym_environment *env)
```

Description:

This routine invokes the user callback `user_start_heurs()` (see Section 6.3.1) to set the priori bound for the problem.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
-----------------------------------	-------	--------------------------------------

Return values:

<code>ERROR_USER</code>	Error. User error detected in <code>user_start_heurs()</code> function.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_load_problem**

```
int sym_load_problem(sym_environment *env)
```

Description:

This routine loads the description of the problem given in MPS or GMPL/AMPL format or in a file read by a custom file parser implemented in the `user_io()` (see Section 6.3.1) callback. If the problem is to be loaded from an MPS or a GMPL/AMPL file whose location is specified on the command line, then the `sym_parse_command_line()` function has to be invoked beforehand. This function also invokes the user callback `user_initialize_root_node()` (see Section 6.3.1). Note that if the user wishes to load the problem manually without implementing a callback or using one of SYMPHONY's built-in parsers (as is typically done in other callable libraries), then the `sym_explicit_load_problem()` routine should be used.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
-----------------------------------	-------	--------------------------------------

Return values:

<code>ERROR__USER</code>	Error. User error detected in one of <code>user_io()</code> and <code>user_init_draw()</code> functions.
<code>ERROR__READING_GMPL_FILE</code>	Error detected in the given GMPL/AMPL file.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_explicit_load_problem**

```
int sym_explicit_load_problem_user(sym_environment * env, int numcols,
                                   int numrows, int *start, int *index, double *value,
                                   double *collb, double *colub, char *is_int,
                                   double *obj, double *obj2, char *rowsen,
                                   double *rowrhs, double *rowrng, char make_copy)
```

Description:

This routine is used to load a problem description into SYMPHONY manually. The constraint matrix is passed in a standard column-ordered format. The arguments here are the same as the fields in the `MIPdesc` data structure discussed in Section 6.3.2.1. Please see the discussion there for a more detailed description of the arguments here.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int numcols</code>	IN	Number of the columns.
<code>int numrows</code>	IN	Number of the rows.
<code>int *start</code>	IN	Array of the starting positions of each of column.
<code>int *index</code>	IN	Array of the row indices corresponding to each entry of <code>value</code> .
<code>int *value</code>	IN	Array of the values of nonzero entries of the constraint matrix in <i>column order</i> .
<code>double *collb</code>	IN	Array of the lower bounds of the columns.
<code>double *colub</code>	IN	Array of the upper bounds of the columns.
<code>double *obj</code>	IN	Array of the objective function coefficients.
<code>double *obj2</code>	IN	Array of the second objective function coefficients when multi criteria solver is to be used.
<code>char *rowsen</code>	IN	Array of the senses of the constraints. 'L': \leq constraint 'E': $=$ constraint 'G': \geq constraint 'R': ranged constraint 'N': free constraint
<code>double *rowrhs</code>	IN	Array of the right hand side values.
<code>double *rowrng</code>	IN	Array of the row ranges. (<code>sym_get_row_upper</code>) - (<code>sym_get_row_lower</code>) if the row sense is 'R', 0 otherwise.
<code>char make_copy</code>	IN	SYMPHONY will create the copies of these arrays for internal usage if this flag is set to true, otherwise, will own them.

Return values:

<code>ERROR_USER</code>	Error. User error detected in <code>user_initialize_root_node</code> function.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_read_mps**

```
int sym_read_mps(sym_environment *env, char *infile)
```

Description: This routine is used to load an instance from an MPS file.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>char *infile</code>	IN	Pointer to a character array indicating the name of the file.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_read_gmpl**

```
int sym_read_gmpl(sym_environment *env, char *modelfile, char *datafile)
```

Description: This routine is used to load an instance from a GMPL file.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>char *modelfile</code>	IN	Pointer to a character array indicating the name of the model file.
<code>char *datafile</code>	IN	Pointer to a character array indicating the name of the data file.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_solve**

```
int sym_solve(sym_environment *env)
```

Description:

This routine solves the currently loaded MILP problem from scratch even in the presence of a loaded warm start. Any warm start information loaded or kept before will be deleted from the environment!

Arguments:

`sym_environment *env` INOUT Pointer to the SYMPHONY environment.

Return values:

<code>ERROR__USER</code>	Error. User error detected in one of <code>user_send_lp_data()</code> , <code>user_send_cg_data()</code> , <code>user_send_cp_data()</code> , <code>user_receive_feasible_solution()</code> , <code>user_display_solution()</code> and <code>user_process_own_messages()</code> functions.
<code>TM_OPTIMAL_SOLUTION_FOUND</code>	Tree Manager (TM) found the optimal solution and stopped.
<code>TM_TIME_LIMIT_EXCEEDED</code>	TM stopped after reaching the predefined time limit.
<code>TM_NODE_LIMIT_EXCEEDED</code>	TM stopped after reaching the predefined node limit.
<code>TM_TARGET_GAP_ACHIEVED</code>	TM stopped after achieving the predefined target gap.
<code>TM_FOUND_FIRST_FEASIBLE</code>	TM stopped after finding the first feasible solution.
<code>TM_ERROR__NO_BRANCHING_CANDIDATE</code>	Error. TM stopped. User didn't select branching candidate in <code>user_select_candidates()</code> callback.
<code>TM_ERROR__ILLEGAL_RETURN_CODE</code>	Error. TM stopped after getting a non-valid return code.
<code>TM_ERROR__NUMERICAL_INSTABILITY</code>	Error. TM stopped due to some numerical difficulties.
<code>TM_ERROR__COMM_ERROR</code>	Error. TM stopped due to communication error.
<code>TM_ERROR__USER</code>	Error. TM stopped. User error detected in one of user callbacks called during TM processes.

▷ **sym_warm_solve**

```
int sym_warm_solve(sym_environment *env)
```

Description:

This routine re-solves the corresponding problem after some of the parameters have been changed or problem data has been modified from a warm start. If the user plans to invoke this routine, the **keep_warm_start** parameter must be set to **TRUE** before the initial call to the **sym_solve()** routine, so that SYMPHONY will collect the necessary warm starting information during the solve procedure.

Arguments:

sym_environment *env INOUT Pointer to the SYMPHONY environment.

Return values:

ERROR__USER	Error. User error detected in one of user_send_lp_data , user_send_cg_data , user_send_cp_data , user_receive_feasible_solution , user_display_solution and user_process_own_messages functions.
TM_OPTIMAL_SOLUTION_FOUND	Tree Manager (TM) found the optimal solution and stopped.
TM_TIME_LIMIT_EXCEEDED	TM stopped after reaching the predefined time limit.
TM_NODE_LIMIT_EXCEEDED	TM stopped after reaching the predefined node limit.
TM_TARGET_GAP_ACHIEVED	TM stopped after achieving the predefined target gap.
TM_FOUND_FIRST_FEASIBLE	TM stopped after finding the first feasible solution.
TM_ERROR__NO_BRANCHING_CANDIDATE	Error. TM stopped. User didn't select branching candidate in user_select_candidates callback
TM_ERROR__ILLEGAL_RETURN_CODE	Error. TM stopped after getting a non-valid return code.
TM_ERROR__NUMERICAL_INSTABILITY	Error. TM stopped due to some numerical difficulties.
TM_ERROR__COMM_ERROR	Error. TM stopped due to communication error.
TM_ERROR__USER	Error. TM stopped. User error detected in one of user callbacks called during TM processes.
FUNCTION_TERMINATED_ABNORMALLY	Function invoked unsuccessfully.

▷ **sym_mc_solve**

```
int sym_mc_solve(sym_environment *env)
```

Description:

This routine is used to solve the loaded problem as a multicriteria problem. For this function, a second objective function must be set either by calling the `sym_set_obj2_coeff()` function or by passing it directly using the `sym_explicit_load_problem()` function.

Arguments:

`sym_environment *env` INOUT Pointer to the SYMPHONY environment.

Return values:

<code>ERROR__USER</code>	Error. User error detected in one of <code>user_send_lp_data()</code> , <code>user_send_cg_data()</code> , <code>user_send_cp_data()</code> , <code>user_receive_feasible_solution()</code> , <code>user_display_solution()</code> , <code>user_process_own_messages()</code> functions.
<code>TM_OPTIMAL_SOLUTION_FOUND</code>	The set of supported or nondominated solutions have been found.
<code>TM_ERROR__NO_BRANCHING_CANDIDATE</code>	Error. TM stopped. User didn't select branching candidate in <code>user_select_candidates</code> callback
<code>TM_ERROR__ILLEGAL_RETURN_CODE</code>	Error. TM stopped after getting a non-valid return code.
<code>TM_ERROR__NUMERICAL_INSTABILITY</code>	Error. TM stopped due to some numerical difficulties.
<code>TM_ERROR__COMM_ERROR</code>	Error. TM stopped due to communication error.
<code>TM_ERROR__USER</code>	Error. TM stopped. User error detected in one of user callbacks activated by user and invoked during TM processes.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_create_permanent_cut_pools**

```
int sym_create_permanent_cut_pools(sym_environment *env, int *cp_num)
```

Description:

This routine is used to create a global cut pool that will be saved even after the solve call exits and can be used to initialize the cut pool for later solve calls. This can be useful when solving a series of related MILPs that share classes of globally valid inequalities. For instance, if only the objective function is varied, as is the case with multicriteria integer programming, then cuts can be saved for use in later solve calls.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int *cp_num</code>	OUT	Pointer to an integer indicating the number of cut pools stored in the environment.

Return values:

INT The number of the cut pools created.

▷ **sym_set_user_data**

```
int sym_set_user_data(sym_environment *env, void *user)
```

Description:

This routine is used to give SYMPHONY a pointer to the user's problem data structure. This pointer will then be handed back to the user during subsequent calls to user callbacks. This allows the user to store static problem data. Note that this pointer can also be stored by filling out the callback function `user_initialize()` (see Section 6.3.1).

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>void *user</code>	IN	Pointer to the user defined problem structure.

Return values:

<code>ERROR__USER</code>	Error in the passed in <code>user</code> structure.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_user_data**

```
int sym_get_user_data(sym_environment *env, void **user)
```

Description: This routine is used to get the user's problem data structure from SYMPHONY environment.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>void **user</code>	OUT	Pointer to the user defined problem structure.

Return values:

<code>ERROR__USER</code>	Error in the passed in <code>user</code> structure.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_close_environment**

```
int sym_close_environment(sym_environment *env)
```

Description:

This routine closes the environment and returns the allocated memory.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
-----------------------------------	-------	--------------------------------------

Return values:

<code>ERROR__USER</code>	Error. User error detected in <code>user_free_master()</code> function.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

6.1.2 Parameter Query and Modification

▷ **sym_set_defaults**

```
int sym_set_defaults(sym_environment *env)
```

Description:

This routine sets all the environment variables and parameters to their default values.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment to be modified.
-----------------------------------	-------	---

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_set_int_param**

```
void sym_set_int_param(sym_environment *env, char *key, int value)
```

Description:

This routine is used to set an integer type parameter.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>char *key</code>	IN	The name of the parameter to be set.
<code>int value</code>	OUT	New value of the corresponding parameter.

▷ **sym_set_dbl_param**

```
void sym_set_int_param(sym_environment *env, char *key, double value)
```

Description:

This routine is used to set a double type parameter.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>char *key</code>	IN	The name of the parameter to be set.
<code>double value</code>	OUT	New value of the corresponding parameter.

▷ **sym_set_str_param**

```
void sym_set_str_param(sym_environment *env, char *key, char *value)
```

Description:

This routine is used to set a string type parameter.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>char *key</code>	IN	The name of the parameter to be set.
<code>char *value</code>	OUT	New value of the corresponding parameter.

▷ **sym_get_int_param**

```
int sym_get_int_param(sym_environment *env, char *key)
```

Description:

This routine is used to get the value of an integer type parameter.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>char *key</code>	IN	The name of the parameter.

Return values:

INT An integer indicating the value of the parameter.

▷ **sym_get_dbl_param**

```
double sym_get_int_param(sym_environment *env, char *key)
```

Description:

This routine is used to get the value of a double type parameter.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>char *key</code>	IN	The name of the parameter.

Return values:

DOUBLE A double indicating the value of the parameter.

▷ **sym_get_str_param**

```
char *sym_get_int_param(sym_environment *env, char *key)
```

Description:

This routine is used to get the value of a string type parameter.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>char *key</code>	IN	The name of the parameter.

Return values:

`CHAR*` A character array indicating the value of the parameter.

6.1.3 Solver Status Query Functions

▷ **sym_get_status**

```
int sym_get_status(sym_environment *env)
```

Description:

This post-solution query routine is used to learn the termination status of the solution procedure.

Arguments: `sym_environment *env` IN Pointer to the SYMPHONY environment.

Return values:

<code>ERROR__USER</code>	Error. User error detected in one of <code>user_send_lp_data()</code> , <code>user_send_cg_data()</code> , <code>user_send_cp_data()</code> , <code>user_receive_feasible_solution()</code> , <code>user_display_solution()</code> , <code>user_process_own_messages()</code> functions.
<code>TM_OPTIMAL_SOLUTION_FOUND</code>	Tree Manager (TM) found the optimal solution and stopped.
<code>TM_TIME_LIMIT_EXCEEDED</code>	TM stopped after reaching the predefined time limit.
<code>TM_NODE_LIMIT_EXCEEDED</code>	TM stopped after reaching the predefined node limit.
<code>TM_TARGET_GAP_ACHIEVED</code>	TM stopped after achieving the predefined target gap.
<code>TM_FOUND_FIRST_FEASIBLE</code>	TM stopped after finding the first feasible solution.
<code>TM_ERROR__NO_BRANCHING_CANDIDATE</code>	Error. TM stopped. User didn't select branching candidate in <code>user_select_candidates()</code> callback
<code>TM_ERROR__ILLEGAL_RETURN_CODE</code>	Error. TM stopped after getting an invalid return code.
<code>TM_ERROR__NUMERICAL_INSTABILITY</code>	Error. TM stopped due to some numerical difficulties.
<code>TM_ERROR__COMM_ERROR</code>	Error. TM stopped due to communication error.
<code>TM_ERROR__USER</code>	Error. TM stopped. User error detected in one of user callbacks called during TM processes.

▷ **sym_is_proven_optimal**

```
int sym_is_proven_optimal(sym_environment *env)
```

Description:

This post-solution query routine is used to learn whether the problem was solved to optimality.

Arguments:

`sym_environment *env` IN Pointer to the SYMPHONY environment.

Return values:

TRUE The problem was solved to optimality.

FALSE The problem was not solved to optimality.

▷ **sym_is_proven_primal_infeasible**

```
int sym_is_proven_primal_infeasible(sym_environment *env)
```

Description:

This post-solution query routine is used to learn whether the problem was proven to be infeasible.

Arguments:

sym_environment *env IN Pointer to the SYMPHONY environment.

Return values:

TRUE The problem was proven to be infeasible.
FALSE The problem was not proven to be infeasible.

▷ **sym_is_iteration_limit_reached**

```
int sym_is_iteration_limit_reached(sym_environment *env)
```

Description:

This post-solution query routine is used to learn whether the iteration (node limit) was reached. It can also be used if “find_first_feasible” parameter was set to true before solving the problem.

Arguments:

`sym_environment *env` IN Pointer to the SYMPHONY environment.

Return values:

TRUE The iteration limit is reached.

FALSE The iteration limit is not reached.

▷ **sym_is_time_limit_reached**

```
int sym_is_time_limit_reached(sym_environment *env)
```

Description:

This post-solution query routine is used to learn whether the time limit was reached.

Arguments:

`sym_environment *env` IN Pointer to the SYMPHONY environment.

Return values:

TRUE Time limit was reached.

FALSE Time limit was not reached.

▷ **sym_is_target_gap_achieved**

```
int sym_is_target_gap_achieved(sym_environment *env)
```

Description:

This post-solution query routine is used to learn whether the target gap was reached.

Arguments:

sym_environment *env IN Pointer to the SYMPHONY environment.

Return values:

TRUE Target gap was reached.

FALSE Target gap was not reached.

▷ **sym_is_abandoned**

```
int sym_is_abandoned(sym_environment *env)
```

Description:

This post-solution query routine is used to learn whether the problem was abandoned for some reason.

Arguments:

sym_environment *env IN Pointer to the SYMPHONY environment.

Return values:

TRUE The problem was abandoned.

FALSE The problem was not abandoned.

6.1.4 Data Query Functions

▷ **sym_create_copy_mip_desc**

```
MIPdesc *sym_create_copy_mip_desc(sym_environment *env)
```

Description:

This routine is used to copy the problem description loaded to the environment.

Arguments:

`sym_environment *env` IN Pointer to the SYMPHONY environment.

Return values:

NULL An empty environment is passed in or there is no problem description loaded to the environment.

MIPdesc * Pointer to the copy of the problem description.

▷ **sym_get_num_cols**

```
int sym_get_num_cols(sym_environment *env, int *numcols)
```

Description:

This routine is used to get the number of the columns of the current problem.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int *numcols</code>	OUT	Pointer to an integer indicating the number of columns.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_num_rows**

```
int sym_get_num_cols(sym_environment *env, int *numrows)
```

Description:

This routine is used to get the number of the rows of the current problem.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int *numrows</code>	OUT	Pointer to an integer indicating the number of rows.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_num_elements**

```
int sym_get_num_elements(sym_environment *env, int *numelems)
```

Description:

This routine is used to get the number of non-zero entries of the constraint matrix of the current problem.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int *numelems</code>	OUT	Pointer to an integer indicating the number of non-zero elements.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_col_lower**

```
int sym_get_col_lower(sym_environment *env, double *collb)
```

Description:

This routine is used to get the lower bounds of the variables.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>double *collb</code>	OUT	Pointer to a double type array to be filled by the column lower bounds. Note that, the size of this array has to be at least the number of columns.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_col_upper**

```
int sym_get_col_upper(sym_environment *env, double *colub)
```

Description:

This routine is used to get the upper bounds of the variables.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>double *colub</code>	OUT	Pointer to a double type array to be filled by the column upper bounds. Note that, the size of this array has to be at least the number of columns.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_row_sense**

```
int sym_get_row_sense(sym_environment *env, char *rowсен)
```

Description:

This routine is used to get the row senses.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>char *rowсен</code>	OUT	Pointer to a char type array to be filled by the row senses. Note that, the size of this array has to be at least the number of rows.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_rhs**

```
int sym_get_rhs(sym_environment *env, double *rowrhs)
```

Description:

This routine is used to get the right hand side vector.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>double *rowrhs</code>	OUT	Pointer to a double type array to be filled by the right hand side vector. Note that, the size of this array has to be at least the number of rows.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_row_range**

```
int sym_get_row_range(sym_environment *env, double *rowrng)
```

Description:

This routine is used to get the row ranges.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>double *rowrng</code>	OUT	Pointer to a double type array to be filled by the row range values. Note that, the size of this array has to be at least the number of rows.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_row_lower**

```
int sym_get_row_lower(sym_environment *env, double *rowlb)
```

Description:

This routine is used to get the lower bounds of the rows.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>double *rowlb</code>	OUT	Pointer to a double type array to be filled by the row lower bounds. Note that, the size of this array has to be at least the number of rows.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_row_upper**

```
int sym_get_row_upper(sym_environment *env, double *rowub)
```

Description:

This routine is used to get the upper bounds of the rows.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>double *rowub</code>	OUT	Pointer to a double type array to be filled by the row upper bounds. Note that, the size of this array has to be at least the number of rows.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_matrix**

```
int sym_get_matrix(sym_environment *env, int *nz, int *matbeg, int *matind,  
double *matval)
```

Description: This routine is used to get the constraint matrix in a standard column-ordered format.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int *nz</code>	OUT	Pointer to integer indicating the non zero elements of the constraint matrix.
<code>int *matbeg</code>	OUT	Pointer to a double type array to be filled by the starting positions of each of column. Note that, the size of this array has to be at least the number of columns+1
<code>int *matind</code>	OUT	Pointer to a double type array to be filled by the row indices corresponding to each entry of <code>matval</code> . Note that, the size of this array has to be at least the number of nonzero elements of the constraint matrix.
<code>int *matval</code>	OUT	Pointer to a double type array of the values of nonzero entries of the constraint matrix in <i>column order</i> . Note that, the size of this array has to be at least the number of nonzero elements of the constraint matrix.

Return values: `FUNCTION_TERMINATED_ABNORMALLY` Function invoked unsuccessfully.
 `FUNCTION_TERMINATED_NORMALLY` Function invoked successfully.

▷ **sym_get_obj_coeff**

```
int sym_get_obj_coeff(sym_environment *env, double *obj)
```

Description:

This routine is used to get the objective vector.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>double *obj</code>	OUT	Pointer to a double type array to be filled by the objective vector. Note that, the size of this array has to be at least the number of columns.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_obj2_coeff**

```
int sym_get_obj2_coeff(sym_environment *env, double *obj2)
```

Description:

This routine is used to get the second objective vector if it exists. By default, it is set to the zero vector.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>double *obj2</code>	OUT	Pointer to a double type array to be filled by the second objective vector. Note that, the size of this array has to be at least the number of columns.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_obj_sense**

```
int sym_get_obj_sense(sym_environment *env, int *sense)
```

Description:

This routine is used to get the objective sense.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int *sense</code>	OUT	Pointer to an integer indicating the objective sense. In return, it will be 1 in case of minimization and -1 in case of maximization.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_is_continuous**

```
int sym_is_continuous(sym_environment *env, int index, int *value)
```

Description:

This routine is used to learn whether the queried variable is continuous.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int index</code>	IN	The index of the queried variable. Note that, it has to be at most the number of columns.
<code>int *value</code>	OUT	Pointer to a boolean indicating the variable status.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_is_binary**

```
int sym_is_binary(sym_environment *env, int index, int *value)
```

Description:

This routine is used to learn whether the queried variable is binary.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int index</code>	IN	The index of the queried variable. Note that, it has to be at most the number of columns.
<code>int *value</code>	OUT	Pointer to a boolean indicating the variable status.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_is_integer**

```
int sym_is_integer(sym_environment *env, int index, int *value)
```

Description:

This routine is used to ask whether the queried variable is integer.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int index</code>	IN	Index of the queried variable. Note that, it has to be at most the number of columns.
<code>int *value</code>	OUT	Pointer to a boolean indicating the variable status.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_infinity**

```
double sym_get_infinity()
```

Description:

This routine returns the infinity value of SYMPHONY.

Arguments:

Return values:

DOUBLE Infinity value of SYMPHONY

▷ **sym_get_col_solution**

```
int sym_get_col_solution(sym_environment *env, double *colsol)
```

Description:

This routine is used to get the post-solution column values.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>double *colsol</code>	OUT	Pointer to a double type array to be filled by the solution vector. Note that, the size of this array has to be at least the number of columns.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_sp_size**

```
int sym_get_sp_size(sym_environment *env, int *size)
```

Description:

This routine is used to get the size of the solution pool available after the solve. the solution pool contains additional (possibly sub optimal) solutions generated during the solution procedure.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int *size</code>	OUT	Pointer to the location in which the size of the solution pool will be stored after the function call.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_sp_solution**

```
int sym_get_sp_solution(sym_environment *env, int index, double *colsol,  
                        double *objval)
```

Description:

This routine is used to get the solution indexed by `index` from the solution pool.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int *size</code>	OUT	Index of the solution in the solution pool,
<code>double *colsol</code>	OUT	Pointer to a double type array to be filled by the solution vector. Note that, the size of this array has to be at least the number of columns.
<code>double *colsol</code>	OUT	Pointer to a double to the location that will contain the value of the solution after the function call.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_row_activity**

```
double *sym_get_row_activity(sym_environment *env, double *rowact)
```

Description:

This routine is used to get the row activities which are defined as the left hand side values, i.e., constraint matrix times the solution.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>double *rowact</code>	OUT	Pointer to a double type array to be filled by the row activity values. Note that, the size of this array has to be at least the number of rows.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_obj_val**

```
double *sym_get_obj_val(sym_environment *env, double *objval)
```

Description:

This routine is used to get the objective value after solving the problem.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>double *objval</code>	OUT	Pointer to a double indicating the post-solution objective value.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_primal_bound**

```
double *sym_get_primal_bound(sym_environment *env, double *ub)
```

Description:

This routine is used to get the a priori upper/lower bound for the problem.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>double *ub</code>	OUT	Pointer to a double indicating the upper (for minimization) or lower (for maximization) bound obtained through user defined primal heuristics.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

▷ **sym_get_iteration_count**

```
double *sym_get_iteration\_count(sym_environment *env, int *numnodes)
```

Description:

This routine is used to get the number of the analyzed nodes of the branching tree after solving the problem. It can also be used to query the status of a loaded warm start.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int *numnodes</code>	OUT	Pointer to an integer indicating the number of nodes analyzed so far.

Return values:

<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.
<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.

6.1.5 Data Modification Functions

▷ **sym_set_obj_coeff**

```
double *sym_set_obj_coeff(sym_environment *env, int index, double value)
```

Description:

This routine is used to set an objective coefficient.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int index</code>	IN	Index of the objective coefficient to be modified. Note that, it has to be at most the number of columns.
<code>double value</code>	IN	New objective value of the corresponding column.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_set_obj2_coeff**

```
double *sym_set_obj2_coeff(sym_environment *env, int index, double value)
```

Description:

This routine is used to set a coefficient of the second objective function of the corresponding bicriteria problem.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int index</code>	IN	Index of the objective coefficient to be modified. Note that, it has to be at most the number of columns.
<code>double value</code>	IN	New value of the objective coefficient to be modified.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_set_col_lower**

```
double *sym_set_col_lower(sym_environment *env, int index, double value)
```

Description:

This routine is used to set the lower bound of a variable.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int index</code>	IN	Index of the variable. Note that, it has to be at most the number of columns.
<code>double value</code>	IN	New lower bound of the variable.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_set_col_upper**

```
double *sym_set_col_upper(sym_environment *env, int index, double value)
```

Description:

This routine is used to set the upper bound of a variable.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int index</code>	IN	Index of the variable. Note that, it has to be at most the number of columns.
<code>double value</code>	IN	New upper bound of the variable.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_set_row_lower**

```
double *sym_set_row_lower(sym_environment *env, int index, double value)
```

Description:

This routine is used to set the lower bound of a row.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int index</code>	IN	Index of the row. Note that, it has to be at most the number of rows.
<code>double value</code>	IN	New lower bound of the row.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_set_row_upper**

```
double *sym_set_row_upper(sym_environment *env, int index, double value)
```

Description:

This routine is used to set the upper bound of a row.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int index</code>	IN	Index of the row. Note that, it has to be at most the number of rows.
<code>double value</code>	IN	New upper bound of the row.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_set_row_type**

```
int sym_set_row_type(sym_environment *env, int index, char rowsense,  
                    double rowrhs, double rowrng)
```

Description:

This routine is used to set the characteristics of a row.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int index</code>	IN	Index of the row. Note that, it has to be at most the number of rows.
<code>char rowsense</code>	IN	New sense of the row.
<code>double rowrhs</code>	IN	New value of the right hand side of the row.
<code>double rowrng</code>	IN	New value of the row range.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_set_obj_sense**

```
int sym_set_obj_sense(sym_environment *env, int sense)
```

Description:

This routine is used to set the objective sense. By default, SYMPHONY will solve a minimization problem.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int sense</code>	IN	New sense of the objective function. It can be 1 and -1 for minimization and maximization. Otherwise, SYMPHONY will assume the objective sense to be minimization.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully

▷ **sym_set_col_solution**

```
int sym_set_col_solution(sym_environment *env, double *colsol)
```

Description:

This routine is used to set the current solution if a known one exists. Note that setting the column solution will not affect or help the treemanager's processes other than setting the best feasible solution and the corresponding upper bound.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>double *colsol</code>	IN	Pointer to a double type array of the known column values. Note that, if the given solution is not feasible or if a better solution was found/loaded before, SYMPHONY will refuse to set the column solution and will leave this function without success.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_set_primal_bound**

```
int sym_set_primal_bound(sym_environment *env, double bound)
```

Description:

This routine is used to set a priori upper/lower bound to the problem.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>double double</code>	IN	The value of the priori upper (for minimization) or lower (for maximization) bound.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_set_continuous**

```
int sym_set_continuous(sym_environment *env, int index)
```

Description:

This routine is used to set the type of a variable to be continuous.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int index</code>	IN	The index of the variable to be modified. Note that, it has to be at most the number of columns.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_set_integer**

```
int sym_set_integer(sym_environment *env, int index)
```

Description:

This routine is used to set the type of a variable to be integer.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int index</code>	IN	The index of the variable to be modified. Note that, it has to be at most the number of columns.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_set_col_names**

```
int sym_set_col_names(sym_environment * env, char **colname)
```

Description:

This routine is used to set the column names.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>char **colname</code>	IN	Pointer to a string array including the column names. Note that, the size of this array has to be at least the number of columns.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_add_col**

```
int sym_add_col(sym_environment *env, int numelems, int *indices,  
                double *elements, double collb, double colub,  
                double obj, char *name)
```

Description:

This routine is used to add a new column to the original problem description.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int numelems</code>	IN	An integer indicating the non zero elements of the column.
<code>int *indices</code>	IN	Pointer to an integer type array indicating the row indices of the non zero elements of the column and having a size of at least <code>numelems</code> .
<code>double *elements</code>	IN	Pointer to a double type array indicating the values of the non zero elements of the column and having a size of at least <code>numelems</code> .
<code>double collb</code>	IN	A double indicating the lower bound of the column.
<code>double colub</code>	IN	A double indicating the upper bound of the column.
<code>double obj</code>	IN	A double indicating the objective coefficient value of the column.
<code>char *name</code>	IN	Pointer to a string of the name of the column.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_add_row**

```
int sym_add_row(sym_environment *env, int numelems, int *indices,  
                double *elements, char rowsen, double rowrhs,  
                double rowrng)
```

Description:

This routine is used to add a new row to the original constraint matrix.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int numelems</code>	IN	An integer indicating the non zero elements of the row.
<code>int *indices</code>	IN	Pointer to an integer type array indicating the column indices of the non zero elements of the row and having a size of at least <code>numelems</code> .
<code>double *elements</code>	IN	Pointer to a double type array indicating the values of the non zero elements of the row and having a size of at least <code>numelems</code> .
<code>char rowsen</code>	IN	A character indicating the sense of the row.
<code>double rowrhs</code>	IN	A double indicating the right hand side of the row.
<code>double rowrng</code>	IN	A double indicating the range value of the row.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_delete_cols**

```
int sym_delete_cols(sym_environment *env, int num, int * indices)
```

Description:

This routine is used to delete columns from the original problem description.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int num</code>	IN	An integer indicating the number of columns to be deleted.
<code>int *indices</code>	IN	Pointer to an integer type array indicating the indices of the columns to be deleted and having a size of at least <code>num</code> .

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully or one of the <code>indices</code> is out of the range of <code>[0, number of variables-1]</code>

▷ **sym_delete_rows**

```
int sym_delete_rows(sym_environment *env, int num, int * indices)
```

Description:

This routine is used to delete rows from the original constraint matrix.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>int num</code>	IN	An integer indicating the number of rows to be deleted.
<code>int *indices</code>	IN	An array indicating the indices of the rows to be deleted and having a size of at least <code>num</code> .

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully or one of the <code>indices</code> is out of the range of <code>[0, number of variables-1]</code>

6.1.6 Warm Starting Functions

▷ **sym_write_warm_start_desc**

```
int sym_write_warm_start_desc(warm_start_desc *ws, char *file)
```

Description:

This routine is used to write the given warm start structure to a file.

Arguments:

<code>warm_start_desc *ws</code>	IN	Pointer to the warm start description to be written.
<code>char *file</code>	IN	The name of the file the warm start is desired to be written to.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_read_warm_start**

```
int sym_read_warm_start(char *file, warm_start_desc *ws)
```

Description:

This routine is used to read in a warm start structure from a file.

Arguments:

<code>char *file</code>	IN	The name of the file the warm start is desired to be read from.
<code>warm_start_desc *ws</code>	OUT	Pointer to a warm start object to be read from the file.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_delete_warm_start**

```
void sym_delete_warm_start(warm_start_desc *ws)
```

Description:

This routine is used to free a warm start structure and return the allocated memory.

Arguments:

`warm_start_desc *ws` IN Pointer to the warm start description to be deleted.

▷ **sym_get_warm_start**

```
warm_start_desc *sym_get_warm_start(sym_environment *env, int copy_warm_start,  
warm_start_desc **ws)
```

Description:

This routine is used to get the warm start description loaded to the environment.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int copy_warm_start</code>	IN	A boolean indicating whether the warm start of the environment is desired to be copied or overtaken.
<code>warm_start_desc **ws</code>	OUT	Pointer to a pointer to be directed to a copy or the itself of the currently loaded warm start.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_set_warm_start**

```
int sym_set_warm_start(sym_environment *env, warm_start_desc *ws)
```

Description:

This routine is used to load a warm start structure to the environment.

Arguments:

<code>sym_environment *env</code>	INOUT	Pointer to the SYMPHONY environment.
<code>warm_start_desc *ws</code>	IN	Pointer to the warm start structure to be loaded to the environment.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ **sym_create_copy_warm_start**

```
warm_start_desc *sym_create_copy_warm_start(warm_start_desc *ws)
```

Description:

This routine is used to copy the given warm start structure.

Arguments:

`warm_start_desc *ws` INOUT Pointer to the warm start structure to be copied.

Return values:

`tt NULL` An empty warm start description is passed in.

`WARM_START_DESC *` Pointer to the copy of the warm start structure.

6.1.7 Sensitivity Analysis Functions

▷ **sym_get_lb_for_new_rhs**

```
int sym_get_lb_for_new_rhs(sym_environment *env, int cnt, int *new_rhs_ind,  
                           double *new_rhs_val, double *lb_for_new_rhs)
```

Description:

This routine is used for a basic sensitivity analysis of the right hand side case. It returns a lower bound for the problem with a modified right hand side using the information gathered from the branching tree of the original solved problem. Note that, in order to use this feature, the `sensitivity_analysis` parameter needs to be set before solving the original problem.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int cnt</code>	IN	The number of the non zero elements in the new right hand side vector.
<code>int *new_rhs_ind</code>	IN	Array of the column indices of these non zero elements.
<code>double *new_rhs_val</code>	IN	Array of the values of these non zero elements.
<code>double *lb_for_new_rhs</code>	OUT	Pointer to a double indicating the lower bound obtained for the new problem.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ `sym_get_ub_for_new_rhs`

```
int sym_get_ub_for_new_rhs(sym_environment *env, int cnt, int *new_rhs_ind,  
                           double *new_rhs_val, double *ub_for_new_rhs)
```

Description:

This routine is used for a basic sensitivity analysis of the right hand side case. It returns a quick upper bound for the problem with a modified right hand side using the information gathered from the branching tree of the original solved problem. Note that, in order to use this feature, the `sensitivity_analysis` parameter needs to be set before solving the original problem.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int cnt</code>	IN	The number of the non zero elements in the new right hand side vector.
<code>int *new_rhs_ind</code>	IN	Array of the column indices of these non zero elements.
<code>double *new_rhs_val</code>	IN	Array of the values of these non zero elements.
<code>double *ub_for_new_rhs</code>	OUT	Pointer to a double indicating the lower bound obtained for the new problem. This value will be set to <code>SYM_INFINITY</code> if an upper bound can not be found.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ `sym_get_lb_for_new_obj`

```
int sym_get_lb_for_new_rhs(sym_environment *env, int cnt, int *new_obj_ind,  
                           double *new_obj_val, double *lb_for_new_obj)
```

Description:

This routine is used for a basic sensitivity analysis of the objective function case. It returns a quick lower bound for the problem with a modified objective vector using the information gathered from the branching tree of the original solved problem. Note that, in order to use this feature, the `sensitivity_analysis` parameter needs to be set before solving the original problem.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int cnt</code>	IN	The number of the non zero elements in the new objective coefficients.
<code>int *new_obj_ind</code>	IN	Array of the column indices of these non zero elements.
<code>double *new_obj_val</code>	IN	Array of the values of these non zero elements.
<code>double *lb_for_new_obj</code>	OUT	Pointer to a double indicating the lower bound obtained for the new problem.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

▷ `sym_get_ub_for_new_obj`

```
int sym_get_ub_for_new_rhs(sym_environment *env, int cnt, int *new_obj_ind,  
                           double *new_obj_val, double *ub_for_new_obj)
```

Description:

This routine is used for a basic sensitivity analysis of the objective function case. It returns a quick lower bound for the problem with a modified objective vector using the information gathered from the branching tree of the original solved problem. Note that, in order to use this feature, the `sensitivity_analysis` parameter needs to be set before solving the original problem.

Arguments:

<code>sym_environment *env</code>	IN	Pointer to the SYMPHONY environment.
<code>int cnt</code>	IN	The number of the non zero elements in the new objective coefficients.
<code>int *new_obj_ind</code>	IN	Array of the column indices of these non zero elements.
<code>double *new_obj_val</code>	IN	Array of the values of these non zero elements.
<code>double *ub_for_new_obj</code>	OUT	Pointer to a double indicating the upper bound obtained for the new problem. This value will be set to <code>SYM_INFINITY</code> if an upper bound can not be found.

Return values:

<code>FUNCTION_TERMINATED_NORMALLY</code>	Function invoked successfully.
<code>FUNCTION_TERMINATED_ABNORMALLY</code>	Function invoked unsuccessfully.

6.2 Callable Library C++ API

SYMPHONY's C++ interface is derived from COIN-OR's Open Solver Interface (OSI). The OSI methods are implemented simply as wrapped calls to the SYMPHONY C callable library just described. For instance, when an instance of the OSI interface class is constructed, a call is made to `sym_open_environment()` and a pointer to the environment is stored in the class and when the OSI object is destroyed, `sym_close_environment` is called to destroy the environment object. Most subsequent calls within the class can then be made without any arguments. To fully support SYMPHONY's capabilities, we have extended the OSI interface to include some other methods not in the base class. For example, we added calls equivalent to our `sym_parse_command_line()` and `sym_find_initial_bounds()`. Additionally, SYMPHONY has a warm start class derived from the CoinWarmStart base class to support the new functionalities of the MILP warm starting such as `sym_get_warm_start` and `sym_set_warm_start`. They are also implemented as wrapped calls to the C interface library.

In order to have the whole list of the methods and information regarding their usage, see the OSI SYMPHONY interface and SYMPHONY warm start header files (`OsiSymSolverInterface.hpp` and `SymWarmStart.hpp`). Here, we will give the table of the C library equivalent calls of the C++ interface routines with brief descriptions:

C++ Interface	C Interface	Description
OsiSymSolverInterface	sym_open_environment	create a new environment.
loadProblem	sym_load_problem	load the problem read trough an MPS or GMPL file
branchAndBound	sym_solve/sym_warm_solve	solve the MILP problem from scratch or from a warm start if loaded.
resolve	sym_warm_solve	re-solve the MILP problem after some modifications.
initialSolve	sym_solve	solve the MILP problem from scratch.
multiCriteriaBranchAndBound	sym_mc_solve	solve the multi criteria problem.
setInitialData	sym_set_defaults	set the parameters to their defaults.
parseCommandLine	sym_parse_command_line	read the command line arguments.
findInitialBounds	sym_find_initial_bounds	find the initial bounds via the user defined heuristics.
createPermanentCutPools	sym_create_permanent_cut_pools	save the global cuts.
loadProblem	sym_explicit_load_problem	load the problem through a set of arrays.
getWarmStart	sym_get_warm_start	get the warm start description.
setWarmStart	sym_set_warm_start	set the warm start description.
getLbForNewRhs	sym_get_lb_for_new_rhs	find a lower bound to the new rhs problem using the post solution info.
getUbForNewRhs	sym_get_lb_for_new_rhs	find an upper bound to the new rhs problem. using the post solution info.
getLbForNewObj	sym_get_lb_for_new_rhs	find a lower bound to the new obj problem. using the post solution info.
getUbForNewObj	sym_get_lb_for_new_rhs	find an upper bound to the new obj problem. using the post solution info.
reset	sym_close_environment	return the allocated memory.
setIntParam	sym_set_int_param	set the integer type OSI parameter.
setSymParam(int)	sym_set_int_param	set the integer type SYMPHONY parameter.
setDblParam	sym_set_dbl_param	set the double type OSI parameter.
setSymParam(double)	sym_set_dbl_param	set the double type SYMPHONY parameter.
setStrParam	sym_set_str_param	set the string type OSI parameter.
setSymParam(string)	sym_set_str_param	set the string type SYMPHONY parameter.
getIntParam	sym_get_int_param	get the value of the integer type OSI parameter.
getSymParam(int &)	sym_get_int_param	get the value of the integer type SYMPHONY parameter.
getDblParam	sym_get_dbl_param	get the value of the double type OSI parameter.
getSymParam(double &)	sym_get_dbl_param	get the value of the double type SYMPHONY parameter.
getStrParam	sym_get_str_param	get the value of the string type OSI parameter.
getSymParam(string &)	sym_get_str_param	get the value of the string type SYMPHONY parameter.
isProvenOptimal	sym_is_proven_optimal	query the problem status.
isProvenPrimalInfeasible	sym_is_proven_primal_infeasible	query the problem status.
isPrimalObjectiveLimitReached	sym_is_target_gap_achieved	query the problem status.
isIterationLimitReached	sym_is_iteration_limit_reached	query the problem status.
isTimeLimitReached	sym_is_time_limit_reached	query the problem status.
isTargetGapReached	sym_is_target_gap_achieved	query the problem status.
getNumCols	sym_get_num_cols	get the number of columns.
getNumRows	sym_get_num_rows	get the number of rows.
getNumElements	sym_get_num_elements	get the number of nonzero elements.
getColLower	sym_get_col_lower	get the column lower bounds.
getColUpper	sym_get_col_upper	get the column upper bounds.
getRowSense	sym_get_row_sense	get the row senses.
getRightHandSide	sym_get_rhs	get the rhs values.
getRowRange	sym_get_row_range	get the row range values.
getRowLower	sym_get_row_lower	get the row lower bounds.
getRowUpper	sym_get_row_upper	get the row upper bounds.
getObjCoefficients	sym_get_obj_coeff	get the objective function vector.

C++ Interface	C Interface	Description
getObjSense	sym_get_obj_sense	get the objective sense.
isContinuous	sym_is_continuous	query the variable type.
isBinary	sym_is_binary	query the variable type.
isInteger	sym_is_integer	query the variable type.
isIntegerNonBinary	-	query the variable type.
isFreeBinary	sym_is_binary	query the variable type.
getMatrixByRow	-	get the constraint matrix by row oriented.
getMatrixByCol	-	get the constraint matrix by column oriented.
getInfinity	-	get the infinity definition of SYMPHONY.
getColSolution	sym_get_col_solution	get the current best column solution.
getRowActivity	sym_get_row_activity	get the current row activity.
getObjValue	sym_get_obj_val	get the current best objective value.
getPrimalBound	sym_get_primal_bound	get the primal upper bound.
getIterationCount	sym_get_iteration_count	get the number of the analyzed tree nodes.
setObjCoeff	sym_set_obj_coeff	set the objective function vector.
setObj2Coeff	sym_set_obj2_coeff	set the second objective function vector.
setColLower	sym_set_col_lower	set the column lower bounds.
setColUpper	sym_set_col_upper	set the column upper bounds.
setRowLower	sym_set_row_lower	set the row lower bounds.
setRowUpper	sym_set_row_upper	set the row upper bounds.
setRowType	sym_set_row_type	set the row characteristics.
setObjSense	sym_set_obj_sense	set the objective sense.
setColSolution	sym_set_col_solution	set the current solution.
setContinuous	sym_set_continuous	set the variable type.
setInteger	sym_set_integer	set the variable type.
setColName	sym_set_col_names	set the column names.
addCol	sym_add_col	add columns to the constraint matrix.
addRow	sym_add_row	add rows to the constraint matrix.
deleteCols	sym_delete_cols	delete some columns from the constraint matrix.
deleteRows	sym_delete_rows	delete some rows from the constraint matrix.
writeMps	-	write the current problem in MPS format.
applyRowCut	-	add some row cuts.
applyColCut	-	add some column cuts.
SymWarmStart(warm_start_desc *)	sym_create_copy_warm_start	create a SYMPHONY warm start by copying the given one.
SymWarmStart(char *)	sym_read_warm_start	create a SYMPHONY warm start reading from file.
getCopyOfWarmStartDesc	sym_create_copy_warm_start	get the copy of the warm start structure.
writeToFile	sym_write_warm_start_desc	write the loaded warm start to a file.

6.3 User Callback API

6.3.1 Master module callbacks

▷ **user_usage**

```
void user_usage()
```

Description:

SYMPHONY's command-line switches are all lower case letters. The user can use any upper case letter (except 'H' and as specified below) for command line switches to control user-defined parameter settings without the use of a parameter file. The function **user_usage()** can optionally print out usage information for the user-defined command line switches. The command line switch -H automatically calls the user's usage sub-routine. The switch -h prints SYMPHONY's own usage information. In its default configuration, the command-line switch -F is used to specify the file in which the instance data is contained (either an MPS file or an GMPL/AMPL file). The -D switch is used to specify the data file if an GMPL/AMPL file is being read in (see the README file). The -L switch is used to specify the data file if a file in LP format is being read in

▷ **user_initialize**

```
int user_initialize(void **user)
```

Description:

The user allocates space for and initializes the user-defined data structures for the master module.

Arguments:

<code>void **user</code>	OUT	Pointer to the user-defined data structure.
--------------------------	-----	---

Return values:

<code>USER_ERROR</code>	Error. SYMPHONY exits.
<code>USER_SUCCESS</code>	Initialization is done.
<code>USER_DEFAULT</code>	There is no user-defined data structure (this can be the case if the default parser is being used to read in either an MPS or GMPL/AMPL file).

▷ **user_readparams**

```
int user_readparams(void *user, char *filename, int argc, char **argv)
```

Description:

The user can optionally read in parameter settings from the file named `filename`, specified on the command line using the `-f` switch. The parameter file `filename` can contain both SYMPHONY's built-in parameters and user-defined parameters. If desired, the user can open this file for reading, scan the file for lines that contain user parameter settings and then read the parameter settings. A shell for doing this is set up in the file `SYMPHONY-5.0/USER/user_master.c`. Also, see the file `Master/master_io.c` to see how SYMPHONY does this.

The user can also parse the command line arguments for user settings. A shell for doing this is also set up in the file `SYMPHONY-5.0/USER/user_master.c`. Upper case letters are reserved for user-defined command line switches. The switch `-H` is reserved for help and calls the user's usage subroutine (see `user_usage()` 6.3.1). If the user returns 'USER_DEFAULT', then SYMPHONY will look for the command-line switches `-F` to specify the file name for reading in the model from either an MPS or a GMPL/AMPL file or `-L` to specify the file name for reading in the model from an LP format file. The `-D` command-line switch is used to specify an additional data file for GMPL/AMPL models. If the `-D` option is not present, SYMPHONY assumes the file is an MPS file.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>char *filename</code>	IN	The name of the parameter file.

Return values:

<code>USER_ERROR</code>	Error. SYMPHONY stops.
<code>USER_SUCCESS</code>	User parameters were read successfully.
<code>USER_DEFAULT</code>	SYMPHONY will read in the problem instance from either an MPS, LP, or GMPL/AMPL file. The command-line switches <code>-F</code> , <code>-L</code> , and <code>-D</code> (as appropriate) will be used to specify the model file.

▷ **user_io**

```
int user_io(void *user)
```

Description:

Here, the user can read in an instance in a custom format and set up appropriate data structures. If the user wants to use the default parsers to read either an MPS file or a GMPL/AMPL file, then the return value `USER_DEFAULT` should be specified (see `user_readparams()` 6.3.1 for the command-line switches to use to specify this behavior).

Arguments:

`void *user` IN Pointer to the user-defined data structure.

Return values:

<code>USER_ERROR</code>	Error. SYMPHONY stops.
<code>USER_SUCCESS</code>	User I/O was completed successfully.
<code>USER_DEFAULT</code>	Input will be read in from an MPS or GMPL/AMPL file.

▷ **user_init_draw_graph**

```
int user_init_draw_graph(void *user, int dg_id)
```

Description:

This function is invoked only if the `do_draw_graph` parameter is set. The user can initialize the graph drawing module by sending some initial information (e.g., the location of the nodes of a graph, like in the TSP.)

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>int dg_id</code>	IN	The process id of the graph drawing module.

Return values:

<code>USER_ERROR</code>	Error. SYMPHONY stops.
<code>USER_SUCCESS</code>	The user completed initialization successfully.
<code>USER_DEFAULT</code>	No action.

▷ **user_start_heurs**

```
int user_start_heurs(void *user, double *ub, double *ub_estimate)
```

Description:

The user invokes heuristics and generates the initial global upper bound and also perhaps an upper bound estimate. This is the last place where the user can do things before the branch and cut algorithm starts. She might do some preprocessing, in addition to generating the upper bound.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>double *ub</code>	OUT	Pointer to the global upper bound. Initially, the upper bound is set to either <code>-MAXDOUBLE</code> or the bound read in from the parameter file, and should be changed by the user only if a better valid upper bound is found.
<code>double *ub_estimate</code>	OUT	Pointer to an estimate of the global upper bound. This is useful if the <code>BEST_ESTIMATE</code> diving strategy is used (see the treeman-ager parameter <code>diving_strategy</code> (Section 6.4.4))

Return values:

<code>USER_ERROR</code>	Error. This error is probably not fatal.
<code>USER_SUCCESS</code>	User executed function successfully.
<code>USER_DEFAULT</code>	No action (someday, there may be a default MIP heuristic here).

▷ `user_initialize_root_node`

```
int user_initialize_root_node(void *user, int *basevarnum, int **basevars,
                             int *basecutnum, int *extravarnum, int **extravars,
                             char *obj_sense, double *obj_offset,
                             char ***colnames, int *colgen_strat)
```

Description:

In this function, the user must specify the list of indices for the base and extra variables. The option to specify a variable as base is provided simply for efficiency reasons. If there is no reasonable way to come up with a set of base variables, then all variables should be specified as extra (see Section 4.3.2.1 for a discussion of base and extra variables). If the function returns `USER_DEFAULT` and sets `extravarnum`, then SYMPHONY will put all variables indexed from 0 to `extravarnum` in the set of extra variables by default. If an MPS or GMPL/AMPL file was read in using SYMPHONY's built-in parser, i.e., the default behavior of `user_io()` 6.3.1 was not modified, then `extravarnum` need not be set.

In this function, the user may also specify column names for display purposes. If the `colnames` array is allocated, then SYMPHONY will use for displaying solutions. If the data was read in from either an MPS or GMPL/AMPL file, then the column names will be set automatically.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>int *basevarnum</code>	OUT	Pointer to the number of base variables.
<code>int **basevars</code>	OUT	Pointer to an array containing a list of user indices of the base variables to be active in the root.
<code>int *basecutnum</code>	OUT	The number of base constraints.
<code>int *extravarnum</code>	OUT	Pointer to the number of extra active variables in the root.
<code>int **extravars</code>	OUT	Pointer to an array containing a list of user indices of the extra variables to be active in the root.
<code>char *obj_sense</code>	INOUT	Whether to negate the objective function value when printing the solution, set to either <code>MAXIMIZE</code> or <code>MINIMIZE</code> . Note that SYMPHONY always minimizes— this only effects the printing of the solution . The default is <code>MINIMIZE</code> .
<code>double *obj_offset</code>	INOUT	A specified constant to be added to the objective function value when printing out the solution.
<code>int ***colnames</code>	OUT	Pointer to an array containing a list of column names to be used for display purposes.
<code>int *colgen_strat</code>	INOUT	The default strategy or one that has been read in from the parameter file is passed in, but the user is free to change it. See <code>colgen_strat</code> in the description of parameters for details on how to set it.

Return values:

USER_ERROR	Error. SYMPHONY stops.
USER_SUCCESS	The required data are filled in.
USER_DEFAULT	All variables indexed 0 to extravarnum are put in the extra set (The user must set extravarnum unless an MPS or GMPL/AMPL file was read in by SYMPHONY).

Post-processing:

The array of base and extra indices are sorted.

▷ `user_receive_feasible_solution`

```
int user_receive_feasible_solution(void *user, int msgtag, double cost,  
                                  int numvars, int *indices, double *values)
```

Description:

This function is only used for parallel execution. Feasible solutions can be sent and/or stored in a user-defined packed form if desired. For instance, the TSP, a tour can be specified simply as a permutation, rather than as a list of variable indices. In the LP module, a feasible solution is packed either by the user or by a default packing routine. If the default packing routine was used, the `msgtag` will be `FEASIBLE_SOLUTION_NONZEROS`. In this case, `cost`, `numvars`, `indices` and `values` will contain the solution value, the number of nonzeros in the feasible solution, and their user indices and values. The user has only to interpret and store the solution. Otherwise, when `msgtag` is `FEASIBLE_SOLUTION_USER`, SYMPHONY will send and receive the solution value only and the user has to unpack exactly what she has packed in the LP module. In this case the contents of the last three arguments are undefined.

In most cases, SYMPHONY's default routines for sending and receiving feasible solutions, as well as displaying them, will suffice. These routines simply display all nonzeros by either index or name, depending on whether the user set the column names. See `user_receive_lp_data()` in Section 6.3.2.2 for more discussion.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>int msgtag</code>	IN	<code>FEASIBLE_SOLUTION_NONZEROS</code> or <code>FEASIBLE_SOLUTION_USER</code>
<code>double cost</code>	IN	The cost of the feasible solution.
<code>int numvars</code>	IN	The number of variables whose user indices and values were sent (length of <code>indices</code> and <code>values</code>).
<code>int *indices</code>	IN	The user indices of the nonzero variables.
<code>double *values</code>	IN	The corresponding values.

Return values:

<code>USER_ERROR</code>	Ignored. This is probably not a fatal error.
<code>USER_SUCCESS</code>	The solution has been unpacked and stored.
<code>USER_DEFAULT</code>	Store the nonzeros in the solutions for later display.

▷ `user_send_lp_data`

```
int user_send_lp_data(void *user, void **user_lp)
```

Description:

If the user wishes to employ parallelism, she has to send all problem-specific data that will be needed to implement user functions in the LP module in order to set up the initial LP relaxation and perform later computations. This could include instance data, as well as user parameter settings (see Section 5.5.1 for a discussion of this). This is one of the few places where the user may need to worry about the configuration of the modules. If either the tree manager or the LP are running as a separate process (either `COMPILE_IN_LP` or `COMPILE_IN_TM` are `FALSE` in the make file), then the data will be sent and received through message-passing. See `user_receive_lp_data()` in Section 6.3.2.2 for more discussion. Otherwise, it can be copied through shared memory. The easiest solution, which is set up by default is to simply copy over a pointer to a single user data structure where instance data is stored. The code for the two cases is put in the same source file by use of `#ifdef` statements. See the comments in the code stub for this function for more details.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>void **user_lp</code>	OUT	Pointer to the user-defined data structure for the LP module.

Return values:

<code>USER_ERROR</code>	Error. SYMPHONY stops.
<code>USER_SUCCESS</code>	Packing is done.
<code>USER_DEFAULT</code>	User has no data to send. This would be used when SYMPHONY has read in an MPS or GMPL/AMPL model file.

▷ **user_send_cg_data**

```
int user_pack_cg_data(void *user, void **user_cg)
```

Description:

If the user wishes to employ parallelism and wants a separate cut generator module, this function can be used to send all problem-specific data that will be needed by the cut generator module to perform separation. This could include instance data, as well as user parameter settings (see Section 5.5.1 for a discussion of this). This is one of the few places where the user may need to worry about the configuration of the modules. If either the tree manager or the LP are running as a separate process (either `COMPILE_IN_LP` or `COMPILE_IN_TM` are `FALSE` in the `make` file), then the data will be sent and received through message-passing. See `user_receive_cg_data()` in Section 6.3.3 for more discussion. Otherwise, it can be copied through shared memory. The easiest solution, which is set up by default is to simply copy over a pointer to a single user data structure where instance data is stored. The code for the two cases is put in the same source file by use of `#ifdef` statements. See the comments in the code stub for this function for more details.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>void **user_cg</code>	OUT	Pointer to the user-defined data structure for the cut generator module.

Return values:

<code>USER_ERROR</code>	Error. SYMPHONY stops.
<code>USER_SUCCESS</code>	Packing is done.
<code>USER_DEFAULT</code>	No data to send to the cut generator (no separation performed).

▷ `user_send_cp_data`

```
int user_pack_cp_data(void *user, void **user_cp)
```

Description:

If the user wishes to employ parallelism and wants to use the cut pool to store user-defined cuts, this function can be used to send all problem-specific data that will be needed by the cut pool module. This could include instance data, as well as user parameter settings (see Section 5.5.1 for a discussion of this). This is one of the few places where the user may need to worry about the configuration of the modules. If either the tree manager or the LP are running as a separate process (either `COMPILE_IN_LP` or `COMPILE_IN_TM` are `FALSE` in the `make` file), then the data will be sent and received through message-passing. See `user_receive_cp_data()` in Section 6.3.4 for more discussion. Otherwise, it can be copied through shared memory. The easiest solution, which is set up by default is to simply copy over a pointer to a single user data structure where instance data is stored. The code for the two cases is put in the same source file by use of `#ifdef` statements. See the comments in the code stub for this function for more details.

Note that there is support for cuts generated and stored as explicit matrix rows. The cut pool module is already configured to deal with such cuts, so no user implementation is required. Only the use of user-defined cuts requires customization of the Cut Pool module.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>void **user_cp</code>	OUT	Pointer to the user-defined data structure for the cut pool module.

Return values:

<code>USER_ERROR</code>	Error. SYMPHONY stops.
<code>USER_SUCCESS</code>	Packing is done.
<code>USER_DEFAULT</code>	No data to send to the cut pool (no user-defined cut classes or cut pool not used).

▷ **user_display_solution**

```
int user_display_solution(void *user, double lpetol, int varnum, int *indices,  
                          double *values, double objval)
```

Description:

This function is invoked when a best solution found so far is to be displayed (after heuristics, after the end of the first phase, or the end of the whole algorithm). This can be done using either a text-based format or using the **drawgraph** module. By default, SYMPHONY displays the indices (or column names, if specified) and values for each nonzero variable in the solution. The user may wish to display a custom version of the solution by interpreting the variables.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure. For sequential computation, a pointer to the user's LP data structure is passed in. For parallel computation, a pointer to the user's Master data structure is passed in.
<code>double lpetol</code>	IN	The LP zero tolerance used.
<code>int varnum</code>	IN	The number of nonzeros in the solution.
<code>int *indices</code>	IN	The indices of the nonzeros.
<code>double *values</code>	IN	The values of the nonzeros.
<code>double objval</code>	IN	The objective function value of the solution.

Return values:

<code>USER_ERROR</code>	Ignored.
<code>USER_SUCCESS</code>	User displayed the solution. SYMPHONY should do nothing.
<code>USER_DEFAULT</code>	SYMPHONY should display the solution in default format.

Post-processing:

If requested, SYMPHONY will display a best solution found so far in the default format.

▷ **user_send_feas_sol**

```
int user_send_feas_sol(void *user, int *feas_sol_size, int **feas_sol)
```

Description:

This function is useful for debugging purposes. It passes a known feasible solution to the tree manager. The tree manager then tracks which current subproblem admits this feasible solution and notifies the user when it gets pruned. It is useful for finding out why a known optimal solution never gets discovered. Usually, this is due to either an invalid cut of an invalid branching. Note that this feature only works when branching on binary variables. See Section 5.6.3 for more on how to use this feature.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>int *feas_sol_size</code>	INOUT	Pointer to size of the feasible solution passed by the user.
<code>int **feas_sol</code>	INOUT	Pointer to the array of user indices containing the feasible solution. This array is simply copied by the tree manager and must be freed by the user.

Return values:

Arguments:	<code>USER_ERROR</code>	Solution tracing is not enabled.
	<code>USER_SUCCESS</code>	Tracing of the given solution is enabled.
	<code>USER_DEFAULT</code>	No feasible solution given.

▷ **user_process_own_messages**

```
int user_process_own_messages(void *user, int msgtag)
```

Description:

The user must receive any message he sends to the master module (independently of SYMPHONY's own messages). An example for such a message is sending feasible solutions from separate heuristics processes fired up in `user_start_heurs()`.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>int msgtag</code>	IN	The message tag of the message.

Return values:

<code>USER_ERROR</code>	Ignored.
<code>USER_SUCCESS</code>	Message is processed.
<code>USER_DEFAULT</code>	No user message types defined.

▷ **user_free_master**

```
int user_free_master(void **user)
```

Description:

The user frees all the data structures within ***user**, and also free ***user** itself. This can be done using the built-in macro **FREE** that checks the existence of a pointer before freeing it.

Arguments:

void **user	INOUT	Pointer to the user-defined data structure (should be NULL on return).
--------------------	-------	---

Return values:

USER_ERROR	Ignored. This is probably not a fatal error.
USER_SUCCESS	Everything was freed successfully.
USER_DEFAULT	There is no user memory to free.

6.3.2 LP module callbacks

6.3.2.1 Data Structures

We first describe a few structures that are used to pass data into and out of the user functions of the LP module.

▷ **MIPdesc**

One of the few internally defined data structures that the user has to deal with frequently is the **MIPdesc** data structure, which holds the data needed to describe a MILP. This data structure is used by SYMPHONY for two purposes. First, it is used to store the description of a generic MILP that is either read in from an MPS or AMPL file. More importantly, it is the data structure the user must use to pass the subproblem descriptions back to SYMPHONY at the time a new search tree node is created in the function **user_create_subproblem()** (see Section 6.3.2.2). The structure has 14 fields (listed below) that must be filled out to describe a subproblem (some fields are optional).

A subproblem is a mixed-integer program defined by a matrix of constraints, an objective function, bounds on both the right hand side values of the constraints and the variables, an array indicating which variables are integer, and (optionally) an array of column names that allows SYMPHONY to report the solution back in terms of column names instead of user indices. If the subproblem has **n** variables and **m** constraints, the constraints are given by a constraint coefficient matrix of size $m \times n$ (described in the next paragraph). The sense of each constraint, the right hand side values and bounds on the right hand side (called *range*) are vectors of size **m**. The objective function coefficients and the lower and upper bounds on the variables are vectors of length **n**. The sense of each constraint can be either 'L' (\leq), 'E' ($=$), 'G' (\geq) or 'R' (ranged). For non-ranged rows the range value is 0, for a ranged row the range value must be non-negative and the constraint means that the row activity level has to be between the right hand side value and the right hand side increased by the range value.

Since the coefficient matrix is very often sparse, only the nonzero entries are stored. Each entry of the matrix has a column index, a row index and a coefficient value associated with it. A matrix is specified in the form of the three arrays **matval**, **matind**, and **matbeg**. The array **matval** contains the values of the nonzero entries of the matrix in *column order*; that is, all the entries for the 0th column come first, then the entries for the 1st column, etc. The row index corresponding to each entry of **matval** is listed in **matind** (both of them are of length **nz**, the number of nonzero entries in the matrix). Finally, **matbeg** contains the starting positions of each of the columns in **matval** and **matind**. Thus, **matbeg[i]** is the position of the first entry of column **i** in both **matval** and **matind**. By convention **matbeg** is allocated to be of length **n+1**, with **matbeg[n]** containing the position after the very last entry in **matval** and **matind** (so it is very conveniently equal to **nz**). This representation of a matrix is known as a *column ordered* or *column major* representation. Below are listed the fields that must be filled out to describe a subproblem.

`int n` – The number of columns.
`int m` – The number of rows.
`int nz` – The number of nonzeros.
`double obj_offset` – Constant to be added to the objective function value when printed.
`char obj_sense` – Objective sense (set to `MAXIMIZE` or `MINIMIZE`).
`char *is_int` – Indicates which variables are required to be integer.
`int *matbeg` – The array containing the starting positions for each column.
`int *matind` – The array containing the indices for each column.
`double *matval` – The array containing the matrix values for each column.
`double *obj` – The objective function coefficients for the second objective (for multicriteria solve).
`double *obj2` – The objective function coefficients.
`double *rhs` – The right hand side values for the constraints.
`double *rngval` – The range values for the constraints (optional).
`char *sense` – The senses of the constraints.
`double *lb` – The lower bounds of the variables.
`double *ub` – The upper bounds of the variables.
`char **colname` – The column names.

▷ **cut_data**

Another of the internally defined data structures that the user has to deal with frequently is the `cut_data` data structure, used to store the packed form of cuts. This structure has 8 fields listed below.

`int size` – The size of the `coef` array.

`char *coef` – An array containing the packed form of the cut, which is defined and constructed by the user. Given this packed form and a list of the variables active in the current relaxation, the user must be able to construct the corresponding constraint.

`double rhs` – The right hand side of the constraint.

`double range` – The range of the constraint. It is zero for a standard form constraint. Otherwise, the row activity level is limited to between `rhs` and `rhs + range`.

`char type` – A user-defined type identifier that represents the general class that the cut belongs to.

`char sense` – The sense of the constraint. Can be either 'L' (\leq), 'E' ($=$), 'G' (\geq) or 'R' (ranged). This may be evident from the `type`.

`char deletable` – Determines whether or not a cut can be deleted once added to the formulation. `TRUE` by default.

`char branch` – Determines whether the cut can be branched on or not. Possible initial values are `DO_NOT_BRANCH_ON_THIS_ROW` and `ALLOWED_TO_BRANCH_ON`.

`int name` – Identifier used by SYMPHONY. The user should not set this.

▷ **waiting_row**

A closely related data structure is the `waiting_row`, essentially the “unpacked” form of a cut. There are six fields.

`source_pid` – Used internally by SYMPHONY.

`cut_data *cut` – Pointer to the cut from which the row was generated.

`int nzcnt, *matind, *matval` – Fields describing the row. `nzcnt` is the number of nonzeros in the row, i.e., the length of the `matind` and `matval` arrays, which are the variable indices (wrt. the current LP relaxation) and nonzero coefficients in the row.

`double violation` – If the constraint corresponding to the cut is violated, this value contains the degree of violation (the absolute value of the difference between the row activity level (i.e., lhs) and the right hand side). This value does not have to be set by the user.

▷ **var_desc**

The `var_desc` structure is used list the variables in the current relaxation. There are four fields.

`int userind` – The user index of the variables,

`int colind` – The column index of the variables (in the current relaxation),

`double lb` – The lower bound of the variable,

`double ub` – The upper bound of the variable.

6.3.2.2 Function Descriptions

Now we describe the functions themselves.

▷ **user_receive_lp_data**

```
int user_receive_lp_data (void **user)
```

Description:

This function only has to be filled out for parallel execution and only if either the TM or LP modules are configured as separate processes. Otherwise, data will have been copied into appropriate locations in the master function **user_send_lp_data()** (see Section 6.3.1). The two cases can be handled by means of **#ifdef** statements. See comments in the source code stubs for more details.

Here, the user must receive all problem-specific information sent from the master, set up necessary data structures, etc. Note that the data must be received in exactly the same order as it was sent in **user_send_lp_data()** (see Section 6.3.1). See Section 5.5.1 for more notes on receiving data.

Arguments:

void **user OUT Pointer to the user-defined LP data structure.

Return values:

USER_ERROR Error. SYMPHONY aborts this LP module.
USER_SUCCESS User received the data successfully.
USER_DEFAULT User did not send any data.

Wrapper invoked from: **lp_initialize()** at process start.

▷ `user_create_subproblem`

```
int user_create_subproblem(void *user, int *indices, MIPdesc *mip,
                           int *maxn, int *maxm, int *maxnz)
```

Description:

Based on the instance data contained in the user data structure and the list of base and extra variables that are active in the current subproblem, the user has to create the subproblem for the search node. The matrix describing the subproblem must contain those variables whose user indices are listed in the array `indices` (in the same order) and the base constraints. The extra (dynamically generated) constraints are added automatically by SYMPHONY after the initial subproblem description is created.

In this function, the user is required to construct a description of the subproblem in column-ordered format and pass it back to SYMPHONY by filling out the `MIPdesc` data structure, described in Section 6.3.2.1. The user is not responsible for allocating extra memory to allow for the addition of dynamically generated cuts and variables. The arrays allocated in `user_create_subproblem()` are owned by SYMPHONY after allocation and are freed as soon as the relaxation is loaded into the solver. However, if the user has an idea as to the maximum number of variables and constraints that are likely to be generated during processing of the subproblem, this information can be passed to SYMPHONY in the variables `*maxn`, `*maxm`, and `*maxnz`. These numbers are only estimates that SYMPHONY can use to perform memory allocation. They do not have to be exact numbers. In fact, these estimates need not be provided at all. The `obj_sense` and `obj_offset` fields are set globally in the function `user_initialize_root_node()` (see Section 6.3.1). Setting them here will have no effect.

Note that, the user should return “USER.DEFAULT” if an MPS or GMPL/AMPL file was read in to describe the original MILP. SYMPHONY will allocate the corresponding arrays and specify the constraint matrix automatically in this case.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>int *indices</code>	IN	The list of the active variables (base and extra).
<code>MIPdesc *mip</code>	OUT	Pointer to the <code>MIPdesc</code> data structure.
<code>int *maxn</code>	OUT	Estimated maximum number of variables.
<code>int *maxm</code>	OUT	Estimated maximum number of constraints.
<code>int *maxnz</code>	OUT	Estimated maximum number of nonzeros.

Return values:

<code>USER_ERROR</code>	Error. The LP module is aborted.
<code>USER_SUCCESS</code>	User created the constraint matrix with the base constraints.
<code>USER_DEFAULT</code>	This return code is used when the default routines for reading in an MPS or AMPL file were used and the user wants to let SYMPHONY set up the subproblem automatically. This return will cause an error if the default I/O routines were not used.

Post-processing:

The extra constraints are added to the matrix by calling the `user_unpack_cuts()` subroutine and then adding the corresponding rows to the matrix. This is easier for the user to implement, but less efficient than adding the cuts at the time the original matrix was being constructed.

Wrapper invoked from: `process_chain()` which is invoked when setting up a the initial search node in a chain.

▷ `user_is_feasible`

```
int user_is_feasible(void *user, double lpetol, int varnum, int
                    *indices, double *values, int *feasible,
                    double *objval, int branching, double *heur_solution)
```

Description:

The user can test the feasibility status of the solution to the current LP relaxation and/or return a feasible solution generated by a primal heuristic. This function is primarily used in cases where a solution satisfying integrality restrictions may not be feasible, i.e., it may violate an inequality not present in the current relaxation that will be generated by the user during the cut generation phase. In such a case, it is not possible for SYMPHONY to determine this and the user should set `*feasible` to `IP_INFEASIBLE` and return `USER_SUCCESS`. If the user tests the feasibility status of the solution and determines that the current solution is feasible, then `*feasible` should be set to `IP_FEASIBLE` instead. In this case, the user can set `*objval` to the true objective function value of the feasible solution to account for any round-off error introduced by the solver if desired. If all integral solution must be feasible, the user can ask SYMPHONY to simply test the integrality status by returning `USER_DEFAULT`.

If the user is able to generate a feasible solution using a primal heuristic, then the solution can be returned in a dense format in the array `heur_solution`. In this case, `*objval` should be set to the value of the solution and `*feasible` should be set to `IP_HEUR_FEASIBLE`.

Arguments:

<code>void *user</code>	INOUT	Pointer to the user-defined LP data structure.
<code>double lpetol</code>	IN	The ϵ tolerance of the LP solver.
<code>int varnum</code>	IN	The length of the <code>indices</code> and <code>values</code> arrays.
<code>int *indices</code>	IN	User indices of variables at nonzero level in the current solution.
<code>double *values</code>	IN	Values of the variables listed in <code>indices</code> .
<code>int *feasible</code>	OUT	Feasibility status of the solution (<code>IP_INFEASIBLE</code> , <code>IP_FEASIBLE</code> , or <code>IP_HEUR_FEASIBLE</code>).
<code>double *objval</code>	INOUT	The user can return the “true” objective function value of the solution in the case it is feasible, i.e., eliminating the round-off error. This variable should also be set to the value of the heuristic solution, if one is found.
<code>int branching</code>	In	Indicates whether the function is being invoked during strong branching (from <code>select_branching_object()</code>) or after solving an Lp relaxation (from <code>fathom_branch()</code>).
<code>double *heur_solution</code>	OUT	Used to return a solution determined heuristically.

Return values:

<code>USER_ERROR</code>	Error. Solution is considered to be not feasible.
<code>USER_SUCCESS</code>	User checked IP feasibility.
<code>USER_DEFAULT</code>	SYMPHONY should test the integrality of the solution to determine if it is feasible.

Wrapper invoked from: `select_branching_object()` after pre-solving the LP relaxation of a child corresponding to a candidate and from `fathom_branch()` after solving an LP relaxation.

▷ **user_send_feasible_solution**

```
int user_send_feasible_solution(void *user, double lpetol,  
                               int varnum, int *indices, double *values)
```

Description:

This function is only used for parallel computation. The user can send a feasible solution in custom format to the master module if desired. However, the default routine suffices in most situations. The solution is sent using the communication functions described in Section 5.5.1 in whatever logical format the user wants to use. The default is to pack the user indices and values of variables at non-zero level. If the user packs the solution herself then the same data must be packed here that will be received in the `user_receive_feasible_solution()` function in the master module. See the description of that function for details. This function will only be called when either the LP or tree manager are running as a separate executable. Otherwise, the solution gets stored within the LP user data structure.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>double lpetol</code>	IN	The ϵ tolerance of the LP solver.
<code>int varnum</code>	IN	The length of the <code>indices</code> and <code>values</code> arrays.
<code>int *indices</code>	IN	User indices of variables at nonzero level in the current solution.
<code>double *values</code>	IN	Values of the variables listed in <code>indices</code> .

Return values:

<code>USER_ERROR</code>	Error. Do the default.
<code>USER_SUCCESS</code>	User packed the solution.
<code>USER_DEFAULT</code>	Regulated by the parameter <code>pack_feasible_solution_default</code> , but set to <code>SEND_NONZEROS</code> unless overridden by the user.
<code>SEND_NONZEROS</code>	Pack the nonzero values and their indices.

Wrapper invoked: as soon as feasibility is detected anywhere.

▷ **user_display_lp_solution**

```
int user_display_lp_solution(void *user, int which_sol,
                             int varnum, int *indices, double *values)
```

Description:

Given a solution to an LP relaxation (the indices and values of the nonzero variables) the user can (graphically) display it. The `which_sol` argument shows what kind of solution is passed to the function: `DISP_FEAS_SOLUTION` indicates a solution feasible to the original IP problem, `DISP_RELAXED_SOLUTION` indicates the solution to any LP relaxation and `DISP_FINAL_RELAXED_SOLUTION` indicates the solution to an LP relaxation when no cut has been found. There is no post-processing. Default options print out user indices and values of nonzero or fractional variables on the standard output.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>int which_sol</code>	IN	The type of solution passed on to the displaying function. Possible values are <code>DISP_FEAS_SOLUTION</code> , <code>DISP_RELAXED_SOLUTION</code> and <code>DISP_FINAL_RELAXED_SOLUTION</code> .
<code>int varnum</code>	IN	The number of variables in the current solution at nonzero level (the length of the <code>indices</code> and <code>values</code> arrays).
<code>int *indices</code>	IN	User indices of variables at nonzero level in the current solution.
<code>double *values</code>	IN	Values of the nonzero variables.

Return values:

<code>USER_ERROR</code>	Error. SYMPHONY ignores error message.
<code>USER_SUCCESS</code>	User displayed whatever she wanted to.
<code>USER_DEFAULT</code>	Regulated by the parameter <code>display_solution_default</code> , but set to <code>DISP_NZ_INT</code> unless overridden by the user.
<code>DISP_NOTHING</code>	Display nothing.
<code>DISP_NZ_INT</code>	Display user indices (as integers) and values of nonzero variables.
<code>DISP_NZ_HEXA</code>	Display user indices (as hexadecimal) and values of nonzero variables.
<code>DISP_FRAC_INT</code>	Display user indices (as integers) and values of variables not at their lower or upper bounds.
<code>DISP_FRAC_HEXA</code>	Display user indices (as hexadecimal) and values of variables not at their lower and upper bounds.

Wrapper invoked from: `fathom_branch()` with `DISP_FEAS_SOLUTION` or `DISP_RELAXED_SOLUTION` after solving an LP relaxation and checking its feasibility status. If it was not feasible and no cut could be added either then the wrapper is invoked once more, now with `DISP_FINAL_RELAXED_SOLUTION`.

▷ `user_shall_we_branch`

```
int user_shall_we_branch(void *user, double lpetol, int cutnum,
                        int slacks_in_matrix_num,
                        cut_data **slacks_in_matrix,
                        int slack_cut_num, cut_data **slack_cuts,
                        int varnum, var_desc **vars, double *x,
                        char *status, int *cand_num,
                        branch_obj ***candidates, int *action)
```

Description:

There are two user-written functions invoked from `select_candidates_u`. The first one (`user_shall_we_branch()`) decides whether to branch at all, the second one (`user_select_candidates()`) chooses the branching objects. The argument lists of the two functions are the same, and if branching occurs (see discussion below) then the contents of `*cand_num` and `*candidates` will not change between the calls to the two functions.

The first of these two functions is invoked in each iteration after solving the LP relaxation and (possibly) generating cuts. Therefore, by the time it is called, some violated cuts might be known. Still, the user might decide to branch anyway. The second function is invoked only when branching is decided on.

Given (1) the number of known violated cuts that can be added to the problem when this function is invoked, (2) the constraints that are slack in the LP relaxation, (3) the slack cuts not in the matrix that could be branched on (more on this later), and (4) the solution to the current LP relaxation, the user must decide whether to branch or not. Branching can be done either on variables or slack cuts. A pool of slack cuts which has been removed from the problem and kept for possible branching is passed to the user. If any of these happen to actually be violated (it is up to the user to determine this), they can be passed back as branching candidate type `VIOLATED_SLACK` and will be added into the current relaxation. In this case, branching does not have to occur (the structure of the `*candidates` array is described below in `user_select_candidates()`).

This function has two outputs. The first output is `*action` which can take four values: `USER_DO_BRANCH` if the user wants to branch, `USER_DO_NOT_BRANCH` if he doesn't want to branch, `USER_BRANCH_IF_MUST` if he wants to branch only if there are no known violated cuts, or finally `USER_BRANCH_IF_TAILOFF` if he wants to branch in case tailing off is detected. The second output is the number of candidates and their description. In this function the only sensible "candidates" are `VIOLATED_SLACKS`.

There is no post processing, but in case branching is selected, the `col_gen_before_branch()` function is invoked before the branching would take place. If that function finds dual infeasible variables then (instead of branching) they are added to the LP relaxation and the problem is resolved. (Note that the behavior of the `col_gen_before_branch()` is governed by the `colgen_strat[]` TM parameters.)

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>double lpetol</code>	IN	The ϵ tolerance of the LP solver.
<code>int cutnum</code>	IN	The number of violated cuts (known before invoking this function) that could be added to the problem (instead of branching).
<code>int slacks_in_matrix_num</code>	IN	Number of slack constraints in the matrix.
<code>cut_data **slacks_in_matrix</code>	IN	The description of the cuts corresponding to these constraints (see Section 6.3.2.1).
<code>int slack_cut_num</code>	IN	The number of slack cuts not in the matrix.
<code>cut_data **slack_cuts</code>	IN	Array of pointers to these cuts (see Section 6.3.2.1).
<code>int varnum</code>	IN	The number of variables in the current lp relaxation (the length of the following three arrays).
<code>var_desc **vars</code>	IN	Description of the variables in the relaxation.
<code>double *x</code>	IN	The corresponding solution values (in the optimal solution to the relaxation).
<code>char *status</code>	IN	The status of the variables. There are five possible status values: <code>NOT_FIXED</code> , <code>TEMP_FIXED_TO_UB</code> , <code>PERM_FIXED_TO_UB</code> , <code>TEMP_FIXED_TO_LB</code> and <code>PERM_FIXED_TO_LB</code> .
<code>int *cand_num</code>	OUT	Pointer to the number of candidates returned (the length of <code>*candidates</code>).
<code>candidate ***candidates</code>	OUT	Pointer to the array of candidates generated (see description below).
<code>int *action</code>	OUT	What to do. Must be one of the four above described values unless the return code is <code>USER_DEFAULT</code> .

Return values:

<code>USER_ERROR</code>	Error. <code>DEFAULT</code> is used.
<code>USER_SUCCESS</code>	The user filled out <code>*action</code> (and possibly <code>*cand_num</code> and <code>*candidates</code>).
<code>USER_DEFAULT</code>	Action taken is controlled by the parameter <code>shall_we_branch_default</code> , which is initially <code>USER__BRANCH_IF_MUST</code> unless overridden by the user.

Notes:

- The user has to allocate the pointer array for the candidates and place the pointer for the array into `***candidates` (if candidates are returned).
- Candidates of type `VIOLATED_SLACK` are always added to the LP relaxation regardless of what `action` is chosen and whether branching will be carried out or not.

- Also note that the user can change his mind in `user_select_candidates()` and not branch after all, even if she chose to branch in this function. A possible scenario: `cut_num` is zero when this function is invoked and the user asks for `USER__BRANCH_IF_MUST` without checking the slack constraints and slack cuts. Afterward no columns are generated (no dual infeasible variables found) and thus SYMPHONY decides branching is called for and invokes `user_select_candidates()`. However, in that function the user checks the slack cuts, finds that some are violated, cancels the branching request and adds the violated cuts to the relaxation instead.

Warning: The cuts the user unpacks and wants to be added to the problem (either because they are of type `VIOLATED_SLACK` or type `CANDIDATE_CUT_NOT_IN_MATRIX`) will be deleted from the list of slack cuts after this routine returns. Therefore the same warning applies here as in the function `user_unpack_cuts()`.

Wrapper invoked from: `select_branching_object()`.

▷ **user_select_candidates**

```
int user_select_candidates(void *user, double lpetol, int cutnum,
                           int slacks_in_matrix_num,
                           cut_data **slacks_in_matrix,
                           int slack_cut_num, cut_data **slack_cuts,
                           int varnum, var_desc **vars, double *x,
                           char *status, int *cand_num,
                           branch_obj ***candidates, int *action,
                           int bc_level)
```

Description:

The purpose of this function is to generate branching candidates. Note that ***action** from **user_shall_we_branch()** is passed on to this function (but its value can be changed here, see notes at the previous function), as well as the candidates in ****candidates** and their number in ***cand_num** if there were any.

Violated cuts found among the slack cuts (not in the matrix) can be added to the candidate list. These violated cuts will be added to the LP relaxation regardless of the value of ***action**.

The **branch_obj** structure contains fields similar to the **cut_data** data structure. Branching is accomplished by imposing inequalities which divide the current subproblem while cutting off the corresponding fractional solution. Branching on cuts and variables is treated symmetrically and branching on a variable can be thought of as imposing a constraint with a single unit entry in the appropriate column. Following is a list of the fields of the **branch_obj** data structure which must be set by the user.

char type Can take five values:

CANDIDATE_VARIABLE The object is a variable.

CANDIDATE_CUT_IN_MATRIX The object is a cut (it must be slack) which is in the current formulation.

CANDIDATE_CUT_NOT_IN_MATRIX The object is a cut (it must be slack) which has been deleted from the formulation and is listed among the slack cuts.

VIOLATED_SLACK The object is not offered as a candidate for branching, but rather it is selected because it was among the slack cuts but became violated again.

SLACK_TO_BE_DISCARDED The object is not selected as a candidate for branching rather it is selected because it is a slack cut which should be discarded even from the list of slack cuts.

int position The position of the object in the appropriate array (which is one of **vars**, **slacks_in_matrix**, or **slack_cuts**).

waiting_row *row Used only if the type is **CANDIDATE_CUT_NOT_IN_MATRIX** or **VIOLATED_SLACK**. In these cases this field holds the row extension corresponding to the cut. This structure can be filled out easily using a call to **user_unpack_cuts()**.

int child_num

The number of children of this branching object.

`char *sense, double *rhs, double *range, int *branch`

The description of the children. These arrays determine the sense, rhs, etc. for the cut to be imposed in each of the children. These are defined and used exactly as in the `cut_data` data structure. **Note:** If a limit is defined on the number of children by defining the `MAX_CHILDREN_NUM` macro to be a number (it is pre-defined to be 4 as a default), then these arrays will be statically defined to be the correct length and don't have to be allocated. This option is highly recommended. Otherwise, the user must allocate them to be of length `child_num`.

`double lhs` The activity level for the row (for branching cuts). This field is purely for the user's convenience. SYMPHONY doesn't use it so it need not be filled out.

`double *objval, int *termcode, int *iterd, int *feasible`

The objective values, termination codes, number of iterations and feasibility status of the children after pre-solving them. These are all filled out by SYMPHONY during strong branching. The user may access them in `user_compare_candidates()` (see below).

There are three default options (see below), each chooses a few variables (the number is determined by the strong branching parameters (see Section 6.4.5)).

Arguments:

Same as for `user_shall_we_branch()`, except that `*action` must be either `USER_DO_BRANCH` or `USER_DO_NOT_BRANCH`, and if branching is asked for, there must be a real candidate in the candidate list (not only `VIOLATED_SLACKS` and `SLACK_TO_BE_DISCARDEDs`). Also, the argument `bc_level` is the level in the tree. This could be used in deciding how many strong branching candidates to use.

Return values:

<code>USER_ERROR</code>	Error. DEFAULT is used.
<code>USER_SUCCESS</code>	User generated branching candidates.
<code>USER_DEFAULT</code>	Regulated by the <code>select_candidates_default</code> parameter, but set to <code>USER_CLOSE_TO_HALF</code> unless overridden by the user.
<code>USER_CLOSE_TO_HALF</code>	Choose variables with values closest to half.
<code>USER_CLOSE_TO_HALF_AND_EXPENSIVE</code>	Choose variables with values close to half and with high objective function coefficients.
<code>USER_CLOSE_TO_ONE_AND_CHEAP</code>	Choose variables with values close to one and with low objective function coefficients.

Wrapper invoked from: `select_branching_object()`.

Notes: See the notes at `user_shall_we_branch()`.

▷ user_compare_candidates

```
int user_compare_candidates(void *user, branch_obj *can1, branch_obj *can2,  
                           double ub, double granularity,  
                           int *which_is_better)
```

Description:

By the time this function is invoked, the children of the current search tree node corresponding to each branching candidate have been pre-solved, i.e., the `objval`, `termcode`, `iterd`, and `feasible` fields of the `can1` and `can2` structures are filled out. Note that if the termination code for a child is `LP_D_UNBOUNDED` or `LP_D_OBJLIM`, i.e., the dual problem is unbounded or the objective limit is reached, then the objective value of that child is set to `MAXDOUBLE / 2`. Similarly, if the termination code is one of `LP_D_ITLIM` (iteration limit reached), `LP_D_INFEASIBLE` (dual infeasible) or `LP_ABANDONED` (because of numerical difficulties) then the objective value of that child is set to that of the parent's.

Based on this information the user must choose which candidate he considers better and whether to branch on this better one immediately without checking the remaining candidates. As such, there are four possible answers: `FIRST_CANDIDATE_BETTER`, `SECOND_CANDIDATE_BETTER`, `FIRST_CANDIDATE_BETTER_AND_BRANCH_ON_IT` and `SECOND_CANDIDATE_BETTER_AND_BRANCH_ON_IT`. An answer ending with `_AND_BRANCH_ON_IT` indicates that the user wants to terminate the strong branching process and select that particular candidate for branching.

There are several default options. In each of them, objective values of the pre-solved LP relaxations are compared.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>branch_obj *can1</code>	IN	One of the candidates to be compared.
<code>branch_obj *can2</code>	IN	The other candidate to be compared.
<code>double ub</code>	IN	The current best upper bound.
<code>double granularity</code>	IN	Defined tolerance
<code>int *which_is_better</code>	OUT	The user's choice. See the description above.

Return values:

USER_ERROR	Error. DEFAULT is used.
USER_SUCCESS	User filled out <code>*which_is_better</code> .
USER_DEFAULT	Regulated by the <code>compare_candidates_default</code> parameter, initially set to <code>LOWEST_LOW_OBJ</code> unless overridden by the user.
BIGGEST_DIFFERENCE_OBJ	Prefer the candidate with the biggest difference between highest and lowest objective function values.
LOWEST_LOW_OBJ	Prefer the candidate with the lowest minimum objective function value. The minimum is taken over the objective function values of all the children.
HIGHEST_LOW_OBJ	Prefer the candidate with the highest minimum objective function value.
LOWEST_HIGH_OBJ	Prefer the candidate with the lowest maximum objective function value.
HIGHEST_HIGH_OBJ	Prefer the candidate with the highest maximum objective function value .
LOWEST_LOW_FRAC	Prefer the candidate with the lowest minimum number of fractional variables. The minimum is taken over the number of fractional variables in all the children. Fractional branching options are only available if the fractional branching compile-time option is set in the makefile.
HIGHEST_LOW_FRAC	Prefer the candidate with the highest minimum number of fractional variables.
LOWEST_HIGH_FRAC	Prefer the candidate with the lowest maximum number of fractional variables.
HIGHEST_HIGH_FRAC	Prefer the candidate with the highest maximum number of fractional variables.

Wrapper invoked from: `select_branching_object()` after the LP relaxations of the children have been pre-solved.

▷ **user_select_child**

```
int user_select_child(void *user, double ub, branch_obj *can, char *action)
```

Description:

By the time this function is invoked, the candidate for branching has been chosen. Based on this information and the current best upper bound, the user has to decide what to do with each child. Possible actions for a child are **KEEP_THIS_CHILD** (the child will be kept at this LP for further processing, i.e., the process *dives* into that child), **PRUNE_THIS_CHILD** (the child will be pruned based on some problem specific property—no questions asked...), **PRUNE_THIS_CHILD_FATHOMABLE** (the child will be pruned based on its pre-solved LP relaxation) and **RETURN_THIS_CHILD** (the child will be sent back to tree manager). Note that at most one child can be kept at the current LP module.

There are two default options—in both of them, objective values of the pre-solved LP relaxations are compared (for those children whose pre-solve did not terminate with primal infeasibility or high cost). One rule prefers the child with the lowest objective function value and the other prefers the child with the higher objective function value.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>int ub</code>	IN	The current best upper bound.
<code>branch_obj *can</code>	IN	The branching candidate.
<code>char *action</code>	OUT	Array of actions for the children. The array is already allocated to length <code>can->number</code> .

Return values:

<code>USER_ERROR</code>	Error. <code>DEFAULT</code> is used.
<code>USER_SUCCESS</code>	User filled out <code>*action</code> .
<code>USER_DEFAULT</code>	Regulated by the <code>select_child.default</code> parameter, which is initially set to <code>PREFER_LOWER_OBJ_VALUE</code> , unless overridden by the user.
<code>PREFER_HIGHER_OBJ_VALUE</code>	Choose child with the highest objective value.
<code>PREFER_LOWER_OBJ_VALUE</code>	Choose child with the lowest objective value.
<code>PREFER_MORE_FRACTIONAL</code>	Choose child with the most fractional variables. Fractional branching options are only available if the fractional branching compile-time option is set in the makefile.
<code>PREFER_LESS_FRACTIONAL</code>	Choose child with the lowest number of fractional variables.

Post-processing:

Checks which children can be fathomed based on the objective value of their pre-solved LP relaxation.

Wrapper invoked from: `branch()`.

▷ **user_print_branch_stat**

```
int user_print_branch_stat(void *user, branch_obj *can, cut_data *cut,  
                           int n, var_desc **vars, char *action)
```

Description:

Print out information about branching candidate **can**, such as a more explicit problem-specific description than SYMPHONY can provide (for instance, end points of an edge). If **verbosity** is set high enough, the identity of the branching object and the children (with objective values and termination codes for the pre-solved LPs) is printed out to the standard output by SYMPHONY.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>branch_obj *can</code>	IN	The branching candidate.
<code>cut_data *cut</code>	IN	The description of the cut if the branching object is a cut.
<code>int n</code>	IN	Number of variables.
<code>var_desc **vars</code>	IN	Array of variables in the current relaxation.
<code>char *action</code>	IN	Array of actions for the children.

Return values:

<code>USER_ERROR</code>	Error. Ignored by SYMPHONY.
<code>USER_SUCCESS</code>	The user printed out whatever she wanted to.
<code>USER_DEFAULT</code>	SYMPHONY prints out its own branching information.

Wrapper invoked from: `branch()` after the best candidate has been selected, pre-solved, and the action is decided on for the children.

▷ **user_add_to_desc**

```
int user_add_to_desc(void *user, int *desc_size, char **desc)
```

Description:

Before a node description is sent to the TM, the user can provide a pointer to a data structure that will be appended to the description for later use by the user in reconstruction of the node. This information must be placed into ***desc**. Its size should be returned in ***desc_size**.

There is only one default option: the description to be added is considered to be of zero length, i.e., there is no additional description.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>int *desc_size</code>	OUT	The size of the additional information, the length of *desc in bytes.
<code>char **desc</code>	OUT	Pointer to the additional information (space must be allocated by the user).

Return values:

<code>USER_ERROR</code>	Error. <code>DEFAULT</code> is used.
<code>USER_SUCCESS</code>	User filled out *desc_size and *desc .
<code>USER_DEFAULT</code>	No description is appended.

Wrapper invoked from: `create_explicit_node_desc()` before a node is sent to the tree manager.

▷ **user_same_cuts**

```
int user_same_cuts (void *user, cut_data *cut1, cut_data *cut2,  
                   int *same_cuts)
```

Description:

Determine whether the two cuts are comparable (the normals of the half-spaces corresponding to the cuts point in the same direction) and if yes, which one is stronger. The default is to declare the cuts comparable only if the **type**, **sense** and **coef** fields of the two cuts are the same byte by byte; and if this is the case to compare the right hand sides to decide which cut is stronger.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>cut_data *cut1</code>	IN	The first cut.
<code>cut_data *cut2</code>	IN	The second cut.
<code>int *same_cuts</code>	OUT	Possible values: <code>SAME</code> , <code>FIRST_CUT_BETTER</code> , <code>SECOND_CUT_BETTER</code> and <code>DIFFERENT</code> (i.e., not comparable).

Return values:

<code>USER_ERROR</code>	Error. <code>DEFAULT</code> is used.
<code>USER_SUCCESS</code>	User did the comparison, filled out <code>*same_cuts</code> .
<code>USER_DEFAULT</code>	Compare byte by byte (see above).

Wrapper invoked from: `process_message()` when a `PACKED_CUT` arrives.

Note:

This function is used to check whether a newly arrived cut is already in the local pool. If so, or if it is weaker than a cut in the local pool, then the new cut is discarded; if it is stronger than a cut in the local pool, then the new cut replaces the old one and if the new is different from all the old ones, then it is added to the local pool.

▷ **user_unpack_cuts**

```
int user_unpack_cuts(void *user, int from, int type, int varnum,
                    var_desc **vars, int cutnum, cut_data **cuts,
                    int *new_row_num, waiting_row ***new_rows)
```

Description:

If the user has defined application-specific cut classes, these cuts must be interpreted as constraints for the current LP relaxation, i.e., the user must decode the compact representation of the cuts (see the `cut_data` structure) into rows for the matrix. A pointer to the array of generated rows must be returned in `***new_rows` (the user has to allocate this array) and their number in `*new_row_num`.

Note that SYMPHONY has built-in support for cuts generated explicitly as matrix rows with no user-defined packed form, i.e., cuts whose indices and coefficients are given explicitly (see the function `user_find_cuts()` in Section 6.3.3. These cuts can be constructed and added using the helper function `cg_add_explicit_cut()` (see the description of `user_find_cuts()` in Section 6.3.3) and are packed and unpacked automatically, so the user does not need to implement this function. In post processing, SYMPHONY unpacks explicitly defined cuts and internally generated generic cuts.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>int from</code>	IN	See below in “Notes”.
<code>int type</code>	IN	UNPACK_CUTS_SINGLE or UNPACK_CUTS_MULTIPLE (see notes below).
<code>int varnum</code>	IN	The number of variables.
<code>var_desc **vars</code>	IN	The variables currently in the problem.
<code>int cutnum</code>	IN	The number of cuts to be decoded.
<code>cut_data **cuts</code>	IN	Cuts that need to be converted to rows for the current LP. See “Warning” below.
<code>int *new_row_num</code>	OUT	Pointer to the number of rows in <code>***new_rows</code> .
<code>waiting_row ***new_rows</code>	OUT	Pointer to the array of pointers to the new rows.

Return values:

USER_ERROR	Error. The cuts are discarded.
USER_SUCCESS	User unpacked the cuts.
USER_DEFAULT	There are no user cut types defined. In this case, SYMPHONY deals with only explicit cuts and internally generated cuts.

Wrapper invoked from: Wherever a cut needs to be unpacked (multiple places).

Post-processing:

Explicit row cuts are processed, as well as SYMPHONY’s internally generated cuts. Also, the pointers to each cut are transferred to the `waiting_rows` data structure (in previous version, this was done by the user).

Notes:

- When decoding the cuts, the expanded constraints have to be adjusted to the current LP, i.e., coefficients corresponding to variables currently not in the LP have to be left out.
- If the `one_row_only` flag is set to `UNPACK_CUTS_MULTIPLE`, then the user can generate as many constraints (even zero!) from a cut as she wants (this way she can lift the cuts, thus adjusting them for the current LP). However, if the flag is set to `UNPACK_CUTS_SINGLE`, then for each cut the user must generate a unique row, the same one that had been generated from the cut before. (The flag is set to this value only when regenerating a search tree node.)
- The `from` argument can take on six different values: `CUT_FROM_CG`, `CUT_FROM_CP`, `CUT_FROM_TM`, `CUT_LEFTOVER` (these are cuts from a previous LP relaxation that are still in the local pool), `CUT_NOT_IN_MATRIX_SLACK` and `CUT_VIOLATED_SLACK` indicating where the cut came from. This might be useful in deciding whether to lift the cut or not.
- The `matind` fields of the rows must be filled with indices with respect to the position of the variables in `**vars`.
- **Warning:** For each row, the user must make sure that the cut the row was generated from (and can be uniquely regenerated from if needed later) is safely stored in the `waiting_row` structure. SYMPHONY will free the entries in `cuts` after this function returns. If a row is generated from a cut in `cuts` (and not from a lifted cut), the user has the option of physically copying the cut into the corresponding part of the `waiting_row` structure, or copying the pointer to the cut into the `waiting_row` structure and erasing the pointer in `cuts`. If a row is generated from a lifted cut, the user should store a copy of the lifted cut in the corresponding part of `waiting_row`.

▷ **user_send_lp_solution**

```
int user_send_lp_solution(void *user, int varnum, var_desc **vars,  
                          double *x, int where)
```

Description:

This function is only used in the case of parallel execution. The user has the option to send the LP solution to either the cut pool or the cut generator in some user-defined form if desired. There are two default options—sending the indices and values for all nonzero variables (`SEND_NONZEROS`) and sending the indices and values for all fractional variables (`SEND_FRACTIONS`).

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>int varnum</code>	IN	The number of variables currently in the LP relaxation. (The length of the <code>*vars</code> and <code>x</code> arrays.)
<code>var_desc **vars</code>	IN	The variables currently in the LP relaxation.
<code>double *x</code>	IN	Values of the above variables.
<code>int where</code>	IN	Where the solution is to be sent— <code>LP_SOL_TO_CG</code> or <code>LP_SOL_TO_CP</code> .

Return values:

<code>USER_ERROR</code>	Error. No message will be sent.
<code>USER_SUCCESS</code>	User packed and sent the message.
<code>USER_DEFAULT</code>	Regulated by the <code>pack_lp_solution.default</code> parameter, initially set to <code>SEND_NOZEROS</code> .
<code>SEND_NONZEROS</code>	Send user indices and values of variables at nonzero level.
<code>SEND_FRACTIONS</code>	Send user indices and values of variables at fractional level.

Wrapper invoked from: `fathom_branch()` after an LP relaxation has been solved. The message is always sent to the cut generator (if there is one). The message is sent to the cut pool if a search tree node at the top of a chain is being processed (except at the root in the first phase), or if a given number (`cut_pool_check_freq`) of LP relaxations have been solved since the last check.

Note:

The wrapper automatically packs the level, index, and iteration number corresponding to the current LP solution within the current search tree node, as well as the objective value and upper bound in case the solution is sent to a cut generator. This data will be unpacked by SYMPHONY on the receiving end, the user will have to unpack there exactly what he has packed here.

▷ user_logical_fixing

```
int user_logical_fixing(void *user, int varnum, var_desc **vars,  
                        double *x, char *status, int *num_fixed)
```

Description:

Logical fixing is modifying the stati of variables based on logical implications derived from problem-specific information. In this function the user can modify the status of any variable. Valid stati are: `NOT_FIXED`, `TEMP_FIXED_TO_LB`, `PERM_FIXED_TO_LB`, `TEMP_FIXED_TO_UB` and `PERM_FIXED_TO_UB`. Be forewarned that fallaciously fixing a variable in this function can cause the algorithm to terminate improperly. Generally, a variable can only be fixed permanently if the matrix is *full* at the time of the fixing (i.e. all variables that are not fixed are in the matrix). There are no default options.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>int varnum</code>	IN	The number of variables currently in the LP relaxation. (The length of the <code>*vars</code> and <code>x</code> arrays.)
<code>var_desc **vars</code>	IN	The variables currently in the LP relaxation.
<code>double *x</code>	IN	Values of the above variables.
<code>char *status</code>	INOUT	Stati of variables currently in the LP relaxation.
<code>int *num_fixed</code>	OUT	Number of fixed variables.

Return values:

<code>USER_ERROR</code>	Error. Ignored by SYMPHONY.
<code>USER_SUCCESS</code>	User changed the stati of the variables she wanted.
<code>USER_DEFAULT</code>	No logical fixing rules are implemented.

Wrapper invoked from: `fix_variables()` after doing reduced cost fixing, but only when a specified number of variables have been fixed by reduced cost (see LP parameter settings).

▷ **user_generate_column**

```
int user_generate_column(void *user, int generate_what, int cutnum,  
                        cut_data **cuts, int prevind, int nextind,  
                        int *real_nextind, double *colval,  
                        int *colind, int *collen, double *obj,  
                        double *lb, double *ub)
```

Description:

This function is called when pricing out the columns that are not already fixed and are not explicitly represented in the matrix. Only the user knows the explicit description of these columns. When a missing variable need to be priced, the user is asked to provide the corresponding column. SYMPHONY scans through the known variables in the order of their user indices. After testing a variable in the matrix (**prevind**), SYMPHONY asks the user if there are any missing variables to be priced before the next variable in the matrix (**nextind**). If there are missing variables before **nextind**, the user has to supply the user index of the real next variable (**real_nextind**) along with the corresponding column. Occasionally SYMPHONY asks the user to simply supply the column corresponding to **nextind**. The **generate_what** flag is used for making a distinction between the two cases: in the former case it is set to **GENERATE_REAL_NEXTIND** and in the latter it is set to **GENERATE_NEXTIND**.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>int generate_what</code>	IN	<code>GENERATE_NEXTIND</code> or <code>GENERATE_REAL_NEXTIND</code> (see description above).
<code>int cutnum</code>	IN	The number of added rows in the LP formulation (i.e., the total number of rows less the number of base constraints). This is the length of the <code>**cuts</code> array.
<code>cut_data **cuts</code>	IN	Description of the cuts corresponding to the added rows of the current LP formulation. The user is supposed to know about the cuts corresponding to the base constraints.
<code>int prevind</code>	IN	The last variable processed (<code>-1</code> if there was none) by SYMPHONY.
<code>int nextind</code>	IN	The next variable (<code>-1</code> if there are none) known to SYMPHONY.
<code>int *real_nextind</code>	OUT	Pointer to the user index of the next variable (<code>-1</code> if there is none).
<code>double *colval</code>	OUT	Values of the nonzero entries in the column of the next variable. (Sufficient space is already allocated for this array.)
<code>int *colind</code>	OUT	Row indices of the nonzero entries in the column. (Sufficient space is already allocated for this array.)
<code>int *collen</code>	OUT	The length of the <code>colval</code> and <code>colind</code> arrays.
<code>double *obj</code>	OUT	Objective coefficient corresponding to the next variable.
<code>double *lb</code>	OUT	Lower bound of the next variable.
<code>double *ub</code>	OUT	Upper bound of the next variable.

Return values:

<code>USER_ERROR</code>	Error. The LP process is aborted.
<code>USER_SUCCESS</code>	User filled out <code>*real_nextind</code> and generated its column if needed.
<code>USER_DEFAULT</code>	No column generation is done.

Wrapper invoked from: `price_all_vars()` and `restore_lp_feasibility()`.

Note:

`colval`, `colind`, `collen` and `obj` do not need to be filled out if `real_nextind` is the same as `nextind` and `generate_what` is `GENERATE_REAL_NEXTIND`.

▷ `user_generate_cuts_in_lp`

```
int user_generate_cuts_in_lp(void *user, LPdata *lp_data, int varnum,
                             var_desc **vars, double *x, int *new_row_num,
                             cut_data ***cuts)
```

Description:

The user might decide to generate cuts directly within the LP module instead of using the cut generator. This can be accomplished either through a call to this function or simply by configuring SYMPHONY such that the cut generator is called directly from the LP solver. One example of when this might be done is when generating Gomory cuts or something else that requires knowledge of the current LP tableau. The user must return the list of generated cuts by allocating an array of `cut_data` structures and setting `*cuts` to point to this array. Post-processing consists of checking if any of the new cuts are already in the local pool (or dominated by a cut in the local pool).

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>LPdata *lp_data</code>	IN	A pointer to SYMPHONY's internal data structure for storing the LP relaxation and related data.
<code>int varnum</code>	IN	The number of variables currently in the LP relaxation. (The length of the <code>*vars</code> and <code>x</code> arrays.)
<code>var_desc **vars</code>	IN	The variables currently in the LP relaxation.
<code>double *x</code>	IN	Values of the above variables.
<code>int *new_row_num</code>	OUT	The number of cuts generated.
<code>cut_data ***cuts</code>	OUT	The cuts and the corresponding rows.

Return values:

<code>USER_ERROR</code>	Error. Interpreted as if no cuts were generated.
<code>USER_SUCCESS</code>	Cuts were generated.
<code>USER_DEFAULT</code>	No cuts were generated. By default, SYMPHONY uses the CGL to generate generic cuts, according to parameter settings.
<code>GENERATE_CGL_CUTS</code>	Generate generic CGL cuts, according to parameter settings.
<code>DO_NOT_GENERATE_CGL_CUTS</code>	No additional cuts are generated.

Post-processing:

SYMPHONY checks if any of the newly generated cuts are already in the local pool.

Wrapper invoked from: `receive_cuts()` before the cuts from the CG module are received. Since the user will probably use this function to generate tableau-dependent cuts, it is highly unlikely that any of the new cuts would already be in the pool. Therefore the user will probably return `USER_AND_PP` to force SYMPHONY to skip post-processing.

Notes:

- Just like in `user_unpack_cuts()`, the user has to allocate space for the rows.

- Unless the `name` field of a cut is explicitly set to `CUT__SEND_TO_CP`, SYMPHONY will assume that the cut is locally valid only and set that field to `CUT__DO_NOT_SEND_TO_CP`.

▷ **user_print_stat_on_cuts_added**

```
int user_print_stat_on_cuts_added(void *user, int rownum, waiting_row **rows)
```

Description:

The user can print out some information (if he wishes to) on the cuts that will be added to the LP formulation. The default is to print out the number of cuts added.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>int rownum</code>	IN	The number of cuts added.
<code>waiting_row **rows</code>	IN	Array of waiting rows containing the cuts added.

Return values:

<code>USER_ERROR</code>	Revert to default.
<code>USER_SUCCESS</code>	User printed whatever he wanted.
<code>USER_DEFAULT</code>	Print out the number of cuts added.

Wrapper invoked from: `add_best_waiting_rows()` after it has been decided how many cuts to add and after the cuts have been selected from the local pool.

▷ **user_purge_waiting_rows**

```
int user_purge_waiting_rows(void *user, int rownum,
                           waiting_row **rows, char *delete_rows)
```

Description:

The local pool is purged from time to time to control its size. In this function the user has the power to decide which cuts to purge from this pool if desired. To mark the i^{th} waiting row (an element of the pre-pool) for removal she has to set `delete_rows[i]` to be TRUE (`delete_rows` is allocated before the function is called and its elements are set to FALSE by default).

Post-processing consists of actually deleting those entries from the waiting row list and compressing the list. The default is to discard the least violated waiting rows and keep no more than what can be added in the next iteration (this is determined by the `max_cut_num_per_iter` parameter).

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined LP data structure.
<code>int rownum</code>	IN	The number of waiting rows.
<code>waiting_row **rows</code>	IN	The array of waiting rows.
<code>char *delete_rows</code>	OUT	An array of indicators showing which waiting rows are to be deleted.

Return values:

<code>USER_ERROR</code>	Purge every single waiting row.
<code>USER_SUCCESS</code>	The user marked in <code>delete</code> the rows to be deleted.
<code>USER_DEFAULT</code>	Described above.

Post-processing:

The marked rows are deleted.

Wrapper invoked from: `receive_cuts()` after cuts have been added.

▷ **user_free_lp**

```
int user_free_lp(void **user)
```

Description:

The user has to free all the data structures within ***user**, and also free **user** itself. The user can use the built-in macro **FREE** that checks the existence of a pointer before freeing it.

Arguments:

void **user INOUT Pointer to the user-defined LP data structure.

Return values:

USER_ERROR Error. SYMPHONY ignores error message.
USER_SUCCESS User freed everything in the user space.
USER_DEFAULT There is no user memory to free.

Wrapper invoked from: `lp_close()` at module shutdown.

6.3.3 Cut generator module callbacks

Due to the relative simplicity of the cut generator, there are no wrapper functions implemented for CG. Consequently, there are no default options and no post-processing.

▷ **user_receive_cg_data**

```
int user_receive_cg_data (void **user)
```

Description:

This function only has to be filled out for parallel execution and only if the TM, LP, and CG modules are all compiled as separate modules. This would not be typical. If needed, the user can use this function to receive problem-specific data needed for computation in the CG module. The same data must be received here that was sent in the **user_send_cg_data()** (see Section 6.3.1) function in the master module. The user has to allocate space for all the data structures, including **user** itself. Note that some or all of this may be done in the function **user_send_cg_data()** if the Tree Manager, LP, and CG are all compiled together. See that function for more information.

Arguments:

void **user INOUT Pointer to the user-defined data structure.

Return values:

USER_ERROR Error. CG exits.
USER_SUCCESS The user received the data properly.
USER_DEFAULT User did not send any data.

Invoked from: **cg_initialize()** at process start.

▷ **user_receive_lp_solution_cg**

```
int user_receive_lp_solution_cg(void *user)
```

Description:

This function is invoked only in the case of parallel computation and only if in the `user_send_lp_solution()` function of the LP module the user opted for packing the current LP solution herself. Here she must receive the data sent from there.

Arguments:

`void *user` IN Pointer to the user-defined data structure.

Invoked from: Whenever an LP solution is received.

Return values:

<code>USER_ERROR</code>	Error. The LP solution was not received and will not be processed.
<code>USER_SUCCESS</code>	The user received the LP solution.
<code>USER_DEFAULT</code>	The solution was sent by SYMPHONY and will be received automatically.

Note:

SYMPHONY automatically unpacks the level, index and iteration number corresponding to the current LP solution within the current search tree node as well as the objective value and upper bound.

▷ `user_find_cuts`

```
int user_find_cuts(void *user, int varnum, int iter_num, int level,  
    int index, double objval, int *indices, double *values,  
    double ub, double lpetol, int *cutnum)
```

Description:

In this function, the user can generate cuts based on the current LP solution stored in `soln`. Cuts found are added to the LP by calling the `cg_add_user_cut(cut_data *new_cut)` function, which then transfers the cut to the LP module, either through message passing or shared memory. The argument of this function is a pointer to the cut to be sent. See Section 6.3.2 for a description of this data structure. Each user-defined cut assigned a type and a designated packed form. Each user-defined type must be recognized by the user's `user_unpack_cuts()` 6.3.2.2 function in the master module. If the user wants a user-defined cut to be added to the cut pool in case it proves to be effective in the LP, then `new_cut->name` should be set to `CUT_SEND_TO_CP`. In this case, the cut must be globally valid. Otherwise, it should be set to `CUT_DO_NOT_SEND_TO_CP`.

Alternatively, SYMPHONY provides automated support for the generation of cuts represented explicitly as matrix rows. These cuts are passed as sparse vectors and can be added by calling the routine `cg_add_explicit_cut()`, which has the following interface.

```
int cg_add_explicit_cut(int nzcnt, int *indices, double *values,  
    double rhs, double range, char sense,  
    char send_to_cp)
```

Here, `nzcnt` is the number of nonzero coefficients in the cut, `indices` is an array containing the indices of the columns with nonzero entries, and `values` is an array of the corresponding values. The right hand side value is passed in through the variable `rhs`, the range is passed in through the variable `range`, and the sense of the inequality is passed through the variable `sense`. Finally, the variable `send_to_cp` indicates to SYMPHONY whether the cut is globally valid and should be sent to the cut pool, or whether it is only to be used locally.

The only output of the `user_find_cuts()` function is the number of cuts generated and this value is returned in the last argument. For options to generate generic cuts automatically using the COIN Cut Generation Library, see the function `user_generate_cuts_in_lp()` 6.3.2.2

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>int iter_num</code>	IN	The iteration number of the current LP solution.
<code>int level</code>	IN	The level in the tree on which the current LP solution was generated.
<code>index</code>	IN	The index of the node in which LP solution was generated.
<code>objval</code>	IN	The objective function value of the current LP solution.
<code>int varnum</code>	IN	The number of nonzeros in the current LP solution.
<code>indices</code>	IN	The column indices of the nonzero variables in the current LP solution.
<code>values</code>	IN	The values of the nonzero variables listed in <code>indices</code> .
<code>double ub</code>	IN	The current global upper bound.
<code>double lpetol</code>	IN	The current error tolerance in the LP.
<code>int *cutnum</code>	OUT	Pointer to the number of cuts generated and sent to the LP.

Return values:

<code>USER_ERROR</code>	Ignored.
<code>USER_SUCCESS</code>	The user function exited properly.
<code>USER_DEFAULT</code>	No cuts were generated.

Invoked from: Whenever an LP solution is received.

▷ **user_check_validity_of_cut**

```
int user_check_validity_of_cut(void *user, cut_data *new_cut)
```

Description:

This function is provided as a debugging tool. Every cut that is to be sent to the LP solver is first passed to this function where the user can independently verify that the cut is valid by testing it against a known feasible solution (usually an optimal one). This is useful for determining why a particular known feasible (optimal) solution was never found. Usually, this is due to an invalid cut being added. See Section 5.6.3 for more on this feature.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>cut_data *new_cut</code>	IN	Pointer to the cut that must be checked.

Return values:

<code>USER_ERROR</code>	Ignored.
<code>USER_SUCCESS</code>	The user is done checking the cut.
<code>USER_DEFAULT</code>	The cut is ignored.

Invoked from: Whenever a cut is being sent to the LP.

▷ **user_free_cg**

```
int user_free_cg(void **user)
```

Description:

The user has to free all the data structures within **user**, and also free **user** itself. The user can use the built-in macro **FREE** that checks the existence of a pointer before freeing it.

Arguments:

<code>void **user</code>	INOUT	Pointer to the user-defined data structure (should be NULL on exit from this function).
--------------------------	-------	---

Return values:

<code>USER_ERROR</code>	Ignored.
<code>USER_SUCCESS</code>	The user freed all data structures.
<code>USER_DEFAULT</code>	The user has no memory to free.

Invoked from: `cg_close()` at module shutdown.

6.3.4 Cut pool module callbacks

Due to the relative simplicity of the cut pool, there are no wrapper functions implemented for CP. Consequently, there are no default options and no post-processing.

▷ **user_receive_cp_data**

```
int user_receive_cp_data(void **user)
```

Description:

The user has to receive here all problem-specific information sent from **user_send_cp_data()** (see Section 6.3.1) function in the master module. The user has to allocate space for all the data structures, including **user** itself. Note that this function is only called if the either the Tree Manager, LP, or CP are running as a separate process (i.e. either **COMPILE_IN_TM**, **COMPILE_IN_LP**, or **COMPILE_IN_CP** are set to **FALSE** in the **make** file). Otherwise, this is done in **user_send_cp_data()**. See the description of that function for more details.

Arguments:

void **user INOUT Pointer to the user-defined data structure.

Return values:

USER_ERROR Error. Cut pool module exits.
USER_SUCCESS The user received data successfully.
USER_DEFAULT The user did not send any data to be received.

Invoked from: **cp_initialize** at module start-up.

▷ **user_receive_lp_solution_cp**

```
void user_receive_lp_solution_cp(void *user)
```

Description:

This function is invoked only in the case parallel computation and only if in the `user_send_lp_solution()` function of the LP module, the user opted for packing the current LP solution in a custom format. Here she must receive the data she sent there.

Arguments:

`void *user` IN Pointer to the user-defined data structure.

Return values:

`USER_ERROR` Cuts are not checked for this LP solution.

`USER_SUCCESS` The user function executed properly.

`USER_DEFAULT` SYMPHONY's default format should be used.

Invoked from: Whenever an LP solution is received.

Note:

SYMPHONY automatically unpacks the level, index and iteration number corresponding to the current LP solution within the current search tree node.

▷ **user_prepare_to_check_cuts**

```
int user_prepare_to_check_cuts(void *user, int varnum, int *indices,  
                               double *values)
```

Description:

This function is invoked after an LP solution is received but before any cuts are tested. Here the user can build up data structures (e.g., a graph representation of the solution) that can make the testing of cuts easier in the `user_check_cuts` function.

Arguments:

<code>void *user</code>	IN	Pointer to the user-defined data structure.
<code>int varnum</code>	IN	The number of nonzero/fractional variables described in <code>indices</code> and <code>values</code> .
<code>int *indices</code>	IN	The user indices of the nonzero/fractional variables.
<code>double *values</code>	IN	The nonzero/fractional values.

Return values:

<code>USER_ERROR</code>	Cuts are not checked for this LP solution.
<code>USER_SUCCESS</code>	The user is prepared to check cuts.
<code>USER_DEFAULT</code>	There are no user-defined cuts in the pool.

Invoked from: Whenever an LP solution is received.

▷ **user_check_cut**

```
int user_check_cut(void *user, double lpetol, int varnum,
                  int *indices, double *values, cut_data *cut,
                  int *is_violated, double *quality)
```

Description:

The user has to determine whether a given cut is violated by the given LP solution (see Section 6.3.2 for a description of the `cut_data` data structure). Also, the user can assign a number to the cut called the *quality*. This number is used in deciding which cuts to check and purge. See the section on Cut Pool Parameters for more information.

Arguments:

<code>void *user</code>	INOUT	The user defined part of p.
<code>double lpetol</code>	IN	The ϵ tolerance in the LP module.
<code>int varnum</code>	IN	Same as the previous function.
<code>int *indices</code>	IN	Same as the previous function.
<code>double *values</code>	IN	Same as the previous function.
<code>cut_data *cut</code>	IN	Pointer to the cut to be tested.
<code>int *is_violated</code>	OUT	TRUE/FALSE based on whether the cut is violated or not.
<code>double *quality</code>	OUT	a number representing the relative strength of the cut.

Return values:

<code>USER_ERROR</code>	Cut is not sent to the LP, regardless of the value of <code>*is_violated</code> .
<code>USER_SUCCESS</code>	The user function exited properly.
<code>USER_DEFAULT</code>	Same as error.

Invoked from: Whenever a cut needs to be checked.

Note:

The same note applies to `number`, `indices` and `values` as in the previous function.

▷ **user_finished_checking_cuts**

```
int user_finished_checking_cuts(void *user)
```

Description:

When this function is invoked there are no more cuts to be checked, so the user can dismantle data structures he created in `user_prepare_to_check_cuts`. Also, if he received and stored the LP solution himself he can delete it now.

Arguments:

void *user IN Pointer to the user-defined data structure.

Return values:

USER_ERROR Ignored.
USER_SUCCESS The user function exited properly.
USER_DEFAULT There are no user-defined cuts in the pool.

Invoked from: After all cuts have been checked.

▷ **user_free_cp**

```
int user_free_cp(void **user)
```

Description:

The user has to free all the data structures within **user**, and also free **user** itself. The user can use the built-in macro **FREE** that checks the existence of a pointer before freeing it.

Arguments:

<code>void **user</code>	INOUT	Pointer to the user-defined data structure (should be NULL on exit).
--------------------------	-------	--

Return values:

<code>USER_ERROR</code>	Ignored.
<code>USER_SUCCESS</code>	The user freed all data structures.
<code>USER_DEFAULT</code>	There is nothing to be freed.

Invoked from: `cp_close()` at module shutdown.

6.3.5 Draw graph module callbacks

Due to the relative simplicity of the cut pool, there are no wrapper functions implemented for DG. Consequently, there are no default options and no post-processing.

▷ **user_dg_process_message**

```
void user_dg_process_message(void *user, window *win, FILE *write_to)
```

Description:

The user has to process whatever user-defined messages are sent to the process. A write-to pipe to the wish process is provided so that the user can directly issue commands there.

Arguments:

<code>void *user</code>	INOUT	Pointer to the user-defined data structure.
<code>window *win</code>	INOUT	The window that received the message.
<code>FILE *write_to</code>	IN	Pipe to the wish process.

Return values:

<code>USER_ERROR</code>	Error. Message ignored.
<code>USER_SUCCESS</code>	The user processed the message.

▷ **user_dg_init_window**

```
void user_dg_init_window(void **user, window *win)
```

Description:

The user must perform whatever initialization is necessary for processing later commands. This usually includes setting up the user's data structure for receiving and storing display data.

Arguments:

<code>void **user</code>	INOUT	Pointer to the user-defined data structure.
<code>window *win</code>	INOUT	

Return values:

<code>USER_ERROR</code>	Error. Ignored.
<code>USER_SUCCESS</code>	The user successfully performed initialization.

▷ **user_dg_free_window**

```
void user_dg_free_window(void **user, window *win)
```

Description:

The user must free any data structures allocated.

Arguments:

<code>void **user</code>	INOUT	Pointer to the user-defined data structure.
<code>window *win</code>	INOUT	

Return values:

<code>USER_ERROR</code>	Error. Ignored.
<code>USER_SUCCESS</code>	The user successfully freed the data structures.

▷ **user_interpret_text**

```
void user_interpret_text(void *user, int text_length,  
    char *text, int owner_tid)
```

Description:

The user can interpret text input from the window.

Arguments:

void *user	INOUT	Pointer to the user-defined data structure.
int text_length	IN	The length of <code>text</code> .
char *text	IN	
int owner_tid	IN	The tid of the process that initiated this window.

Return values:

USER_ERROR	Error. Ignored.
USER_SUCCESS	The user successfully interpreted the text.

6.4 Run-time Parameters

Parameters can be set in one of two ways. Some commonly-used parameters can be set on the command line. To see a list of these, run SYMPHONY with no command-line arguments. Other parameters must be set in a parameter file. The name of this file is specified on the command line with “-f”. Each line of the parameter file contains either a comment or two words – a keyword and a value, separated by white space. If the first word (sequence of non-white-space characters) on a line is not a keyword, then the line is considered a comment line. Otherwise the parameter corresponding to the keyword is set to the listed value. Usually the keyword is the same as the parameter name in the source code. Here we list the keywords, the type of value that should be given with the keywords and the default value. A parameter corresponding to keyword “K” in module “P” can also be set by using the keyword “P_K”.

To make this list shorter, occasionally a comma separated list of parameters is given if the meanings of those parameters are strongly connected. For clarity, the constant name is sometimes given instead of the numerical value for default settings and options. The corresponding value is given in curly braces for convenience.

6.4.1 Global parameters

verbosity – **integer (0)**. Sets the verbosity of all modules to the given value. In general, the greater this number the more verbose each module is. Experiment to find out what this means.

random_seed – **integer (17)**. A random seed.

granularity – **double (1e-6)**. Should be set to “the minimum difference between two distinct objective function values” less the epsilon tolerance. E.g., if every variable is integral and the objective coefficients are integral then for any feasible solution the objective value is integer, so **granularity** could be correctly set to .99999.

upper_bound – **double (none)**. The value of the best known upper bound.

probrname – **string (empty string)**. The name of the problem name.

infile_name – **string (empty string)**. The name of the input file that was read by “-F” or the “-L” flag.

6.4.2 Master module parameters

M_verbosity – **integer (0)**.

M_random_seed – **integer (17)**. A random seed just for the Master module.

upper_bound – **double (no upper bound)**. This parameter is used if the user wants to artificially impose an upper bound (for instance if a solution of that value is already known).

lower_bound – **double (no lower bound)**. This parameter is used if the user wants to artificially impose a lower bound.

upper_bound_estimate – **double (no estimate)**. This parameter is used if the user wants to provide an estimate of the optimal value which will help guide the search. This is used in conjunction with the diving strategy **BEST_ESTIMATE**.

tm_exe, dg_exe – **strings (“tm”, “dg”)**. The name of the executable files of the TM and DG modules. Note that the TM executable name may have extensions that depend on the configuration of the modules, but the default is always set to the file name produced by the makefile. If you change the name of the treemanager executable from the default, you must set this parameter to the new name.

tm_debug, dg_debug – **boolean (both FALSE)**. Whether these modules should be started under a debugger or not (see 5.6.2 for more details on this).

tm_machine – **string (empty string)**. On which processor of the virtual machine the TM should be run. Leaving this parameter as an empty string means arbitrary selection.

do_draw_graph – **boolean (FALSE)**. Whether to start up the DG module or not (see Section 5.6.4 for an introduction to this).

do_branch_and_cut – **boolean (TRUE)**. Whether to run the branch and cut algorithm or not. (Set this to **FALSE** to run the user’s heuristics only.)

mc_search_order – **integer (MC_FIFO)**. Use the fifo (**MC_FIFO**) or lifo (**MC_LIFO**) search order during the multi criteria solution procedure.

mc_warm_start – **boolean (FALSE)**. Whether to solve the corresponding problem of each iteration from a warm start loaded from a base iteration (which is the first iteration where $\gamma = 1.0$ and $\tau = 0.0$) or from scratch. Currently, this option is supported if only the supported solutions are desired to be found.

trim_warm_tree – **boolean (FALSE)**. Whether to trim the warm start tree before re-solving. This consists of locating nodes whose descendants are all likely to be pruned in the resolve and eliminating those descendants in favor of processing the parent node itself.

mc_compare_solution_tolerance – **double (0.001)**. If the difference between the objective values of two solutions to be compared, during the bicriteria solution procedure, are less than this tolerance, then assume them to be equal.

mc_binary_search_tolerance – **double (0)**. The tolerance to be used to differentiate the gamma values if binary search is used during the bicriteria solution procedure. A value greater than zero will cause the binary search to be activated.

prep_level – **integer (5)**. Determines the level of preprocessing that should be done on the current MILP instance. A level of less than 0 means that no preprocessing will be done. At level 2 basic presolve routines are used. At higher levels more advanced routines are deployed. At level 5, valid implications are derived.

prep_dive_level – **integer**(5). When a variable has been modified by preprocessing, then these changes can be used to improve other variables and constraints in the instance as well. This parameter controls how many times can we recursively try to improve the instance if a change is made.

prep_impl_dive_level – **integer**(0). In some advanced preprocessing routines, a variable or constraint is modified to check what implications can be derived from that change. When such an implication is derived, it can recursively lead to more implications. This parameter controls how many levels of recursion are allowed.

prep_impl_limit – **integer**(50). Determines the maximum number of implications that can be derived from preprocessing.

prep_do_probing – **integer**(1). Determines if probing is used while preprocessing. Probing is not yet implemented and this parameter does not have any effect.

prep_verbosity – **integer**(1). Determines the verbosity of messages from the preprocessing stage. Higher levels will produce more verbose messages.

prep_reduce_mip – **boolean** (1). If some variables and constraints have been eliminated in preprocessing and if **prep_reduce_mip** is 1, then the memory allocated for these deleted variables and constraints is freed. Otherwise, these are retained in the instance but are never used.

prep_probing_verbosity – **integer**(0). Determines the verbosity of messages from probing stage. Probing is not yet implemented and this parameter does not have any effect.

prep_probing_level – **integer**(1). Determines the maximum level of probing that is carried out before preprocessing is stopped. Probing is not yet implemented and this parameter does not have any effect.

prep_display_stats – **boolean** (0). Determines if statistics on how many of each type of changes were made in the preprocessing stage are displayed (1) or not (0).

keep_row_ordered – **integer**(1). When the value of this parameter is 1, a row ordered matrix is also retained for use after the preprocessing stage. This capability is not yet implemented and this parameter does not have any effect.

prep_do_sr – **boolean** (0). When the value of this parameter is 1, additional preprocessing is performed by solving an LP with one constraint. This procedure is not thoroughly tested.

max_sr_cnt – **integer**(5). This parameter controls the number of single-constraint LPs that are solved for each constraint in the preprocessing stage. This procedure is not thoroughly tested.

max_aggr_row_cnt – **integer**(0). This parameter is not used and has no effect.

prep_iter_limit – **integer**(10). Determines the maximum number of times preprocessing can be done on an instance. If an instance has been modified by preprocessing, then the new problem can be preprocessed again to get an even better formulation. This parameter puts a limit on the number of times such preprocessing can be done.

`write_mps` – **boolean** (0). Determines if an MPS file be written after all preprocessing has been performed. This can be used for debugging or if the user wants to save the preprocessed instance.

`write_lp` – **boolean** (0). Determines if an LP file be written after all preprocessing has been performed. This can be used for debugging or if the user wants to save the preprocessed instance.

`prep_time_limit` – **integer**(50). Determines the maximum time in seconds that can be spent in preprocessing.

6.4.3 Draw Graph parameters

`source_path` – **string** (“.”). The directory where the DG tcl/tk scripts reside.

`echo_commands` – **boolean** (FALSE). Whether to echo the tcl/tk commands on the screen or not.

`canvas_width`, `canvas_height` – **integers** (1000, 700). The default width and height of the drawing canvas in pixels.

`viewable_width`, `viewable_height` – **integers** (600, 400). The default viewable width and height of the drawing canvas in pixels.

`interactive_mode` – **integer** (TRUE). Whether it is allowable to change things interactively on the canvas or not.

`node_radius` – **integer** (8). The default radius of a displayed graph node.

`disp_nodelabels`, `disp_nodeweights`, `disp_edgeweights` – **integers** (all TRUE). Whether to display node labels, node weights, and edge weights or not.

`nodelabel_font`, `nodeweight_font`, `edgeweight_font` – **strings** (all “-adobe-helvetica-...”). The default character font for displaying node labels, node weights and edge weights.

`node_dash`, `edge_dash` – **strings** (both empty string). The dash pattern of the circles drawn around dashed nodes and that of dashed edges.

6.4.4 Tree Manager parameters

`TM_verbosity` – **integer** (0). The verbosity of the TM module.

`lp_exe`, `cg_exe`, `cp_exe` – **strings** (“lp”, “cg”, “cp”). The name of the LP, CG, and CP module binaries. Note: when running in parallel using PVM, these executables (or links to them) must reside in the `PVM_ROOT/bin/PVM_ARCH/` directory. Also, be sure to note that the executable names may have extensions that depend on the configuration of the modules, but the defaults will always be set to the name that the makefile produces.

`lp_debug`, `cg_debug`, `cp_debug` – **boolean** (all FALSE). Whether the modules should be started under a debugger or not.

`max_active_nodes` – **integer (1)**. The maximum number of processors/threads to be used.

`max_cp_num` – **integer (0)**. The maximum number of cut pools to be used.

`lp_mach_num`, `cg_mach_num`, `cp_mach_num` – **integers (all 0)**. The number of processors in the virtual machine to run LP (CG, CP) processes. If this value is 0 then the processes will be assigned to processors in round-robin order. Otherwise the next `xx_mach_num` lines describe the processors where the LP (CG, CP) modules must run. The keyword – value pairs on these lines must be **TM_**`xx_machine` and the name or IP address of a processor (the processor names need not be distinct). In this case the actual processes are assigned in a round robin fashion to the processors on this list.

This feature is useful if a specific software package is needed for some module, but that software is not licensed for every node of the virtual machine or if a certain process must run on a certain type of machine due to resource requirements.

`use_cg` – **boolean (FALSE)**. Whether to use a cut generator or not.

`TM_random_seed` – **integer (17)**. The random seed used in the TM.

`unconditional_dive_frac` – **double (0.0)**. The fraction of the nodes on which SYMPHONY randomly dives unconditionally into one of the children.

`diving_strategy` – **integer (BEST_ESTIMATE{0})**. The strategy employed when deciding whether to dive or not.

The `BEST_ESTIMATE{0}` strategy continues to dive until the lower bound in the child to be dived into exceeds the parameter `upper_bound_estimate`, which is given by the user.

The `COMP_BEST_K{1}` strategy computes the average lower bound on the best `diving_k` search tree nodes and decides to dive if the lower bound of the child to be dived into does not exceed this average by more than the fraction `diving_threshold`.

The `COMP_BEST_K.GAP{2}` strategy takes the size of the gap into account when deciding whether to dive. After the average lower bound of the best `diving_k` nodes is computed, the gap between this average lower bound and the current upper bound is computed. Diving only occurs if the difference between the computed average lower bound and the lower bound of the child to be dived into is at most the fraction `diving_threshold` of the gap.

Note that fractional diving settings can override these strategies. See below.

`diving_k`, `diving_threshold` – **integer, double (1, 0.05)**. See above.

`fractional_diving_ratio`, `fractional_diving_num` – **integer (0.02, 0)**. Diving occurs automatically if the number of fractional variables in the child to be dived into is less than `fractional_diving_num` or the fraction of total variables that are fractional is less than `fractional_diving_ratio`. This overrides the other diving rules. Note that in order for this option to work, the code must be compiled with `FRACTIONAL_BRANCHING` defined. This is the default. See the makefile for more details.

node_selection_rule – integer (LOWEST_LP_FIRST{0}). The rule for selecting the next search tree node to be processed. This rule selects the one with lowest lower bound. Other possible values are: HIGHEST_LP_FIRST{1}, BREADTH_FIRST_SEARCH{2} and DEPTH_FIRST_SEARCH{3}.

load_balance_level – integer (-1).] A naive attempt at load balancing on problems where significant time is spent in the root node, contributing to a lack of parallel speed-up. Only a prescribed number of iterations (**load_balance_iter**) are performed in the root node (and in each subsequent node on a level less than or equal to **load_balance_level**) before branching is forced in order to provide additional subproblems for the idle processors to work on. This doesn't work well in general.

load_balance_iter – integer (-1).] Works in tandem with the **load_balance_level** to attempt some simple load balancing. See the above description.

tighten_root_bounds – boolean (TRUE).] Whether to tighten the global bounds on variables using the reduced costs from the root node.

status_interval – integer (5).] How often to print status messages to the screen.

keep_description_of_pruned – integer (DISCARD{0}). Whether to keep the description of pruned search tree nodes or not. The reasons to do this are (1) if the user wants to write out a proof of optimality using the logging function, (2) for debugging, or (3) to get a visual picture of the tree using the software VBCTOOL. Otherwise, keeping the pruned nodes around just takes up memory.

There are three options if it is desired to keep some description of the pruned nodes around. First, their full description can be written out to disk and freed from memory (KEEP_ON_DISK_FULL{1}). There is not really too much you can do with this kind of file, but theoretically, it contains a full record of the solution process and could be used to provide a certificate of optimality (if we were using exact arithmetic) using an independent verifier. In this case, the line following **keep_description_of_pruned** should be a line containing the keyword **pruned_node_file_name** with its corresponding value being the name of a file to which a description of the pruned nodes can be written. The file does not need to exist and will be over-written if it does exist.

If you have the software VBCTOOL, then you can alternatively just write out the information VBCTOOL needs to display the tree (KEEP_ON_DISK_VBC_TOOL{2}).

Finally, the user can set the value to of this parameter to KEEP_IN_MEMORY{2}, in which case all pruned nodes will be kept in memory and written out to the regular log file if that option is chosen. This is really only useful for debugging. Otherwise, pruned nodes should be flushed.

keep_warm_start – boolean (FALSE). Turning this parameter on will have exactly the same impact with setting the **keep_description_of_pruned** to KEEP_IN_MEMORY{2}. This will allow SYMPHONY to keep all the necessary information obtained from the branching tree of the original problem to be able to warm start after a parameter or problem data modification. Thus, if it is intended to warm start later, the user should set this parameter before solving the original problem.

warm_start_node_limit – **integer** (SYM_INFINITY). Setting this parameter will start the warm start routine using only the first **warm_start_node_limit** nodes generated during the previous solve procedure. The rest of the tree will be trimmed.

warm_start_node_ratio – **double** (0.0). Setting this parameter will start the warm start routine using only the first **warm_start_node_ratio**% of the nodes generated during the previous solve procedure.

warm_start_node_level – **integer** (SYM_INFINITY). Setting this parameter will start the warm start routine using all the nodes above the level **warm_start_node_level** of the tree generated during the previous solve procedure. The rest of the tree will be trimmed.

warm_start_node_level_ratio – **double** (0.0). Setting this parameter will start the warm start routine using all the nodes above the level **warm_start_node_level**% of the warm start tree depth. The rest of the tree will be trimmed

logging – **integer** (NO_LOGGING{0}). Whether or not to write out the state of the search tree and all other necessary data to disk periodically in order to allow a warm start in the case of a system crash or to allow periodic viewing with VBCTOOL.

If the value of this parameter is set to FULL_LOGGING{1}, then all information needed to warm start the calculation will be written out periodically. The next two lines of the parameter file following should contain the keywords **tree_log_file_name** and **cut_log_file_name** along with corresponding file names as values. These will be the files used to record the search tree and related data and the list of cuts needed to reconstruct the tree.

If the value of the parameter is set to VBC_TOOL{2}, then only the information VBCTOOL needs to display the tree will be logged. This is not really a very useful option since a “live” picture of the tree can be obtained using the **vbc_emulation** parameter described below.

logging_interval – **integer** (1800). Interval (in seconds) between writing out the above log files.

warm_start – **boolean** (0). Used to allow the tree manager to make a warm start by reading in previously written log files. If this option is set, then the two lines following must start with the keywords **warm_start_tree_file_name** and **warm_start_cut_file_name** and include the appropriate file names as the corresponding values.

vbc_emulation – **integer** (NO_VBC_EMULATION{0}). Determines whether or not to employ the VBC-TOOL emulation mode. If one of these modes is chosen, then the tree will be displayed in “real time” using the VBCTOOL Software. When using the option VBC_EMULATION_LIVE{2} and piping the output directly to VBCTOOL, the tree will be displayed as it is constructed, with color coding indicating the status of each node. With VBC_EMULATION_FILE{1} selected, a log file will be produced which can later be read into VBCTOOL to produce an emulation of the solution process at any desired speed. If VBC_EMULATION_FILE is selected, the following line should contain the keyword **vbc_emulation_file_name** along with the corresponding file name for a value.

price_in_root – **boolean** (FALSE). Whether to price out variables in the root node before the second phase starts (called *repricing the root*).

trim_search_tree – **boolean (FALSE)**. Whether to trim the search tree before the second phase starts or not. Useful only if there are two phases. (It is very useful then.)

colgen_in_first_phase, colgen_in_second_phase – **integers (both 4)**. These parameters determine if and when to do column generation in the first and second phase of the algorithm. The value of each parameter is obtained by setting the last four bits. The last two bits refer to what to do when attempting to prune a node. If neither of the last two bits are set, then we don't do anything—we just prune it. If only the last bit is set, then we simply save the node for the second phase without doing any column generation (yet). If only the second to last bit is set, then we do column generation immediately and resolve if any new columns are found. The next two higher bits determine whether or not to do column generation before branching. If only the third lowest bit is set, then no column generation occurs before branching. If only the fourth lowest bit is set, then column generation is attempted before branching. The default is not to generate columns before branching or fathoming, which corresponds to only the third lowest bit being set, resulting in a default value of 4.

time_limit – **double (-1.0)**. Number of seconds of wall-clock time allowed for solution. When this time limit is reached, the solution process will stop and the best solution found to that point, along with other relevant data, will be output. A time limit less than 0.0 means there is no limit.

node_limit – **integer (-1)**. Number of nodes allowed to be analyzed during the solution. When this node limit is reached, the solution process will stop and the best solution found to that point, along with other relevant data, will be output. A node limit less than 0 means there is no limit.

gap_limit – **double (-1.0)**. Target gap limit used to determine when the solver should terminate. When the gap between the lower and the upper bound reaches this point, the solution process will stop and the best solution found to that point, along with other relevant data, will be output. This parameter expresses the target gap as a percentage relative to the current lower bound. For example, a gap limit equal to 10 means that the solution process will stop when the best solution is within 10% of the current lower bound (which is a proxy for the optimal value itself). Additionally, a gap limit less than or equal to 0 means the solution produced is the exact optimum.

find_first_feasible – **boolean (FALSE)**. Whether to stop after finding the first feasible solution or not.

sensitivity_analysis – **boolean (FALSE)**. If the user wants to do the rudimentary sensitivity analysis, which will give a lower bound for the problem modified by the right hand side, then, this parameter has to be set before solving the original problem. If it is set, SYMPHONY will keep the necessary information from the solution processes of the original problem to be able to do the sensitivity analysis later.

6.4.5 LP parameters

LP_verbosity – **integer (0)**. Verbosity level of the LP module.

set_obj_upper_lim – **boolean** (FALSE). Whether to stop solving the LP relaxation when it's optimal value is provably higher than the global upper bound. There are some advantages to continuing the solution process anyway. For instance, this results in the highest possible lower bound. On the other hand, if the matrix is full, this node will be pruned anyway and the rest of the computation is pointless. This option should be set at **FALSE** for column generation since the LP dual values may not be reliable otherwise.

try_to_recover_from_error – **boolean** (TRUE). Indicates what should be done in case the LP solver is unable to solve a particular LP relaxation because of numerical problems. It is possible to recover from this situation but further results may be suspect. On the other hand, the entire solution process can be abandoned.

problem_type – **integer** (ZERO_ONE_PROBLEM{0}). The type of problem being solved. Other values are INTEGER_PROBLEM{1} or MIXED_INTEGER_PROBLEM{2}. (Caution: The mixed-integer option is not well tested.)

cut_pool_check_frequency – **integer** (10). The number of iterations between sending LP solutions to the cut pool to find violated cuts. It is not advisable to check the cut pool too frequently as the cut pool module can get bogged down and the LP solution generally do not change that drastically from one iteration to the next anyway.

not_fixed_storage_size – **integer** (2048). The *not fixed list* is a partial list of indices of variables not in the matrix that have not been fixed by reduced cost. Keeping this list allows SYMPHONY to avoid repricing variables (an expensive operation) that are not in the matrix because they have already been permanently fixed. When this array reaches its maximum size, no more variable indices can be stored. It is therefore advisable to keep the maximum size of this array as large as possible, given memory limitations.

max_non_dual_feas_to_add_min,

max_non_dual_feas_to_add_max,

max_non_dual_feas_to_add_frac – **integer, integer, double** (20, 200, .05). These three parameters determine the maximum number of non-dual-feasible columns that can be added in any one iteration after pricing. This maximum is set to the indicated fraction of the current number of active columns unless this number exceeds the given maximum or is less than the given minimum, in which case, it is set to the max or min, respectively.

max_not_fixable_to_add_min,

max_not_fixable_to_add_max,

max_not_fixable_to_add_frac – **integer, integer, double** (100, 500, .1) As above, these three parameters determine the maximum number of new columns to be added to the problem because they cannot be priced out. These variables are only added when trying to restore infeasibility and usually, this does not require many variables anyway.

mat_col_compress_num, mat_col_compress_ratio – **integer, double** (50, .05). Determines when the matrix should be physically compressed. This only happens when the number of columns is high enough to make it “worthwhile.” The matrix is physically compressed when

the number of deleted columns exceeds either an absolute number *and* a specified fraction of the current number of active columns.

`mat_row_compress_num`, `mat_row_compress_ratio` – **integer**, **double** (**20**, **.05**). Same as above except for rows.

`tailoff_gap_backsteps`, `tailoff_gap_frac` – **integer**, **double** (**2**, **.99**). Determines when tailoff is detected in the LP module. Tailoff is reported if the average ratio of the current gap to the previous iteration's gap over the last `tailoff_gap_backsteps` iterations wasn't at least `tailoff_gap_frac`.

`tailoff_obj_backsteps`, `tailoff_obj_frac` – **integer**, **double** (**2**, **.99**). Same as above, only the ratio is taken with respect to the change in objective function values instead of the change in the gap.

`ineff_cnt_to_delete` – **integer** (**0**). Determines after how many iterations of being deemed ineffective a constraint is removed from the current relaxation.

`eff_cnt_before_cutpool` – **integer** (**3**). Determines after how many iterations of being deemed effective each cut will be sent to the global pool.

`ineffective_constraints` – **integer** (`BASIC_SLACKS_ARE_INEFFECTIVE{2}`). Determines under what condition a constraint is deemed ineffective in the current relaxation. Other possible values are `NO_CONSTRAINT_IS_INEFFECTIVE{0}`, `NONZERO_SLACKS_ARE_INEFFECTIVE{1}`, and `ZERO_DUAL_VALUES_ARE_INEFFECTIVE{3}`.

`base_constraints_always_effective` – **boolean** (**TRUE**). Determines whether the base constraints can ever be removed from the relaxation. In some case, removing the base constraints from the problem can be disastrous depending on the assumptions made by the cut generator.

`branch_on_cuts` – **boolean** (**FALSE**). This informs the framework whether the user plans on branching on cuts or not. If so, there is additional bookkeeping to be done, such as maintaining a pool of slack cuts to be used for branching. Therefore, the user should not set this flag unless he actually plans on using this feature.

`discard_slack_cuts` – **integer** (`DISCARD_SLACKS_BEFORE_NEW_ITERATION{0}`).

Determines when the pool of slack cuts is discarded. The other option is `DISCARD_SLACKS_WHEN_STARTING_NEW_NODE{1}`.

`first_lp_first_cut_time_out`,

`first_lp_all_cuts_time_out`,

`later_lp_first_cut_time_out`,

`later_lp_all_cuts_time_out` – **double** (**0**, **0**, **5**, **1**). The next group of parameters determines when the LP should give up waiting for cuts from the cut generator and start to solve the relaxation in its current form or possibly branch if necessary. There are two factors that contribute to determining this timeout. First is whether this is the first LP in the search node or whether it is a later LP. Second is whether any cuts have been added already in this

iteration. The four timeout parameters correspond to the four possible combinations of these two variables.

`no_cut_timeout` – This keyword does not have an associated value. If this keyword appears on a line by itself or with a value, this tells the framework not to time out while waiting for cuts. This is useful for debugging since it enables runs with a single LP module to be duplicated.

`all_cut_timeout` – **double (no default)**. This keyword tells the framework to set all of the above timeout parameters to the value indicated.

`max_cut_num_per_iter` – **integer (20)**. The maximum number of cuts that can be added to the LP in an iteration. The remaining cuts stay in the local pool to be added in subsequent iterations, if they are strong enough.

`do_reduced_cost_fixing` – **boolean (FALSE)**. Whether or not to attempt to fix variables by reduced cost. This option is highly recommended

`gap_as_ub_frac`, `gap_as_last_gap_frac` – **double (.1, .7)**. Determines when reduced cost fixing should be attempted. It is only done when the gap is within the fraction `gap_as_ub_frac` of the upper bound or when the gap has decreased by the fraction `gap_as_last_gap_frac` since the last time variables were fixed.

`do_logical_fixing` – **boolean (FALSE)**. Determines whether the user's logical fixing routine should be used.

`fixed_to_ub_before_logical_fixing`,

`fixed_to_ub_frac_before_logical_fixing` – **integer, double (1, .01)**. Determines when logical fixing should be attempted. It will be called only when a certain absolute number *and* a certain number of variables have been fixed to their upper bounds by reduced cost. This is because it is typically only after fixing variables to their upper bound that other variables can be logically fixed.

`max_presolve_iter` – **integer (10)**. Number of simplex iterations to be performed in the pre-solve for strong branching.

`strong_branching_cand_num_max`,

`strong_branching_cand_num_min`,

`strong_branching_red_ratio` – **integer (10, 5, 1)**. These three parameters together determine the number of strong branching candidates to be used by default. In the root node, `strong_branching_cand_num_max` candidates are used. On each succeeding level, this number is reduced by the number `strong_branching_red_ratio` multiplied by the square of the level. This continues until the number of candidates is reduced to `strong_branching_cand_num_min` and then that number of candidates is used in all lower levels of the tree.

`strong_branching_high_low_weight` – **double (0.8)**. This parameter is used to calculate the score of each branching candidate. The candidate with the highest score is then selected for branching. Let z_i^+ , z_i^- be the estimated change in objective function value when we branch on

the candidate i . Then the score of candidate i is $s_i = \alpha \times \min\{z_i^+, z_i^-\} + (1 - \alpha) \times \max\{z_i^+, z_i^-\}$, where α is the value of `strong_branching_high_low_weight`. This value should always lie in the interval $[0, 1]$.

`limit_strong_branching_time` – **boolean** (TRUE).] Whether to heuristically limit the amount of time spent in strong branching.

`use_hot_starts` – **boolean** (TRUE). Determines if the LP solver is asked to make special arrangements for doing dual-simplex iterations when bounds on a variable are changed for strong branching. Some LP solvers provide such options so that strong branching can be performed much faster than the regular dual-simplex procedure.

`should_use_rel_br` – **boolean** (TRUE). Determines if reliability braching is used to determine branching candidates or not. This parameter is set to `FALSE` if OPENMP is used. When this branching rule is disabled, strong branching is used to select a candidate.

`rel_br_override_default` – **boolean** (TRUE). If reliability branching is enabled and this paramter is set to `TRUE` then the policy of selecting branching candidates is automatically adjusted on the basis of bounds on solution value and the time elapsed. If this parameter is set to `FALSE`, the policy is based on the values of the following three parameters.

`rel_br_threshold` – **integer** (8). It is assumed that the score obtained by branching on a given variable these many times is reliable for estimating the pseudocosts of this variable in the rest of the branch-and-bound algorithm. In other words, if reliability branching is enabled, strong branching is used on a variable at most `rel_br_threshold` many times.

`rel_br_max_solves` – **integer** (20). If reliability branching is enabled, this parameter determines the maximum number of strong branching LPs that are solved in each node. If some branching candidates have reliable estimates, the number of LPs can be less than the value of this parameter.

`rel_br_cand_threshold` – **integer** (10). If reliability branching is enabled, then strong branching is stopped if the last `rel_br_cand_threshold` LPs did not give a better improvement in the lower bound.

`is_feasible_default` – **integer** (TEST_INTEGRALITY{1}). Determines the default test to be used to determine feasibility. This parameter is provided so that the user can change the default behavior without recompiling. The only other option is `TEST_ZERO_ONE{0}`.

`send_feasible_solution_default` – **integer** (SEND_NONZEROS{0}). Determines the form in which to send the feasible solution. This parameter is provided so that the user can change the default behavior without recompiling. This is currently the only option.

`send_lp_solution_default` – **integer** (SEND_NONZEROS{0}). Determines the default form in which to send the LP solution to the cut generator and cut pool. This parameter is provided so that the user can change the default behavior without recompiling. The other option is `SEND_FRACTIONS{1}`.

`display_solution_default` – **integer** (DISP_NOTHING{0}). Determines how to display the current LP solution if desired. See the description of `user_display_solution()` for other possible

values. This parameter is provided so that the user can change the default behavior without recompiling.

shall_we_branch_default – **integer** (USER_BRANCH_IF_MUST{2}). Determines the default branching behavior. Other values are USER_DO_NOT_BRANCH{0} (not recommended as a default), USER_DO_BRANCH{1} (also not recommended as a default), and USER_BRANCH_IF_TAILOFF{3}. This parameter is provided so that the user can change the default behavior without recompiling.

select_candidates_default – **integer** (USER_CLOSE_TO_HALF_AND_EXPENSIVE{10}). Determines the default rule for selecting strong branching candidates. Other values are USER_CLOSE_TO_HALF{10} and USER_CLOSE_TO_ONE_AND_CHEAP{12}. This parameter is provided so that the user can change the default behavior without recompiling.

compare_candidates_default – **integer** (HIGHEST_LOW_OBJ{2}). Determines the default rule for comparing candidates. See the description of `user_compare_candidates()` for other values. This parameter is provided so that the user can change the default behavior without recompiling.

select_child_default – **integer** (PREFER_LOWER_OBJ_VALUE{0}). Determines the default rule for selecting the child to be processed next. For other possible values, see the description `user_select_child()`. This parameter is provided so that the user can change the default behavior without recompiling.

mc_find_supported_solutions – **boolean** (FALSE). By default, `sym_mc_solve` routine will find all the non-dominated solutions if the problem to be solved is a bicriteria problem. However, if the user plans to find only the supported solutions, then, this parameter has to be set before calling `sym_mc_solve` routine.

mc_rho – **double** (0.00001). The value used in augmented Chebyshev norm during the bicriteria solution procedure.

generate_cgl_cuts – **boolean** (TRUE). Whether or not to generate cuts using COIN's cut generation library. Note that, to use CGL cuts, OSI interface has to be used and moreover the corresponding flags have to be set during installation. See the makefile for more details.

generate_cgl_gomory_cuts – **integer** (GENERATE_DEFAULT).

generate_cgl_knapsack_cuts – **integer** (GENERATE_DEFAULT).

generate_cgl_twomir_cuts – **integer** (GENERATE_ONLY_IN_ROOT).

generate_cgl_clique_cuts – **integer** (GENERATE_DEFAULT).

generate_cgl_oddhole_cuts – **integer** (DO_NOT_GENERATE).

generate_cgl_flowcover_cuts – **integer** (GENERATE_DEFAULT).

This set of parameters is for controlling when certain classes of cuts are to be generated. It can take on the values enumerated below (see values in `symphony.h`).

- `DO_NOT_GENERATE` (-1): Do not generate cuts of this class.
- `GENERATE_DEFAULT` (0): Use SYMPHONY's automatic tuning for generation of this class of cuts.
- `GENERATE_IF_IN_ROOT` (1): Generate these cuts in the tree only if at least one violated cut is generated in the root node.
- `GENERATE_ONLY_IN_ROOT` (2): Generate these cuts only in the root node.
- `GENERATE_ALWAYS` (3): Generate these cuts every time cut generation is done.
- `GENERATE_PERIODICALLY` (4): Generate these cuts every n^{th} time the cut generator is called, where n is set by the parameter `generate_cgl_xxx_cuts_freq`.

See <https://projects.coin-or.org/Cgl> for a description of these classes and others that could be added.

`generate_cgl_gomory_cuts_freq` – integer (5).

`generate_cgl_knapsack_cuts_freq` – integer (5).

`generate_cgl_twomir_cuts_freq` – integer (5).

`generate_cgl_clique_cuts_freq` – integer (5).

`generate_cgl_oddhole_cuts_freq` – integer (5).

`generate_cgl_flowcover_cuts_freq` – integer (5).

The frequency with which to generate cuts if the parameter `cgl.generate_xxx_cuts` for the corresponding class is set to `GENERATE_PERIODICALLY`.

`fp_enabled` – integer (`SYM_FEAS_PUMP_DEFAULT{1}`). Determines the overall policy of using the feasibility pump heuristic to find feasible solutions. `SYM_FEAS_PUMP_DEFAULT{1}` indicates that the decision to use the heuristic is determined on the basis of current values of lower bound, upper bound, the time used etc., based on some preset rules. `SYM_FEAS_PUMP_REPEATED{2}` indicates that the heuristic will be used every few iterations until the problem is solved. The frequency can be adjusted through the `fp_frequency` parameter. `SYM_FEAS_PUMP_TILL_SOL{3}` indicates that the heuristic is used only until the first feasible solution is found. `SYM_FEAS_PUMP_DISABLE{-1}` indicates that the heuristic is not used.

`fp_frequency` – integer (10). Determines the number of LPs that are solved before which the feasibility pump heuristic is called again. This parameter is used only if the parameter `fp_enabled` is set to `SYM_FEAS_PUMP_REPEATED{2}`. Otherwise, the frequency is determined automatically based on some preset rules.

`fp_max_cycles` – integer (100). Determines the maximum number of LPs that can be solved in a call to the feasibility pump heuristic. A higher number might be helpful in finding a better feasible solution but may result in more time spent in the heuristic.

fp_time_limit – **double (50)**. If a feasible solution has been found, this parameter determines the time in seconds that can be spent on the feasibility pump heuristic. If a solution has not been found yet, the parameter **fp_max_initial_time** is used.

fp_max_initial_time – **double (100)**. If a feasible solution has not been found, this parameter determines the time in seconds that can be spent on the feasibility pump heuristic. If a solution has been found, the parameter **fp_time_limit** is used.

fp_min_gap – **double (0.5)**. If the relative (%) gap between the lower and the upper bounds falls below the value of this parameter, feasibility pump is not called.

fp_flip_fraction – **double (0.1)**. When the feasibility pump gets stuck in a cycle, this fraction of binary variables are flipped. The variables are selected randomly. Increasing the value of this parameter may result in the pump getting stuck fewer number of times, but the time to solve LPs after flipping may increase substantially.

fp_poor_sol_lim_fac – **integer (10)**. Sometimes the feasibility pump keeps generating solutions that have high objective function values. When the number of such solutions becomes more than **fp_poor_sol_lim_fac** times the number of “good” solutions, the pump is disabled.

6.4.6 Cut Generator Parameters

CG_verbosity – **integer (0)**. Verbosity level for the cut generator module.

6.4.7 Cut Pool Parameters

CP_verbosity – **integer (0)**. Verbosity of the cut pool module.

cp_logging – **boolean (0)**. Determines whether the logging option is enabled. In this case, the entire contents of the cut pool are written out periodically to disk (at the same interval as the tree manager log files are written). If this option is set, then the line following must start with the keyword **cp_log_file_name** and include the appropriate file name as the value.

cp_warm_start – **boolean (0)**. Used to allow the cut pool to make a warm start by reading in a previously written log file. If this option is set, then the line following must start with the keyword **cp_warm_start_file_name** and include the appropriate file name as the value.

block_size – **integer (5000)**. Indicates the size of the blocks to allocate when more space is needed in the cut list.

max_size – **integer (2000000)**. Indicates the maximum size of the cut pool in bytes. This is the total memory taken up by the cut list, including all data structures and the array of pointers itself.

max_number_of_cuts – **integer (10000)**. Indicates the maximum number of cuts allowed to be stored. When this max is reached, cuts are forcibly purged, starting with duplicates and then those indicated by the parameter **delete_which** (see below), until the list is below the allowable size.

min_to_delete – **integer (1000)**. Indicates the number of cuts required to be deleted when the pool reaches it's maximum size.

touches_until_deletion – **integer (10)**. When using the number of touches a cut has as a measure of its quality, this parameter indicates the number of touches a cut can have before being deleted from the pool. The number of touches is the number of times in a row that a cut has been checked without being found to be violated. It is a measure of a cut's relevance or effectiveness.

delete_which – **integer (DELETE_BY_TOUCHES{2})**. Indicates which cuts to delete when purging the pool. `DELETE_BY_TOUCHES` indicates that cuts whose number of touches is above the threshold (see **touches_until_deletion** above) should be purged if the pool gets too large. `DELETE_BY_QUALITY{1}` indicates that a user-defined measure of quality should be used (see the function `user_check_cuts` in Section 6.3.4).

check_which – **integer (CHECK_ALL_CUTS{0})**. Indicates which cuts should be checked for violation. The choices are to check all cuts (`CHECK_ALL_CUTS{0}`); only those that have number of touches below the threshold (`CHECK_TOUCHES{2}`); only those that were generated at a level higher in the tree than the current one (`CHECK_LEVEL{1}`); or both (`CHECK_LEVEL_AND_TOUCHES{3}`). Note that with `CHECK_ALL_CUTS` set, SYMPHONY will still only check the first **cuts_to_check** cuts in the list ordered by quality (see the function `user_check_cut`).

cuts_to_check – **integer (1000)**. Indicates how many cuts in the pool to actually check. The list is ordered by quality and the first **cuts_to_check** cuts are checked for violation.

6.4.8 C++ Interface/OSI Parameters

As the implementation of the whole interface, there exists a matching C interface parameter to each of the C++ Interface/OSI parameter and the parameter setting functions are designed to set the corresponding C interface parameter. Thus, we will just give a table of the parameter names, their C interface complements and the values they can be set to, rather than their detailed descriptions. For each parameter, the user can see the C interface complement for further explanation.

C++ Interface	C Interface	Value
OsiSymVerbosity	verbosity	-user defined-
OsiSymWarmStart	warm_start	-boolean-
OsiSymNodeLimit OsiMaxNumIteration OsiMaxNumIterationHotStart	node_limit	-user defined-
OsiSymFindFirstFeasible	find_first_feasible	-boolean-
OsiSymSearchStrategy	node_selection_rule	LOWEST_LP_FIRST HIGHEST_LP_FIRST BREADTH_FIRST_SEARCH DEPTH_FIRST_SEARCH
OsiSymUsePermanentCutPools	use_permanent_cut_pools	-boolean-
OsiSymGenerateCglGomoryCuts	generate_cgl_gomory_cuts	-boolean-
OsiSymGenerateCglKnapsackCuts	generate_cgl_knapsack_cuts	-boolean-
OsiSymGenerateCglOddHoleCuts	generate_cgl_oddhole_cuts	-boolean-
OsiSymGenerateCglProbingCuts	generate_cgl_probing_cuts	-boolean-
OsiSymGenerateCglCliqueCuts	generate_cgl_clique_cuts	-boolean-
OsiSymGenerateCglFlowAndCoverCuts	generate_cgl_flow_and_cover_cuts	-boolean-
OsiSymGenerateCglRoundingCuts	generate_cgl_rounding_cuts	-boolean-
OsiSymGenerateCglLiftAndProjectCuts	generate_cgl_lift_and_project_cuts	-boolean-
OsiSymKeepWarmStart	keep_warm_start	-boolean-
OsiSymTrimWarmTree	trim_warm_tree * -boolean-	
OsiSymDoReducedCostFixing	do_reduced_cost_fixing	-boolean-
OsiSymMCFindSupportedSolutions	mc_find_supported_solutions	-boolean-
OsiSymSensitivityAnalysis	sensitivity_analysis	-boolean-
OsiSymRandomSeed	random_seed	-user defined-
OsiSymDivingStrategy	diving_strategy	BEST_ESTIMATE COMP_BEST_K COMP_BEST_K_GAP
OsiSymDivingK	diving_k	-user defined-
OsiSymDivingThreshold	diving_threshold	-user defined-
OsiSymGranularity	granularity	-user defined-
OsiSymTimeLimit	time_limit	-user defined-
OsiSymGapLimit	gap_limit	-user defined-
OsiObjOffset	-	-user defined-
OsiProbName	problem_name	-user defined-

However, as it is seen, only some of the C interface parameters have their matches. If the other parameters are required to be modified, the user can always set them directly by their C interface names, using the overlapping functions: `setSymParam(string, int)`, `setSymParam(string, double)` and `setSymParam(string, string)`. For instance, the `verbosity` parameter can be set, let's say, to 2 either by `setSymParam(OsiSymVerbosity, 2)` or by `setSymParam("verbosity", 2)`. Note that, this flexibility is also supported for parameter querying functions.

Bibliography

- [1] D. Applegate, R. Bixby, V. Chvátal, and W. Cook. CONCORDE TSP solver. <http://www.tsp.gatech.edu/concorde.html>.
- [2] D. Applegate, R. Bixby, V. Chvátal, and W. Cook. On the solution of traveling salesman problems. *Documenta Mathematica*, Extra Volume Proceedings ICM III (1998):645–656, 1998.
- [3] E. Balas, S. Ceria, and G. Cornuéjols. Mixed 0-1 programming by lift-and-project in a branch-and-cut framework. *Management Science*, 42:1229–1246, 1996.
- [4] M. Benchouche, V.-D. Cung, S. Dowaji, B. Le Cun, T. Mautor, and C. Roucairol. Building a parallel branch and bound library. In *Solving Combinatorial Optimization Problems in Parallel, Lecture Notes in Computer Science* **1054**, page 201. Springer, Berlin, 1996.
- [5] V. J. Bowman. On the relationship of the Tchebycheff norm and the efficient frontier of multiple-criteria objectives. In H. Thieriez, editor, *Multiple Criteria Decision Making*, pages 248–258. Springer, Berlin, 1976.
- [6] Q. Chen, M. C. Ferris, and J. T. Linderoth. Fatcop 2.0: Advanced features in an opportunistic mixed integer programming solver. *Annals of Operations Research*, 103:17–32, 2001.
- [7] J. Climaco, C. Ferreira, and M.E. Captivo. Multicriteria integer programming: an overview of different algorithmic approaches. In J. Climaco, editor, *Multicriteria Analysis*, pages 248–258. Springer, Berlin, 1997.
- [8] C. Cordier, H. Marchand, R. Laundy, and L.A. Wolsey. bc-opt: A branch-and-cut code for mixed integer programs. *Mathematical Programming*, 86:335–353, 1999.
- [9] ILOG CPLEX Division. <http://www.cplex.com>.
- [10] J. Eckstein, C.A. Phillips, and W.E. Hart. PICO: An object-oriented framework for parallel branch and bound. Technical Report RRR 40-2000, Rutgers University, 2000.
- [11] M. Ehrgott and X. Gandibleux. A survey and annotated bibliography of multiobjective combinatorial optimization. *OR Spektrum*, 22:425–460, 2000.
- [12] M. Ehrgott and X. Gandibleux. Multiobjective combinatorial optimization—theory, methodology and applications. In M. Ehrgott and X. Gandibleux, editors, *Multiple Criteria Optimization—State of the Art Annotated Bibliographic Surveys*, pages 369–444. Kluwer Academic Publishers, Boston, MA, 2002.

- [13] M. Ehrgott and M.M. Wiecek. Multiobjective programming. In M. Ehrgott, J. Figueira, and S. Greco, editors, *State of the Art of Multiple Criteria Decision Analysis*, Boston, MA, 2004. Kluwer Academic Publishers.
- [14] P. K. Eswaran, A. Ravindran, and H. Moskowitz. Algorithms for nonlinear integer bicriterion problems. *Journal of Optimization Theory and Applications*, 63(2):261–279, 1989.
- [15] A. Geist, A. Beguelin, J. Dongarra, W. Jiang, R. Manchek, and V. Sunderam. *PVM: Parallel Virtual Machine*. The MIT Press, Cambridge, MA, 1994.
- [16] B. Gendron and T. G. Crainic. Parallel branch and bound algorithms: Survey and synthesis. *Operations Research*, 42:1042–1066, 1994.
- [17] A. M. Geoffrion. Proper efficiency and the theory of vector maximization. *Journal of Mathematical Analysis and Applications*, 22:618–630, 1968.
- [18] A.Y. Grama and V. Kumar. State of the art in parallel search techniques for discrete optimization problems. *IEEE Transactions on Knowledge and Data Engineering*, 11:1–8, 1999.
- [19] M. Grötschel, M. Jünger, and G. Reinelt. A cutting plane algorithm for the linear ordering problem. *Operations Research*, 32(6):1195–1220, 1984.
- [20] K. Hoffman and M. Padberg. LP-based combinatorial problem solving. *Annals of Operations Research*, 4:145–194, 1985.
- [21] M. Jünger and S. Thienel. The ABACUS system for branch and cut and price algorithms in integer programming and combinatorial optimization. *Software Practice and Experience*, 30:1325–1352, 2001.
- [22] V. Kumar and V. N. Rao. Parallel depth-first search, part II: Analysis. *International Journal of Parallel Programming*, 16:501–519, 1987.
- [23] J. Linderoth. *Topics in Parallel Integer Optimization*. PhD thesis, School of Industrial and Systems Engineering, Georgia Institute of Technology, Atlanta, GA, 1998.
- [24] R. Lougee-Heimer. The Common Optimization INterface for Operations Research. *IBM Journal of Research and Development*, 47:57–66, 2003.
- [25] A. Makhorin. Introduction to GLPK, 2004. Available from <http://www.gnu.org/software/glpk/glpk.html>.
- [26] A. Martin. Integer programs with block structure. Habilitation Thesis, Technical University of Berlin, Berlin, Germany, 1998.
- [27] G.L. Nemhauser, M.W.P. Savelsbergh, and G.S. Sigismondi. MINTO, a Mixed INTegeR Optimizer. *Operations Research Letters*, 15:47–58, 1994.
- [28] G.L. Nemhauser and L.A. Wolsey. *Integer and Combinatorial Optimization*. Wiley, New York, 1988.
- [29] M. W. Padberg and G. Rinaldi. A branch and cut algorithm for the solution of large scale traveling salesman problems. *SIAM Review*, 33:60–100, 1991.

- [30] T.K. Ralphs and L. Ladányi. *SYMPHONY Version 4.0 User's Manual*, 2004. <http://www.brandandcut.org>.
- [31] T.K. Ralphs, M.J. Saltzman, and M.M. Wiecek. An improved algorithm for biobjective integer programming. To appear in *Annals of Operations Research*, 2005.
- [32] V. N. Rao and V. Kumar. Parallel depth-first search, part I: Implementation. *International Journal of Parallel Programming*, 16:479–499, 1987.
- [33] Y. Shinano, K. Harada, and R. Hirabayashi. A generalized utility for parallel branch and bound algorithms. In *Proceedings of the 1995 Seventh Symposium on Parallel and Distributed Processing*, pages 392–401, Los Alamitos, CA, 1995. IEEE Computer Society Press.
- [34] R. Solanki. Generating the noninferior set in mixed integer biobjective linear programs: an application to a location problem. *Computers and Operations Research*, 18:1–15, 1991.
- [35] S. Tschoke and T. Polzer. *Portable Parallel Branch and Bound Library User Manual: Library Version 2.0*. Department of Computer Science, University of Paderborn.
- [36] Y. Xu, T.K. Ralphs, L. Ladányi, and M.J. Saltzman. ALPS: A framework for implementing parallel search algorithms. In *Proceedings of the Ninth INFORMS Computing Society Conference*, pages 319–334, 2005.

